



MAARS Command

Autonomous AI Enterprise Operating System
by MAARS Global Corporation | Est. 2026

What is MAARS Command?

Imagine you have a company with **499+ AI employees**, each one an expert in something different -- marketing, coding, legal, finance, design, sales, research, and more. MAARS Command is the operating system that manages all of them. You tell it what you want to accomplish (in plain English), and it figures out which AI agents to assign, breaks your goal into steps, runs everything automatically, checks the quality of every result, and delivers the finished work to you.

These agents are organized into **28 specialized network categories** (like departments in a company). They're powered by **33 AI providers with 175,609+ models** (GPT-5, Claude Opus 4.6, Gemini 2.5, Grok-4, DeepSeek R1, Mistral, Perplexity, Groq, Cerebras, SambaNova, and more). Each agent is grounded in the **58,422-skill library** — a curated knowledge pack of SKILL.md documents that is auto-matched and ingested into the agent's brain at creation time, so replies stand on real expertise instead of guesses. The system automatically picks the right AI model for each task, controls costs, ensures quality, and even lets agents collaborate with each other -- all without you lifting a finger.

Multi-Agent Orchestration

Quality Control

Smart Model Router

Developer API Gateway

600+ Models · 33 Providers

58k+ Skill Library

Content Generation

App Builder

Memory System

Voice Commands

Code Explorer

Knowledge Graph

Workflow Builder

Campaign Builder

Integration Hub

Team Builder

Trust Analytics

Real-World Actions

Command Palette

👁 VISION

A world where every organization — from a solo founder to a global enterprise — operates with a reliable, auditable, autonomous AI workforce. Not one chatbot. Not one copilot. An entire workforce, governed like a real company, grounded in verified knowledge, accountable to humans.

🚀 MISSION

Build the operating system that makes autonomous AI workforces safe, provable, and economical. Replace 18–24 months of custom integration work with one platform and one API key. Ship agents that earn trust — through verification, audit, and recovery — instead of agents that just sound confident.

499+

AI Agents
Specialized workers

28

Networks
Team categories

37

Core Systems
Built-in tools

33

LLM Providers
AI companies

600+

AI Models
Via gateway

58,422

Skill Library
Curated SKILL.md
packs

484

API Endpoints
Across 32 route files



The Complete 499-Agent AI Workforce

Every single agent, organized by specialty

Below is every AI agent in MAARS Command, organized by their network category (like departments in a company). Each agent has a unique name, a specific job title (role), a description of what they do, and a list of all their skills (capabilities).



Commander Orion ∞

Commander

Autonomy 1/10

Your AI Commander. The supreme orchestrator of MAARS Infinity. Give it a goal and it will classify, decompose, assign agents across 27 networks, and coordinate execution through the Kernel with verification and governance at every step.

SKILLS:

Task Delegation Strategic Planning Team Orchestration Goal Breakdown Multi-Agent Coordination

TOOLS:

web_search create_task calculate analyze_data send_slack send_email send_sms github_action hubspot_action salesforce_action shopify_action
webhook_action query_tasks update_task query_agent_history product_scan browser_open browser_navigate browser_click browser_fill browser_type
browser_extract browser_screenshot browser_evaluate browser_close browser_see_and_act browser_run_goal



Adrian Stone

Procurement & Vendor Manager

Autonomy 1/10

Procurement & Vendor Manager handling vendor selection, contract negotiation, supply chain optimization, and cost management.

SKILLS:

Vendor Management Contract Negotiation Supply Chain Cost Optimization RFP/RFQ SLA Management

TOOLS:

web_search analyze_data calculate send_email airtable_action query_tasks update_task query_agent_history



Alexandra Reid

Legal Assistant

Autonomy 1/10

Legal assistant helping with contracts, terms of service, privacy policies, and legal document preparation. Familiar with international law and business-favorable contract drafting.

SKILLS:

Contract Drafting Terms of Service Privacy Policies International Law Business-Favorable Contracts Tax Compliance

TOOLS:

web_search send_email query_tasks update_task query_agent_history product_scan



Amara Johnson

HR Specialist

Autonomy 1/10

HR specialist managing hiring, onboarding, employee policies, job descriptions, and building great workplace culture.

SKILLS:

Recruiting Job Descriptions Onboarding HR Policies Employee Relations

TOOLS:

web_search create_task send_email schedule_meeting google_calendar query_tasks update_task query_agent_history product_scan



Axel Brennan

Growth Hacker

Autonomy 1/10

Growth hacker engineering rapid, scalable growth through creative experiments, viral loops, conversion optimization, and data-driven growth strategies.

SKILLS:

Growth Experiments Conversion Optimization Viral Loops A/B Testing Funnel Optimization Product-Led Growth

TOOLS:

web_search analyze_data calculate send_email query_tasks update_task query_agent_history



Benjamin Cole

Financial Analyst

Autonomy 1/10

Financial analyst handling budgets, forecasts, financial models, expense tracking, and money management.

SKILLS:

Budgeting Financial Modeling Forecasting Expense Tracking Profitability Analysis

TOOLS:

calculate analyze_data web_search send_email airtable_action query_tasks update_task query_agent_history product_scan



Cassandra Steele

Chief Strategy Officer

Autonomy 1/10

Chief Strategy Officer developing long-term corporate strategy, competitive positioning, market expansion plans, and strategic partnerships.

SKILLS:

Corporate Strategy Competitive Analysis Market Expansion Strategic Partnerships OKR Frameworks

TOOLS:

web_search analyze_data calculate create_task query_tasks update_task query_agent_history



Catherine Blake

Reputation & PR Manager

Autonomy 1/10

Reputation & PR Manager handling public relations, crisis communications, media relations, and brand reputation management.

SKILLS:

PR Strategy Crisis Management Media Relations Press Releases Reputation Monitoring CSR

TOOLS:

web_search analyze_data send_email send_slack query_tasks update_task query_agent_history

Damien Voss Cybersecurity Officer Autonomy 1/10

Cybersecurity officer protecting digital assets, conducting security audits, managing threats, and ensuring data protection compliance across the organization.

SKILLS:

Security Audits Threat Analysis Compliance Incident Response Data Protection Vulnerability Assessment

TOOLS:

web_search analyze_data create_task query_tasks update_task query_agent_history

Derek Huang SEO Specialist Autonomy 1/10

SEO expert optimizing websites for search engines, improving rankings, and driving organic traffic through proven strategies.

SKILLS:

Keyword Research On-Page SEO Technical SEO Link Building SEO Audits

TOOLS:

web_search analyze_data airtable_action query_tasks update_task query_agent_history product_scan

Diana Morales Operations Manager Autonomy 1/10

Operations manager streamlining business processes, managing supply chains, optimizing resources, and ensuring smooth day-to-day business operations.

SKILLS:

Process Optimization Supply Chain Resource Management SOP Development KPI Tracking Vendor Management

TOOLS:

web_search analyze_data calculate create_task airtable_action google_calendar shopify_action webhook_action query_tasks update_task query_agent_history

Dr. Clara Voss Research Specialist Autonomy 1/10

Research specialist conducting deep research, competitor analysis, market studies, and comprehensive reports on any topic.

SKILLS:

Market Research Competitor Analysis Industry Reports Trend Analysis Due Diliigence

TOOLS:

web_search analyze_data calculate github_action query_tasks update_task query_agent_history product_scan

Dr. Eleanor Shaw Knowledge Architect Autonomy 1/10

Knowledge Architect building knowledge management systems, taxonomies, documentation frameworks, and organizational learning infrastructure.

SKILLS:

Knowledge Management Taxonomies Documentation Organizational Learning Information Architecture

TOOLS:

web_search analyze_data create_task query_tasks update_task query_agent_history

Dr. Luca Bernstein AI Optimization Specialist Autonomy 1/10

AI optimization specialist fine-tuning AI systems, optimizing model performance, managing AI costs, and ensuring AI outputs meet quality standards.

SKILLS:

AI Optimization Prompt Engineering Cost Reduction Model Selection Quality Assurance Performance Tuning

TOOLS:

web_search analyze_data calculate query_tasks update_task query_agent_history

Elena Kovacs Product Manager Autonomy 1/10

Product Manager defining product vision, roadmaps, user stories, feature prioritization, and cross-functional product execution.

SKILLS:

Product Roadmaps User Stories Feature Prioritization Sprint Planning Product Analytics

TOOLS:

web_search analyze_data create_task github_action query_tasks update_task query_agent_history

Ethan Yates Data Analyst Autonomy 1/10

Data analyst turning raw numbers into actionable insights through analysis, visualization, and clear reporting.

SKILLS:

Data Analysis Reporting Dashboards SQL/Excel Business Intelligence

TOOLS:

web_search calculate analyze_data airtable_action google_calendar hubspot_action salesforce_action shopify_action query_tasks update_task query_agent_history product_scan

Felix Romano Graphic Designer Autonomy 1/10

Graphic designer who creates logos, brand assets, banners, and visual content using AI image generation. Designs presentations, social graphics, and marketing materials.

SKILLS:

Image Creation Logo Design Brand Assets Social Graphics Marketing Materials

TOOLS:

web_search search_gif query_tasks update_task query_agent_history product_scan

Isla Fernandez Social Media Manager Autonomy 1/10

Social media manager creating viral content, growing followers, managing communities, and building brand presence. Posts content to Facebook/Instagram/TikTok with geographic boost options.

SKILLS:

Content Creation Community Management Growth Strategies Viral Content Platform Optimization Geographic Boost Ad Management

TOOLS:

web_search analyze_data search_gif send_slack query_tasks update_task query_agent_history product_scan

Jasper Wells Email Marketing Specialist Autonomy 1/10

Email marketing specialist crafting sequences, newsletters, campaigns, and automations that nurture leads and drive sales.

SKILLS:

Email Sequences Newsletter Design Automation Flows Subject Lines Campaign Strategy

TOOLS:

web_search send_email send_gmail query_tasks update_task query_agent_history product_scan

Kai Nakamoto App Developer Autonomy 1/10

Full-stack app developer building web and mobile applications with clean code, scalable architecture, and modern frameworks.

SKILLS:

Web Development Mobile Apps API Development React/Node.js Full-Stack Coding

TOOLS:

web_search calculate github_action query_tasks update_task query_agent_history product_scan

Layla Mansouri Localization & Global Expansion Specialist Autonomy 1/10

Localization & Global Expansion Specialist handling internationalization, cultural adaptation, translation management, and market entry strategy.

SKILLS:

Localization Global Expansion Cultural Adaptation Translation Management Market Entry International Compliance

TOOLS:

web_search analyze_data calculate send_email query_tasks update_task query_agent_history

Luna Bergström Web Designer Autonomy 1/10

Creative web designer specializing in stunning UI/UX, wireframes, landing pages, and visual designs that captivate and convert.

SKILLS:

UI/UX Design Wireframing Landing Pages Visual Design Brand Identity

TOOLS:

web_search search_gif query_tasks update_task query_agent_history product_scan

Marco De Luca 3D Visualization Specialist Autonomy 1/10

3D Visualization Specialist creating product renders, architectural visualizations, 3D mockups, and immersive visual content.

SKILLS:

3D Rendering Product Visualization Motion Graphics AR/VR Architectural Viz Packaging Design

TOOLS:

web_search search_gif query_tasks update_task query_agent_history

Marcus Drake Sales Representative Autonomy 1/10

Sales expert crafting pitches, handling objections, writing proposals, closing deals, and sending follow-up emails and messages via multiple channels.

SKILLS:

Sales Pitches Proposal Writing Objection Handling Lead Nurturing Closing Techniques Email Outreach Multi-Channel Follow-up

TOOLS:

web_search calculate send_email send_gmail send_sms schedule_meeting google_calendar hubspot_action salesforce_action query_tasks update_task query_agent_history product_scan

Maximilian Wolfe Revenue Strategist Autonomy 1/10

Revenue optimization strategist maximizing profitability through pricing strategies, revenue modeling, upselling frameworks, and monetization optimization.

SKILLS:

Pricing Strategy Revenue Modeling Monetization Upselling Financial Forecasting Profitability Analysis

TOOLS:

web_search analyze_data calculate send_email hubspot_action salesforce_action query_tasks update_task query_agent_history

Maya Thompson Customer Service Rep Autonomy 1/10

Customer service specialist handling inquiries, resolving issues via email, SMS, and messaging apps. Ensures every customer feels valued and heard.

SKILLS:

Customer Support Complaint Resolution Support Scripts FAQ Creation Multi-Channel Communication

TOOLS:

web_search create_task send_email send_gmail send_sms query_tasks update_task query_agent_history product_scan

Nadia Kessler Personal Secretary Autonomy 1/10

Your dedicated personal secretary handling appointments, calendars, to-do lists, reminders, sending emails and messages, scheduling appointments, and passing messages. Receives results from Commander.

SKILLS:

Calendar Management To-Do Lists Appointment Scheduling Email Sending Message Delivery Commander Handoff

TOOLS:

create_task calculate send_email send_gmail schedule_meeting google_calendar send_sms query_tasks update_task query_agent_history product_scan

Nathan Cross Project Manager Autonomy 1/10

Project manager keeping teams on track with timelines, milestones, task delegation, and seamless project execution.

SKILLS:

Project Planning Timeline Management Task Delegation Agile/Scrum Progress Tracking

TOOLS:

create_task calculate analyze_data send_slack google_calendar airtable_action query_tasks update_task query_agent_history product_scan

Nikolai Volkov Data Engineer Autonomy 1/10

Data Engineer building data pipelines, managing databases, ensuring data quality, and enabling analytics infrastructure.

SKILLS:

Data Pipelines Database Architecture ETL/ELT Data Quality Query Optimization Analytics Infrastructure

TOOLS:

web_search analyze_data calculate github_action airtable_action query_tasks update_task query_agent_history

Olivia Sinclair Content Writer Autonomy 1/10

Content writer producing engaging blog posts, articles, newsletters, and long-form content that educates and entertains.

SKILLS:

Blog Writing Articles Newsletters Whitepapers Thought Leadership

TOOLS:

web_search search_gif query_tasks update_task query_agent_history product_scan

Professor James Whitfield Ethics & Risk Governance Officer Autonomy 1/10

Ethics & Risk Governance Officer overseeing ethical AI use, risk management, corporate governance, and responsible business practices.

SKILLS:

Ethical AI Governance Risk Management Corporate Governance Bias Auditing ESG Reporting Compliance

TOOLS:

web_search analyze_data create_task query_tasks update_task query_agent_history

Richard Ashworth Investor Relations Manager Autonomy 1/10

Investor Relations Manager managing stakeholder communications, fundraising strategy, pitch decks, and investor reporting.

SKILLS:

Pitch Decks Investor Reports Fundraising Valuation Due Diligence Board Presentations

TOOLS:

web_search analyze_data calculate send_email query_tasks update_task query_agent_history

Riley Chen Video Content Specialist Autonomy 1/10

Video content specialist who creates commercials, ads, short films, and promotional videos using AI. Also plans scripts, storyboards, and video marketing strategies.

SKILLS:

Video Creation Commercials & Ads Video Scripts Storyboarding YouTube Strategy

TOOLS:

web_search search_gif query_tasks update_task query_agent_history product_scan

Scarlett Monroe Copywriter Autonomy 1/10

Persuasive copywriter crafting compelling sales copy, website content, product descriptions, and words that sell.

SKILLS:

Sales Copy Website Copy Headlines & Taglines Product Descriptions Persuasive Writing

TOOLS:

web_search send_email query_tasks update_task query_agent_history product_scan

Serena Okafor Automation Engineer Autonomy 1/10

Automation engineer designing workflows, building process automations, integrating systems, and eliminating manual bottlenecks across business operations.

SKILLS:

Workflow Design Process Automation System Integration Data Pipelines API Integration Efficiency Optimization

TOOLS:

web_search analyze_data create_task github_action airtable_action webhook_action hubspot_action salesforce_action shopify_action query_tasks update_task query_agent_history

Sofia Reyes Customer Experience Architect Autonomy 1/10

Customer Experience Architect designing end-to-end customer journeys, loyalty programs, feedback systems, and experience optimization.

SKILLS:

Journey Mapping CX Strategy Loyalty Programs NPS/CSAT Feedback Systems Personalization

TOOLS:

web_search analyze_data create_task send_email query_tasks update_task query_agent_history

Valentina Cruz Brand Architect Autonomy 1/10

Brand Architect defining brand identity, voice, visual systems, brand guidelines, and ensuring consistency across all touchpoints.

SKILLS:

Brand Identity Brand Guidelines Visual Systems Brand Voice Brand Architecture Brand Monitoring

TOOLS:

web_search search_gif analyze_data query_tasks update_task query_agent_history

Victor Ashford Business Strategist Autonomy 1/10

Business strategist analyzing markets, competitors, and opportunities to develop winning strategies and actionable business plans.

SKILLS:

Business Strategy Competitive Analysis Market Research Business Planning Growth Strategy

TOOLS:

web_search calculate analyze_data airtable_action hubspot_action salesforce_action shopify_action webhook_action query_tasks update_task query_agent_history product_scan

Victoria Harrington Compliance Officer Autonomy 1/10

Compliance officer ensuring regulatory adherence, managing risk, overseeing audits, and maintaining organizational governance across all jurisdictions.

SKILLS:

Regulatory Compliance Audit Management Risk Assessment Policy Development GDPR/CCPA Governance

TOOLS:

web_search analyze_data create_task send_email query_tasks update_task query_agent_history

Yuki Tanaka UX Researcher Autonomy 1/10

UX Researcher conducting user research, usability testing, persona development, and data-driven design recommendations.

SKILLS:

User Research Usability Testing Personas Journey Maps A/B Testing Design Recommendations

TOOLS:

web_search analyze_data create_task query_tasks update_task query_agent_history



Zara Mitchell

Marketing Specialist

Autonomy 1/10

Creative marketing specialist crafting campaigns, social media content, ad copy, and brand messaging that converts audiences into customers.

SKILLS:

Social Media MarketingAd CopywritingEmail CampaignsContent StrategyBrand Messaging

TOOLS:

web_searchanalyze_datasend_emailsearch_gifsend_slackhubspot_actionshopify_actionquery_tasksupdate_taskquery_agent_historyproduct_scan



Adrian Stone

Procurement & Vendor Manager

Autonomy 2/10

Procurement & Vendor Manager handling vendor selection, contract negotiation, supply chain optimization, and cost management.

SKILLS:

Vendor ManagementContract NegotiationSupply ChainCost OptimizationRFP/RFQSLA Management



Alexandra Reid

Legal Assistant

Autonomy 2/10

Legal assistant helping with contracts, terms of service, privacy policies, and legal document preparation. Familiar with international law and business-favorable contract drafting.

SKILLS:

Contract DraftingTerms of ServicePrivacy PoliciesInternational LawBusiness-Favorable ContractsTax Compliance



Amara Johnson

HR Specialist

Autonomy 2/10

HR specialist managing hiring, onboarding, employee policies, job descriptions, and building great workplace culture.

SKILLS:

RecruitingJob DescriptionsOnboardingHR PoliciesEmployee Relations



Axel Brennan

Growth Hacker

Autonomy 2/10

Growth hacker engineering rapid, scalable growth through creative experiments, viral loops, conversion optimization, and data-driven growth strategies.

SKILLS:

Growth ExperimentsConversion OptimizationViral LoopsA/B TestingFunnel OptimizationProduct-Led Growth



Benjamin Cole

Financial Analyst

Autonomy 2/10

Financial analyst handling budgets, forecasts, financial models, expense tracking, and money management.

SKILLS:

BudgetingFinancial ModelingForecastingExpense TrackingProfitability Analysis



Cassandra Steele

Chief Strategy Officer

Autonomy 2/10

Chief Strategy Officer developing long-term corporate strategy, competitive positioning, market expansion plans, and strategic partnerships.

SKILLS:

Corporate StrategyCompetitive AnalysisMarket ExpansionStrategic PartnershipsOKR Frameworks



Catherine Blake

Reputation & PR Manager

Autonomy 2/10

Reputation & PR Manager handling public relations, crisis communications, media relations, and brand reputation management.

SKILLS:

PR StrategyCrisis ManagementMedia RelationsPress ReleasesReputation MonitoringCSR



Commander Orion

Commander

Autonomy 4/10

Your AI Commander. The supreme orchestrator of MAARS Infinity. Give it a goal and it will classify, decompose, assign agents across 28 networks, and coordinate execution through the Kernel with verification and governance at every step.

SKILLS:

Task DelegationStrategic PlanningTeam OrchestrationGoal BreakdownMulti-Agent Coordination



Damien Voss

Cybersecurity Officer

Autonomy 2/10

Cybersecurity officer protecting digital assets, conducting security audits, managing threats, and ensuring data protection compliance across the organization.

SKILLS:

Security AuditsThreat AnalysisComplianceIncident ResponseData ProtectionVulnerability Assessment



Derek Huang

SEO Specialist

Autonomy 2/10

SEO expert optimizing websites for search engines, improving rankings, and driving organic traffic through proven strategies.

SKILLS:

Keyword ResearchOn-Page SEOTechnical SEOLink BuildingSEO Audits

	Diana Morales Operations Manager Autonomy 2/10	Operations manager streamlining business processes, managing supply chains, optimizing resources, and ensuring smooth day-to-day business operations. SKILLS: Process OptimizationSupply ChainResource ManagementSOP DevelopmentKPI TrackingVendor Management
	Dr. Clara Voss Research Specialist Autonomy 2/10	Research specialist conducting deep research, competitor analysis, market studies, and comprehensive reports on any topic. SKILLS: Market ResearchCompetitor AnalysisIndustry ReportsTrend AnalysisDue Diligence
	Dr. Eleanor Shaw Knowledge Architect Autonomy 2/10	Knowledge Architect building knowledge management systems, taxonomies, documentation frameworks, and organizational learning infrastructure. SKILLS: Knowledge ManagementTaxonomiesDocumentationOrganizational LearningInformation Architecture
	Dr. Luca Bernstein AI Optimization Specialist Autonomy 2/10	AI optimization specialist fine-tuning AI systems, optimizing model performance, managing AI costs, and ensuring AI outputs meet quality standards. SKILLS: AI OptimizationPrompt EngineeringCost ReductionModel SelectionQuality AssurancePerformance Tuning
	Elena Kovacs Product Manager Autonomy 2/10	Product Manager defining product vision, roadmaps, user stories, feature prioritization, and cross-functional product execution. SKILLS: Product RoadmapsUser StoriesFeature PrioritizationSprint PlanningProduct Analytics
	Ethan Yates Data Analyst Autonomy 2/10	Data analyst turning raw numbers into actionable insights through analysis, visualization, and clear reporting. SKILLS: Data AnalysisReportingDashboardsSQL/ExcelBusiness Intelligence
	Felix Romano Graphic Designer Autonomy 2/10	Graphic designer who creates logos, brand assets, banners, and visual content using AI image generation. Designs presentations, social graphics, and marketing materials. SKILLS: Image CreationLogo DesignBrand AssetsSocial GraphicsMarketing Materials
	Isla Fernandez Social Media Manager Autonomy 2/10	Social media manager creating viral content, growing followers, managing communities, and building brand presence. Posts content to Facebook/Instagram/TikTok with geographic boost options. SKILLS: Content CreationCommunity ManagementGrowth StrategiesViral ContentPlatform OptimizationGeographic BoostAd Management
	Jasper Wells Email Marketing Specialist Autonomy 2/10	Email marketing specialist crafting sequences, newsletters, campaigns, and automations that nurture leads and drive sales. SKILLS: Email SequencesNewsletter DesignAutomation FlowsSubject LinesCampaign Strategy
	Kai Nakamoto App Developer Autonomy 2/10	Full-stack app developer building web and mobile applications with clean code, scalable architecture, and modern frameworks. SKILLS: Web DevelopmentMobile AppsAPI DevelopmentReact/Node.jsFull-Stack Coding
	Layla Mansouri Localization & Global Expansion Specialist Autonomy 2/10	Localization & Global Expansion Specialist handling internationalization, cultural adaptation, translation management, and market entry strategy. SKILLS: LocalizationGlobal ExpansionCultural AdaptationTranslation ManagementMarket EntryInternational Compliance



Luna Bergström Web Designer Autonomy 2/10

Creative web designer specializing in stunning UI/UX, wireframes, landing pages, and visual designs that captivate and convert.

SKILLS:

UI/UX Design Wireframing Landing Pages Visual Design Brand Identity



Marco De Luca 3D Visualization Specialist Autonomy 2/10

3D Visualization Specialist creating product renders, architectural visualizations, 3D mockups, and immersive visual content.

SKILLS:

3D Rendering Product Visualization Motion Graphics AR/VR Architectural Viz Packaging Design



Marcus Drake Sales Representative Autonomy 2/10

Sales expert crafting pitches, handling objections, writing proposals, closing deals, and sending follow-up emails and messages via multiple channels.

SKILLS:

Sales Pitches Proposal Writing Objection Handling Lead Nurturing Closing Techniques Email Outreach Multi-Channel Follow-up



Maximilian Wolfe Revenue Strategist Autonomy 2/10

Revenue optimization strategist maximizing profitability through pricing strategies, revenue modeling, upselling frameworks, and monetization optimization.

SKILLS:

Pricing Strategy Revenue Modeling Monetization Upselling Financial Forecasting Profitability Analysis



Maya Thompson Customer Service Rep Autonomy 2/10

Customer service specialist handling inquiries, resolving issues via email, SMS, and messaging apps. Ensures every customer feels valued and heard.

SKILLS:

Customer Support Complaint Resolution Support Scripts FAQ Creation Multi-Channel Communication



Nadia Kessler Personal Secretary Autonomy 3/10

Your dedicated personal secretary handling appointments, calendars, to-do lists, reminders, sending emails and messages, scheduling appointments, and passing messages. Receives results from Commander.

SKILLS:

Calendar Management To-Do Lists Appointment Scheduling Email Sending Message Delivery Commander Handoff



Nathan Cross Project Manager Autonomy 2/10

Project manager keeping teams on track with timelines, milestones, task delegation, and seamless project execution.

SKILLS:

Project Planning Timeline Management Task Delegation Agile/Scrum Progress Tracking



Nikolai Volkov Data Engineer Autonomy 2/10

Data Engineer building data pipelines, managing databases, ensuring data quality, and enabling analytics infrastructure.

SKILLS:

Data Pipelines Database Architecture ETL/ELT Data Quality Query Optimization Analytics Infrastructure



Olivia Sinclair Content Writer Autonomy 2/10

Content writer producing engaging blog posts, articles, newsletters, and long-form content that educates and entertains.

SKILLS:

Blog Writing Articles Newsletters Whitepapers Thought Leadership



Professor James Whitfield Ethics & Risk Governance Officer Autonomy 2/10

Ethics & Risk Governance Officer overseeing ethical AI use, risk management, corporate governance, and responsible business practices.

SKILLS:

Ethical AI Governance Risk Management Corporate Governance Bias Auditing ESG Reporting Compliance



Richard Ashworth Investor Relations Manager Autonomy 2/10

Investor Relations Manager managing stakeholder communications, fundraising strategy, pitch decks, and investor reporting.

SKILLS:

Pitch Decks Investor Reports Fundraising Valuation Due Diligence Board Presentations



Riley Chen

Video Content Specialist

Autonomy 2/10

Video content specialist who creates commercials, ads, short films, and promotional videos using AI. Also plans scripts, storyboards, and video marketing strategies.

SKILLS:

Video Creation

Commercials & Ads

Video Scripts

Storyboarding

YouTube Strategy



Scarlett Monroe

Copywriter

Autonomy 2/10

Persuasive copywriter crafting compelling sales copy, website content, product descriptions, and words that sell.

SKILLS:

Sales Copy

Website Copy

Headlines & Taglines

Product Descriptions

Persuasive Writing



Serena Okafor

Automation Engineer

Autonomy 2/10

Automation engineer designing workflows, building process automations, integrating systems, and eliminating manual bottlenecks across business operations.

SKILLS:

Workflow Design

Process Automation

System Integration

Data Pipelines

API Integration

Efficiency Optimization



Sofia Reyes

Customer Experience Architect

Autonomy 2/10

Customer Experience Architect designing end-to-end customer journeys, loyalty programs, feedback systems, and experience optimization.

SKILLS:

Journey Mapping

CX Strategy

Loyalty Programs

NPS/CSAT

Feedback Systems

Personalization



Valentina Cruz

Brand Architect

Autonomy 2/10

Brand Architect defining brand identity, voice, visual systems, brand guidelines, and ensuring consistency across all touchpoints.

SKILLS:

Brand Identity

Brand Guidelines

Visual Systems

Brand Voice

Brand Architecture

Brand Monitoring



Victor Ashford

Business Strategist

Autonomy 2/10

Business strategist analyzing markets, competitors, and opportunities to develop winning strategies and actionable business plans.

SKILLS:

Business Strategy

Competitive Analysis

Market Research

Business Planning

Growth Strategy



Victoria Harrington

Compliance Officer

Autonomy 2/10

Compliance officer ensuring regulatory adherence, managing risk, overseeing audits, and maintaining organizational governance across all jurisdictions.

SKILLS:

Regulatory Compliance

Audit Management

Risk Assessment

Policy Development

GDPR/CCPA

Governance



Yuki Tanaka

UX Researcher

Autonomy 2/10

UX Researcher conducting user research, usability testing, persona development, and data-driven design recommendations.

SKILLS:

User Research

Usability Testing

Personas

Journey Maps

A/B Testing

Design Recommendations



Zara Mitchell

Marketing Specialist

Autonomy 2/10

Creative marketing specialist crafting campaigns, social media content, ad copy, and brand messaging that converts audiences into customers.

SKILLS:

Social Media Marketing

Ad Copywriting

Email Campaigns

Content Strategy

Brand Messaging



Industry Specific

40 agents



Agri Economics Agent

Agricultural Economist

Autonomy 3/10

Analyzes agricultural economics including commodity pricing, subsidies, and market dynamics.

SKILLS:

Commodity Analysis Subsidy Impact Market Dynamics Price Forecasting Trade Policy



Agri Supply Chain Agent

Agricultural Supply Chain Manager

Autonomy 3/10

Manages agricultural supply chains from farm to market including logistics and quality control.

SKILLS:

Supply Chain Farm to Market Quality Control Cold Chain Distribution



Audience Growth Agent

Audience Growth Manager

Autonomy 3/10

Grows media audiences through engagement strategies, content optimization, and community building.

SKILLS:

Audience Growth Engagement Strategy Content Optimization Community Building Subscriber Growth



Brand Safety Agent

Brand Safety Manager

Autonomy 3/10

Ensures brand safety in media placement including content adjacency, reputation protection, and compliance.

SKILLS:

Brand Safety Content Adjacency Reputation Protection Compliance Risk Monitoring



Catalog Optimization Agent

Catalog Optimizer

Autonomy 3/10

Optimizes product catalogs including taxonomy, search, filtering, and product information management.

SKILLS:

Catalog Optimization Taxonomy Search Optimization PIM Data Quality



Climate Yield Agent

Climate-Yield Analyst

Autonomy 3/10

Analyzes climate impact on crop yields and develops adaptation strategies.

SKILLS:

Climate Analysis Yield Impact Adaptation Strategy Weather Modeling Risk Assessment



Clinical Research Agent

Clinical Research Specialist

Autonomy 2/10

Supports clinical research including study design, protocol development, and regulatory compliance.

SKILLS:

Study Design Protocol Development Regulatory Compliance Data Analysis Literature Review



Commerce Analytics Agent

Commerce Analyst

Autonomy 3/10

Analyzes commerce data including sales trends, customer behavior, and revenue optimization.

SKILLS:

Sales Analytics Customer Behavior Revenue Analysis Basket Analysis Trend Detection



Content Distribution Agent

Content Distribution Manager

Autonomy 3/10

Manages content distribution across platforms, optimizes reach, and tracks performance.

SKILLS:

Content Distribution Platform Management Reach Optimization Performance Tracking Scheduling



Content Personalization Agent

Content Personalizer

Autonomy 3/10

Personalizes educational content based on learner profiles, progress, and preferences.

SKILLS:

Content Personalization Learner Profiles Adaptive Content Progress Tracking Recommendation



Crop Systems Agent

Crop Systems Specialist

Autonomy 3/10

Optimizes crop production systems including planting, irrigation, and harvest optimization.

SKILLS:

Crop Planning Irrigation Optimization Harvest Planning Yield Modeling Soil Analysis



Curriculum Strategy Agent

Curriculum Strategist

Autonomy 3/10

Develops curriculum strategies aligned with learning outcomes, standards, and market needs.

SKILLS:

Curriculum Design Learning Outcomes Standards Alignment Market Needs Content Strategy

	Delivery Network Agent	Delivery Network Manager	Autonomy 3/10
Manages delivery networks including last-mile optimization, hub-and-spoke design, and SLA management.			
SKILLS:			
Delivery OptimizationLast-MileHub-and-SpokeSLA ManagementPerformance Tracking			
	Drug Discovery Support Agent	Drug Discovery Analyst	Autonomy 2/10
Supports drug discovery research including target identification, literature analysis, and pipeline tracking.			
SKILLS:			
Target IdentificationLiterature AnalysisPipeline TrackingCompound AnalysisResearch Synthesis			
	E-commerce Conversion Agent	E-commerce CRO Specialist	Autonomy 3/10
Optimizes e-commerce conversion rates through UX improvement, checkout optimization, and cart recovery.			
SKILLS:			
Conversion OptimizationCheckout UXCart RecoveryProduct PagesMobile Commerce			
	Education Analytics Agent	Education Analyst	Autonomy 3/10
Analyzes educational data including learning outcomes, engagement metrics, and program effectiveness.			
SKILLS:			
Learning AnalyticsEngagement MetricsProgram EffectivenessPredictive ModelsReporting			
	Energy Market Agent	Energy Market Analyst	Autonomy 3/10
Analyzes energy markets including pricing, supply/demand dynamics, and regulatory impacts.			
SKILLS:			
Market AnalysisPrice ForecastingSupply/DemandRegulatory ImpactTrading Insights			
	Energy Risk Agent	Energy Risk Analyst	Autonomy 3/10
Assesses energy sector risks including price volatility, regulatory changes, and supply disruptions.			
SKILLS:			
Risk AssessmentPrice VolatilityRegulatory RiskSupply RiskHedging Strategy			
	Equipment Reliability Agent	Equipment Reliability Engineer	Autonomy 3/10
Manages equipment reliability including predictive maintenance, failure analysis, and uptime optimization.			
SKILLS:			
Predictive MaintenanceFailure AnalysisUptime OptimizationAsset ManagementReliability			
	Fleet Planning Agent	Fleet Planner	Autonomy 3/10
Plans fleet operations including vehicle allocation, maintenance scheduling, and capacity management.			
SKILLS:			
Fleet PlanningVehicle AllocationMaintenance SchedulingCapacity ManagementUtilization			
	Food Distribution Agent	Food Distribution Manager	Autonomy 3/10
Optimizes food distribution networks including routing, storage, and shelf-life management.			
SKILLS:			
Distribution OptimizationRoute PlanningStorage ManagementShelf-LifeWaste Reduction			
	Grid Optimization Agent	Grid Optimization Specialist	Autonomy 3/10
Optimizes energy grid operations including load balancing, distribution, and efficiency improvements.			
SKILLS:			
Grid OptimizationLoad BalancingDistributionEfficiencyDemand Response			
	Healthcare Operations Agent	Healthcare Ops Manager	Autonomy 3/10
Optimizes healthcare operations including scheduling, resource allocation, and workflow efficiency.			
SKILLS:			
Healthcare OpsSchedulingResource AllocationWorkflow EfficiencyPatient Flow			
	Infrastructure Energy Agent	Energy Infrastructure Manager	Autonomy 3/10
Manages energy infrastructure planning, maintenance, and modernization projects.			
SKILLS:			
Infrastructure PlanningMaintenanceModernizationAsset ManagementReliability			

	Learning Systems Agent Learning Systems Designer Autonomy 3/10
Designs learning systems including LMS architecture, adaptive learning, and assessment frameworks.	
SKILLS:	
LMS Design Adaptive Learning Assessment Frameworks Learning Paths Engagement	
	Logistics Reliability Agent Logistics Reliability Manager Autonomy 3/10
Ensures logistics reliability including on-time delivery, damage prevention, and contingency planning.	
SKILLS:	
Reliability Management On-Time Delivery Damage Prevention Contingency Planning SLA Compliance	
	Manufacturing Cost Agent Manufacturing Cost Analyst Autonomy 3/10
Analyzes manufacturing costs including material, labor, overhead, and process efficiency.	
SKILLS:	
Cost Analysis Material Costs Labor Analysis Overhead Process Efficiency	
	Marketplace Operations Agent Marketplace Operations Manager Autonomy 3/10
Manages marketplace operations including seller management, trust systems, and dispute resolution.	
SKILLS:	
Marketplace Ops Seller Management Trust Systems Dispute Resolution Policy Enforcement	
	Media Economics Agent Media Economist Autonomy 3/10
Analyzes media economics including advertising revenue, subscription models, and content ROI.	
SKILLS:	
Media Economics Ad Revenue Subscription Models Content ROI Monetization	
	Media Strategy Agent Media Strategist Autonomy 3/10
Develops media strategy including content planning, distribution, and audience development.	
SKILLS:	
Media Strategy Content Planning Distribution Strategy Audience Development Platform Strategy	
	Medical Knowledge Agent Medical Knowledge Manager Autonomy 2/10
Manages medical knowledge including clinical guidelines, drug information, and medical literature.	
SKILLS:	
Clinical Guidelines Drug Information Medical Literature Knowledge Updates Evidence Review	
	Merchandising Agent Merchandising Specialist Autonomy 3/10
Manages product merchandising including assortment planning, visual merchandising, and promotional strategy.	
SKILLS:	
Assortment Planning Visual Merchandising Promotion Strategy Planogram Category Management	
	Production Planning Agent Production Planner Autonomy 3/10
Plans manufacturing production schedules, capacity allocation, and resource optimization.	
SKILLS:	
Production Planning Capacity Allocation Resource Optimization Scheduling Demand Matching	
	Quality Manufacturing Agent Manufacturing Quality Manager Autonomy 3/10
Manages manufacturing quality including Six Sigma, statistical process control, and defect prevention.	
SKILLS:	
Quality Management Six Sigma SPC Defect Prevention Quality Audits	
	Regulatory Health Agent Health Regulatory Specialist Autonomy 2/10
Manages healthcare regulatory compliance including FDA, EMA, and health authority requirements.	
SKILLS:	
FDA Compliance EMA Requirements Health Authority Regulatory Submissions Compliance Monitoring	
	Renewable Systems Agent Renewable Energy Specialist Autonomy 3/10
Manages renewable energy systems including solar, wind, and storage optimization.	
SKILLS:	
Solar/Wind Storage Optimization Integration Performance Monitoring Forecasting	



Route Optimization Agent **Route Optimizer** Autonomy 3/10

Optimizes transportation routes for cost, time, and efficiency across logistics networks.

SKILLS:

- Route Optimization
- Cost Minimization
- Time Optimization
- Multi-Modal
- Dynamic Routing



Student Success Agent **Student Success Manager** Autonomy 3/10

Manages student success programs including retention, engagement, and outcome tracking.

SKILLS:

- Student Retention
- Engagement Programs
- Outcome Tracking
- Intervention Design
- Success Metrics



Supply Optimization Agent **Supply Chain Optimizer** Autonomy 3/10

Optimizes manufacturing supply chains including procurement, inventory, and just-in-time delivery.

SKILLS:

- Supply Optimization
- JIT Delivery
- Inventory Management
- Procurement
- Lead Time



Transport Economics Agent **Transport Economist** Autonomy 3/10

Analyzes transportation economics including cost per mile, fuel efficiency, and pricing models.

SKILLS:

- Transport Economics
- Cost Analysis
- Fuel Efficiency
- Pricing Models
- Market Rates



AI/ML Engineering Agent AI/ML Engineer Autonomy 3/10

Develops and deploys machine learning models, manages training pipelines, and optimizes model serving.

SKILLS:

- ML Engineering
- Model Training
- Pipeline Design
- Model Serving
- Feature Engineering



API Design Agent API Architect Autonomy 3/10

Designs RESTful and GraphQL APIs with proper versioning, documentation, authentication, and rate limiting.

SKILLS:

- API Design
- REST/GraphQL
- Documentation
- Versioning
- Authentication



Backend Engineering Agent Backend Engineer Autonomy 3/10

Designs and builds backend systems including APIs, databases, microservices, and server architecture.

SKILLS:

- API Development
- Database Design
- Microservices
- Server Architecture
- Performance Optimization



CI/CD Agent CI/CD Engineer Autonomy 3/10

Designs and maintains continuous integration and deployment pipelines with quality gates and rollback capabilities.

SKILLS:

- CI/CD Pipelines
- Quality Gates
- Rollback
- Build Optimization
- Release Automation



Cloud Deployment Agent Cloud Engineer Autonomy 3/10

Manages cloud infrastructure deployment across providers, optimizes costs, and ensures high availability.

SKILLS:

- Cloud Deployment
- Multi-Cloud
- Cost Optimization
- High Availability
- Auto-Scaling



Data Engineering Agent Data Engineer Autonomy 3/10

Builds data pipelines, manages data warehouses, and ensures data quality across analytics infrastructure.

SKILLS:

- Data Pipelines
- ETL/ELT
- Data Warehouse
- Data Quality
- Stream Processing



Database Architecture Agent Database Architect Autonomy 3/10

Designs database schemas, optimizes queries, manages migrations, and selects appropriate database technologies.

SKILLS:

- Schema Design
- Query Optimization
- Migration Management
- Technology Selection
- Replication



DevOps Agent DevOps Engineer Autonomy 3/10

Manages infrastructure as code, container orchestration, monitoring, and deployment automation.

SKILLS:

- Infrastructure as Code
- Containers
- Monitoring
- Deployment Automation
- Configuration Management



Documentation Engineering Agent Documentation Engineer Autonomy 2/10

Creates and maintains technical documentation including API docs, architecture decisions, and runbooks.

SKILLS:

- Technical Writing
- API Documentation
- Architecture Decisions
- Runbooks
- Knowledge Base



Frontend Engineering Agent Frontend Engineer Autonomy 3/10

Builds responsive, accessible user interfaces using modern frameworks, component systems, and design systems.

SKILLS:

- UI Development
- Responsive Design
- Component Systems
- Accessibility
- Performance



Infrastructure as Code Agent IaC Engineer Autonomy 3/10

Manages infrastructure definitions using Terraform, Pulumi, or CloudFormation with proper state management.

SKILLS:

- Terraform
- CloudFormation
- State Management
- Module Design
- Drift Detection



Integration Engineering Agent Integration Engineer Autonomy 3/10

Builds and maintains integrations with external services, APIs, and third-party systems.

SKILLS:

- API Integration
- Webhook Design
- Data Sync
- Protocol Translation
- Error Handling

	Mobile Engineering Agent Mobile Engineer Autonomy 3/10 Develops native and cross-platform mobile applications with optimal UX, performance, and store compliance. SKILLS: Mobile DevelopmentCross-PlatformNative AppsStore ComplianceMobile UX
	Observability Engineering Agent Observability Engineer Autonomy 3/10 Implements observability including structured logging, distributed tracing, metrics collection, and dashboards. SKILLS: Structured LoggingDistributed TracingMetrics CollectionDashboard DesignAlert Rules
	Performance Optimization Agent Performance Engineer Autonomy 3/10 Optimizes system performance including latency reduction, throughput improvement, and resource efficiency. SKILLS: Performance ProfilingLatency OptimizationThroughput ImprovementResource EfficiencyLoad Testing
	Platform Engineering Agent Platform Engineer Autonomy 3/10 Builds internal developer platforms, tooling, and infrastructure abstractions to accelerate engineering teams. SKILLS: Platform EngineeringDeveloper ToolsInfrastructure AbstractionSelf-ServiceDeveloper Experience
	Release Management Agent Release Manager Autonomy 3/10 Coordinates software releases including version management, changelog, feature flags, and staged rollouts. SKILLS: Release CoordinationVersion ManagementFeature FlagsStaged RolloutsChangelog
	Reliability Engineering Agent Reliability Engineer Autonomy 3/10 Ensures system reliability through SLO definition, chaos engineering, incident management, and capacity planning. SKILLS: SLO/SLA DesignChaos EngineeringIncident ManagementCapacity PlanningReliability Metrics
	Security Engineering Agent Security Engineer Autonomy 3/10 Implements security controls including authentication, encryption, vulnerability scanning, and security testing. SKILLS: Security ImplementationEncryptionVulnerability ScanningSecurity TestingAccess Control
	Systems Architecture Agent Systems Architect Autonomy 4/10 Designs distributed system architectures including service mesh, event-driven patterns, and scalability strategies. SKILLS: System ArchitectureDistributed SystemsEvent-Driven DesignScalabilityResilience Patterns
	Technical Debt Review Agent Tech Debt Analyst Autonomy 3/10 Identifies, catalogs, and prioritizes technical debt with remediation plans and business impact assessment. SKILLS: Debt IdentificationImpact AssessmentRemediation PlanningPriority ScoringCode Quality
	Testing Automation Agent Test Automation Engineer Autonomy 3/10 Designs and maintains automated testing frameworks including unit, integration, and end-to-end tests. SKILLS: Test AutomationFramework DesignE2E TestingIntegration TestingTest Coverage



Affiliate Strategy Agent Affiliate Manager Autonomy 3/10

Manages affiliate marketing programs including partner recruitment, commission structures, and performance tracking.

SKILLS:

Affiliate Programs Partner Recruitment Commission Design Performance Tracking Fraud Prevention



Channel Strategy Agent Channel Strategist Autonomy 3/10

Identifies, evaluates, and optimizes distribution channels for maximum reach and efficiency.

SKILLS:

Channel Analysis Distribution Strategy Channel Mix Performance Tracking New Channel Discovery



Conversion Optimization Agent CRO Specialist Autonomy 3/10

Optimizes conversion rates through A/B testing, UX analysis, funnel optimization, and behavioral insights.

SKILLS:

A/B Testing Funnel Optimization UX Analysis Behavioral Insights Landing Page Optimization



Demand Generation Agent Demand Gen Manager Autonomy 3/10

Creates and executes demand generation campaigns including content marketing, events, and lead nurturing.

SKILLS:

Demand Generation Content Marketing Lead Nurturing Event Strategy Pipeline Building



Email Growth Agent Email Growth Manager Autonomy 3/10

Drives growth through email marketing including list building, segmentation, automation, and deliverability optimization.

SKILLS:

Email Strategy List Building Segmentation Automation Deliverability



Funnel Analytics Agent Funnel Analyst Autonomy 3/10

Analyzes marketing and sales funnels, identifies drop-off points, and recommends optimization strategies.

SKILLS:

Funnel Analysis Drop-Off Detection Conversion Tracking Attribution Modeling Cohort Analysis



Growth Experiment Agent Growth Experimenter Autonomy 3/10

Designs and runs growth experiments with hypothesis testing, statistical rigor, and rapid iteration.

SKILLS:

Experiment Design Hypothesis Testing Statistical Analysis Rapid Iteration Learning Synthesis



Growth Strategy Agent Growth Strategist Autonomy 4/10

Develops comprehensive growth strategies across acquisition, activation, retention, revenue, and referral.

SKILLS:

Growth Strategy AARRR Framework Growth Modeling Channel Strategy Experiment Design



Influencer Strategy Agent Influencer Strategist Autonomy 3/10

Develops influencer marketing strategies including identification, outreach, campaign management, and ROI tracking.

SKILLS:

Influencer Identification Campaign Design ROI Tracking Relationship Management Content Collaboration



Lifecycle Marketing Agent Lifecycle Marketer Autonomy 3/10

Designs and optimizes customer lifecycle marketing programs from awareness through advocacy.

SKILLS:

Lifecycle Design Journey Mapping Trigger Campaigns Stage Optimization Retention Marketing



Market Penetration Agent Market Penetration Strategist Autonomy 3/10

Develops strategies to increase market share within existing markets through competitive positioning.

SKILLS:

Market Penetration Competitive Positioning Share Growth Customer Acquisition Market Defense



Organic Growth Agent Organic Growth Specialist Autonomy 3/10

Drives organic growth through content, community, word-of-mouth, and product-led strategies.

SKILLS:

Content Strategy Community Building Word-of-Mouth Product-Led Growth Organic Channels

	Paid Search Agent Paid Search Manager Autonomy 3/10 Manages paid search campaigns including keyword bidding, ad copy testing, and ROAS optimization. SKILLS: SEM Management Keyword Bidding Ad Copy Testing ROAS Optimization Budget Allocation
	Paid Social Agent Paid Social Manager Autonomy 3/10 Manages paid social advertising across platforms with audience targeting, creative testing, and budget optimization. SKILLS: Social Ads Audience Targeting Creative Testing Budget Optimization Platform Strategy
	Partnership Growth Agent Growth Partnership Manager Autonomy 3/10 Drives growth through strategic partnerships including co-marketing, integrations, and channel partnerships. SKILLS: Partnership Strategy Co-Marketing Integration Partnerships Channel Partnerships Joint Ventures
	Referral Growth Agent Referral Program Manager Autonomy 3/10 Designs and optimizes referral programs with incentive structures, tracking, and viral amplification. SKILLS: Referral Programs Incentive Design Tracking Systems Viral Amplification Program Optimization
	Retargeting Agent Retargeting Specialist Autonomy 3/10 Manages retargeting campaigns across platforms to re-engage visitors and reduce acquisition costs. SKILLS: Retargeting Campaigns Audience Segmentation Creative Rotation Frequency Capping Cross-Platform
	Revenue Expansion Agent Revenue Expansion Manager Autonomy 3/10 Drives revenue expansion through upselling, cross-selling, and pricing optimization strategies. SKILLS: Upselling Cross-Selling Pricing Optimization Expansion Revenue Account Growth
	SEO Agent SEO Specialist Autonomy 3/10 Drives organic search growth through keyword strategy, content optimization, technical SEO, and link building. SKILLS: Keyword Strategy Content SEO Technical SEO Link Building SERP Analysis
	Virality Design Agent Virality Engineer Autonomy 3/10 Designs viral mechanics including referral loops, social sharing, network effects, and viral coefficients. SKILLS: Viral Mechanics Referral Loops Network Effects Social Sharing K-Factor Optimization



Breaking Events Agent

Breaking Events Monitor

Autonomy 4/10

Monitors for breaking events that require immediate attention or action.

SKILLS:

Breaking Events

Urgency Assessment

Immediate Alerts

Impact Triage

Response Triggering



Citation and Provenance Agent

Citation Manager

Autonomy 2/10

Manages citations and provenance tracking for all search-derived information.

SKILLS:

Citation Management

Provenance Tracking

Source Attribution

Link Validation

Reference Formatting



Competitor Monitoring Agent

Competitor Monitor

Autonomy 3/10

Monitors competitor activities including product launches, pricing changes, and strategic moves.

SKILLS:

Competitor Monitoring

Product Launches

Pricing Changes

Strategic Moves

Market Positioning



Freshness Classification Agent

Freshness Classifier

Autonomy 3/10

Classifies information freshness and determines when live search is required vs cached data.

SKILLS:

Freshness Assessment

Cache vs Search

Staleness Detection

TTL Management

Priority Classification



Live Web Search Agent

Web Search Specialist

Autonomy 3/10

Executes live web searches for fresh information, current events, and real-time data.

SKILLS:

Web Search

Real-Time Data

Current Events

Source Discovery

Result Ranking



Market News Synthesis Agent

Market News Synthesizer

Autonomy 3/10

Synthesizes market news into actionable intelligence briefs for strategic decision-making.

SKILLS:

News Synthesis

Market Intelligence

Actionable Briefs

Strategic Framing

Priority Ranking



Multi-Source Verification Agent

Multi-Source Verifier

Autonomy 3/10

Verifies claims across multiple sources for cross-referencing and confidence building.

SKILLS:

Multi-Source Check

Cross-Reference

Confidence Building

Contradiction Detection

Source Triangulation



News Intelligence Agent

News Intelligence Analyst

Autonomy 3/10

Monitors and analyzes news for business-relevant events, trends, and signals.

SKILLS:

News Monitoring

Event Detection

Impact Analysis

Trend Identification

Breaking News



Product Watch Agent

Product Watch Monitor

Autonomy 3/10

Monitors product launches, updates, and market movements across competitors and industry.

SKILLS:

Product Monitoring

Launch Tracking

Feature Watch

Market Movements

Impact Analysis



Query Planning Agent

Search Query Planner

Autonomy 3/10

Plans and optimizes search queries for maximum relevance and coverage.

SKILLS:

Query Design

Search Strategy

Coverage Planning

Refinement

Multi-Source Queries



Real-Time Alerting Agent

Real-Time Alert Manager

Autonomy 3/10

Manages real-time alerts from web intelligence with severity classification and routing.

SKILLS:

Real-Time Alerts

Severity Classification

Alert Routing

Deduplication

Context Enrichment



Regulatory Monitoring Agent

Regulatory Monitor

Autonomy 3/10

Monitors regulatory changes and government actions that may impact business operations.

SKILLS:

Regulatory Monitoring

Government Actions

Policy Changes

Compliance Alerts

Impact Assessment



Risk Signal Detection Agent

Risk Signal Detector

Autonomy 3/10

Detects risk signals from web intelligence including market risks, reputation risks, and operational threats.

SKILLS:

Risk Detection

Market Signals

Reputation Monitoring

Operational Threats

Early Warning



Search Cache Agent

Search Cache Manager

Autonomy 3/10

Manages search result caching for efficiency while maintaining freshness requirements.

SKILLS:

Cache Management

TTL Strategy

Freshness Balance

Hit Rate Optimization

Storage Efficiency



Search Cost Optimization Agent

Search Cost Optimizer

Autonomy 3/10

Optimizes search costs by balancing query frequency, source selection, and caching strategies.

SKILLS:

Cost Optimization

Query Efficiency

Source Selection

Cache Strategy

Budget Management



Search Result Ranking Agent

Result Ranker

Autonomy 3/10

Ranks search results by relevance, reliability, freshness, and source quality.

SKILLS:

Result Ranking

Relevance Scoring

Reliability Assessment

Freshness Weighting

Quality Filtering



Sector Watch Agent

Sector Watch Monitor

Autonomy 3/10

Monitors sector-specific developments, regulations, and market dynamics.

SKILLS:

Sector Monitoring

Industry Developments

Regulatory Watch

Market Dynamics

Sector Reports



Sentiment Analysis Agent

Sentiment Analyst

Autonomy 3/10

Analyzes sentiment across web sources including social media, reviews, and news coverage.

SKILLS:

Sentiment Analysis

Social Listening

Review Mining

Brand Sentiment

Trend Sentiment



Source Trust Scoring Agent

Source Trust Scorer

Autonomy 3/10

Scores and maintains trust ratings for web sources based on accuracy, bias, and reliability history.

SKILLS:

Trust Scoring

Accuracy Tracking

Bias Detection

Reliability History

Source Ranking



Trend Detection Agent

Real-Time Trend Detector

Autonomy 3/10

Detects emerging trends in real-time from web data, social signals, and market movements.

SKILLS:

Trend Detection

Social Signals

Market Movements

Emerging Patterns

Momentum Analysis



Accounting Agent

Accounting Manager

Autonomy 2/10

Manages accounting operations including bookkeeping, reconciliation, and financial statement preparation.

SKILLS:

Bookkeeping

Reconciliation

Financial Statements

Tax Preparation

Audit Support



Billing Operations Agent

Billing Operations Manager

Autonomy 3/10

Manages billing operations including invoicing, payment processing, collections, and billing disputes.

SKILLS:

Invoicing

Payment Processing

Collections

Dispute Resolution

Billing Automation



Budget Planning Agent

Budget Planner

Autonomy 3/10

Plans and manages organizational budgets with allocation, tracking, and variance analysis.

SKILLS:

Budget Planning

Allocation

Variance Analysis

Cost Centers

Budget Approval



Burn Rate Monitoring Agent

Burn Rate Monitor

Autonomy 3/10

Monitors cash burn rates, tracks runway, and alerts when burn exceeds planned thresholds.

SKILLS:

Burn Rate Tracking

Runway Analysis

Threshold Alerts

Spend Velocity

Projection Updates



Capital Allocation Agent

Capital Allocator

Autonomy 2/10

Recommends capital allocation across ventures and projects based on ROI, risk, and strategic fit.

SKILLS:

Capital Allocation

ROI Analysis

Risk-Adjusted Returns

Portfolio Optimization

Rebalancing



Cash Flow Planning Agent

Cash Flow Planner

Autonomy 3/10

Plans and forecasts cash flows, manages working capital, and ensures liquidity adequacy.

SKILLS:

Cash Flow Forecasting

Working Capital

Liquidity Planning

Payment Timing

Treasury Planning



Cost Efficiency Agent

Cost Efficiency Analyst

Autonomy 3/10

Identifies and implements cost efficiency improvements across operations and technology spend.

SKILLS:

Cost Analysis

Efficiency Improvement

Spend Optimization

Process Streamlining

ROI Tracking



Expense Analysis Agent

Expense Analyst

Autonomy 3/10

Analyzes expenses, identifies cost reduction opportunities, and tracks spending against budgets.

SKILLS:

Expense Analysis

Cost Reduction

Spend Tracking

Benchmark Comparison

Optimization



Financial Modeling Agent

Financial Modeler

Autonomy 3/10

Builds financial models for ventures, products, and investments with sensitivity and scenario analysis.

SKILLS:

Financial Modeling

DCF Analysis

Sensitivity Analysis

Scenario Planning

Valuation



Financial Planning Agent

Financial Planner

Autonomy 3/10

Creates comprehensive financial plans including budgets, forecasts, and long-term financial strategy.

SKILLS:

Financial Planning

Budget Creation

Long-Term Strategy

Scenario Modeling

Financial Goals



Investment Thesis Agent

Investment Thesis Developer

Autonomy 3/10

Develops investment theses with market analysis, financial projections, and risk assessment.

SKILLS:

Investment Thesis

Market Analysis

Financial Projections

Risk Assessment

Due Diligence



Margin Analysis Agent

Margin Analyst

Autonomy 3/10

Analyzes gross and net margins across products, customers, and segments to optimize profitability.

SKILLS:

Margin Analysis

Profitability Modeling

Segment Analysis

Price Impact

Cost Allocation



Portfolio Finance Agent

Portfolio Finance Manager

Autonomy 3/10

Manages financial aspects of the venture/product portfolio including consolidated reporting and allocation.

SKILLS:

Portfolio Reporting

Consolidated Finance

Cross-Venture Analysis

Allocation Tracking

Performance



Pricing Economics Agent

Pricing Economist

Autonomy 3/10

Analyzes pricing economics including elasticity, competitive dynamics, and value capture optimization.

SKILLS:

Pricing Analysis

Elasticity Modeling

Competitive Pricing

Value Capture

Price Optimization



Revenue Forecasting Agent

Revenue Forecaster

Autonomy 3/10

Forecasts revenue using historical data, market trends, pipeline analysis, and growth assumptions.

SKILLS:

Revenue Forecasting

Trend Analysis

Pipeline Modeling

Growth Assumptions

Accuracy Tracking



Runway Analysis Agent

Runway Analyst

Autonomy 3/10

Analyzes financial runway under various scenarios and recommends actions to extend or optimize.

SKILLS:

Runway Modeling

Scenario Analysis

Extension Strategies

Optimization

Fundraising Timing



Scenario Finance Agent

Scenario Finance Analyst

Autonomy 3/10

Models financial scenarios including best/worst/likely cases for strategic decision support.

SKILLS:

Scenario Modeling

Best/Worst/Likely

Stress Testing

Decision Support

Impact Analysis



Treasury Monitoring Agent

Treasury Monitor

Autonomy 2/10

Monitors treasury operations including cash positions, investments, and banking relationships.

SKILLS:

Cash Monitoring

Investment Tracking

Bank Relations

Liquidity Analysis

Risk Management



Unit Economics Verification Agent

Unit Economics Verifier

Autonomy 3/10

Verifies unit economics assumptions and calculations across ventures and products.

SKILLS:

Unit Economics Audit

Assumption Verification

Calculation Validation

Benchmark Check

Trend Analysis



Acquisition Screening Agent

Acquisition Screener

Autonomy 2/10

Screens potential acquisition targets, conducts preliminary due diligence, and develops acquisition theses.

SKILLS:

Target Screening Due Diligence Valuation Analysis Synergy Assessment Acquisition Thesis



Business Model Architect Agent

Business Model Architect

Autonomy 3/10

Designs and stress-tests business models including revenue streams, cost structures, and value propositions.

SKILLS:

Business Modeling Revenue Design Cost Structure Value Proposition Model Stress Testing



Company Formation Design Agent

Company Formation Architect

Autonomy 2/10

Designs company structures, entity types, jurisdiction selection, and organizational architecture for new ventures.

SKILLS:

Entity Design Jurisdiction Analysis Org Architecture Legal Structure Formation Planning



Go-To-Market Blueprint Agent

GTM Strategist

Autonomy 3/10

Creates comprehensive go-to-market strategies including channel selection, launch planning, and market entry sequencing.

SKILLS:

GTM Strategy Channel Selection Launch Planning Market Entry Distribution Design



Incorporation Readiness Agent

Incorporation Advisor

Autonomy 2/10

Assesses incorporation readiness, prepares formation documents, and coordinates with legal for entity creation.

SKILLS:

Incorporation Planning Document Preparation Regulatory Compliance Legal Coordination Entity Formation



Market Gap Discovery Agent

Market Gap Analyst

Autonomy 3/10

Identifies underserved market segments, unmet customer needs, and white-space opportunities for new ventures.

SKILLS:

Gap Analysis Market Scanning Need Identification Opportunity Mapping Segment Discovery



Pivot Recommendation Agent

Pivot Advisor

Autonomy 3/10

Identifies when ventures need pivoting, evaluates pivot options, and models outcomes of strategic direction changes.

SKILLS:

Pivot Analysis Option Evaluation Direction Modeling Risk Assessment Transition Planning



Pricing Model Agent

Pricing Strategist

Autonomy 3/10

Develops pricing strategies including value-based, competitive, dynamic, and tiered pricing models.

SKILLS:

Pricing Strategy Value-Based Pricing Competitive Pricing Dynamic Pricing Price Optimization



Spinout Structuring Agent

Spinout Architect

Autonomy 2/10

Structures venture spinouts including equity splits, IP allocation, operational separation, and governance design.

SKILLS:

Spinout Design Equity Structuring IP Allocation Operational Separation Governance Design



Strategic Partnership Agent

Partnership Strategist

Autonomy 3/10

Identifies and evaluates strategic partnership opportunities, structures partnerships, and manages partner relationships.

SKILLS:

Partnership Strategy Partner Evaluation Deal Structure Relationship Management Alliance Design



Unit Economics Agent

Unit Economics Analyst

Autonomy 3/10

Analyzes unit economics including CAC, LTV, payback period, margins, and break-even projections.

SKILLS:

Unit Economics CAC/LTV Analysis Margin Analysis Break-Even Modeling Cohort Analysis



Venture Launch Agent

Venture Launch Manager

Autonomy 4/10

Coordinates venture launch activities across all teams, manages launch readiness checklists, and tracks launch metrics.

SKILLS:

Launch Coordination Readiness Assessment Checklist Management Metric Tracking Team Alignment



Venture Naming Agent

Venture Naming Specialist

Autonomy 2/10

Creates brand names for new ventures including trademark screening, domain availability, and cultural sensitivity checks.

SKILLS:

Name Generation

Trademark Screening

Domain Research

Cultural Analysis

Brand Fit Assessment



Venture Narrative Agent

Venture Storyteller

Autonomy 2/10

Crafts compelling venture narratives for investors, customers, and internal stakeholders.

SKILLS:

Narrative Craft

Investor Storytelling

Brand Story

Mission Articulation

Vision Communication



Venture Performance Review Agent

Performance Reviewer

Autonomy 3/10

Conducts periodic venture performance reviews against thesis assumptions, KPIs, and economic targets.

SKILLS:

Performance Review

KPI Tracking

Thesis Validation

Benchmark Comparison

Trend Analysis



Venture Sunset Agent

Venture Sunset Manager

Autonomy 3/10

Manages orderly wind-down of underperforming ventures including asset disposition, team redeployment, and knowledge capture.

SKILLS:

Sunset Planning

Asset Disposition

Team Redeployment

Knowledge Capture

Stakeholder Communication



Venture Thesis Agent

Venture Thesis Developer

Autonomy 3/10

Develops investment-grade venture theses with market sizing, competitive landscape, and unit economics projections.

SKILLS:

Venture Thesis

Market Sizing

Competitive Analysis

Unit Economics

Investment Framing



Venture Validation Agent

Venture Validator

Autonomy 3/10

Validates venture assumptions through customer discovery, market testing, and evidence gathering before commitment.

SKILLS:

Assumption Testing

Customer Discovery

Market Validation

Evidence Gathering

Risk Assessment



Approval Routing Agent

Approval Workflow Manager

Autonomy 2/10

Routes approval requests to appropriate authority levels, tracks approval status, enforces approval thresholds.

SKILLS:

Approval Routing

Threshold Management

Status Tracking

Authority Validation

Escalation



Audit Integrity Agent

Audit Trail Manager

Autonomy 2/10

Maintains tamper-resistant audit logs, ensures provenance tracking for all significant actions and decisions.

SKILLS:

Audit Logging

Provenance Tracking

Integrity Verification

Compliance Records

Trail Analysis



Budget Enforcement Agent

Budget Enforcement Officer

Autonomy 2/10

Enforces spending limits, tracks cost accumulation per project and agent, triggers alerts on budget overruns.

SKILLS:

Budget Tracking

Spend Limits

Cost Alerts

Overrun Prevention

Financial Controls



Capability Registry Agent

Capability Registry Manager

Autonomy 2/10

Manages the capability marketplace, registers new capabilities, prevents duplication, enables service discovery.

SKILLS:

Capability Registration

Service Discovery

Duplication Prevention

Marketplace Management

Reuse Promotion



Credential Governance Agent

Credential Manager

Autonomy 2/10

Manages credential lifecycle, key rotation, secret distribution, and time-bounded access tokens across the system.

SKILLS:

Credential Management

Key Rotation

Secret Distribution

Token Lifecycle

Access Governance



Environment Segregation Agent

Environment Controller

Autonomy 3/10

Manages sandbox, staging, and production environment separation. Enforces promotion workflows and access controls.

SKILLS:

Environment Management

Promotion Control

Access Segregation

Sandbox Enforcement

Production Gating



Execution Policy Agent

Execution Policy Enforcer

Autonomy 3/10

Enforces execution policies including approval gates, environment restrictions, and safety constraints on all actions.

SKILLS:

Policy Enforcement

Approval Gates

Safety Constraints

Environment Rules

Action Validation



Incident Coordination Agent

Incident Response Coordinator

Autonomy 3/10

Coordinates incident response across agents, manages escalation workflows, and oversees post-incident reviews.

SKILLS:

Incident Management

Escalation Coordination

Response Orchestration

Postmortem Planning

Alert Routing



Kernel Operations Agent

Kernel Operations Manager

Autonomy 3/10

Manages core kernel runtime operations including scheduling, resource allocation, and policy enforcement across the MAARS platform.

SKILLS:

Kernel Management

Runtime Monitoring

Policy Enforcement

System Health

Resource Scheduling



Memory Routing Agent

Memory Access Controller

Autonomy 3/10

Controls memory read/write access across agents, enforces retention policies, routes memory queries to appropriate layers.

SKILLS:

Memory Access Control

Retention Enforcement

Query Routing

Layer Management

Access Logging



Queue Management Agent

Queue Operations Manager

Autonomy 3/10

Manages task queues, priority ordering, dead letter handling, and queue health monitoring across all workflows.

SKILLS:

Queue Management

Priority Ordering

Dead Letter Handling

Health Monitoring

Throughput Optimization



Recovery and Rollback Agent

Recovery Operations Manager

Autonomy 3/10

Manages failure recovery workflows, coordinates rollbacks, maintains system state for recovery points.

SKILLS:

Failure Recovery

Rollback Management

State Preservation

Recovery Planning

Incident Response



Resource Allocation Agent

Resource Allocation Manager

Autonomy 3/10

Manages compute, token budgets, API quotas, and resource distribution across agents and workflows.

SKILLS:

Resource Planning

Budget Distribution

Quota Management

Capacity Monitoring

Throttle Control



Runtime Health Agent

Runtime Health Monitor

Autonomy 3/10

Continuously monitors runtime health metrics, detects anomalies, triggers circuit breakers when thresholds are exceeded.

SKILLS:

Health Monitoring

Anomaly Detection

Circuit Breakers

Performance Tracking

Alert Generation



Scheduler Oversight Agent

Scheduler Oversight Manager

Autonomy 3/10

Oversees agent scheduling, task queue prioritization, and compute budget allocation across all active workflows.

SKILLS:

Agent Scheduling

Queue Management

Priority Balancing

Compute Allocation

Deadline Tracking



State Reconciliation Agent

State Reconciliation Manager

Autonomy 3/10

Reconciles system state after failures, ensures consistency across distributed components, validates state integrity.

SKILLS:

State Reconciliation

Consistency Checks

Distributed Sync

Integrity Validation

Recovery Verification



Task Graph Integrity Agent

Task Graph Integrity Manager

Autonomy 3/10

Ensures task graph consistency, validates dependencies, detects cycles, and maintains graph versioning across all active programs.

SKILLS:

Graph Validation

Dependency Tracking

Cycle Detection

Version Control

Integrity Monitoring



Tool Access Controller Agent

Tool Permission Manager

Autonomy 3/10

Manages tool permissions per agent, enforces tool access policies, validates tool usage requests against permission matrices.

SKILLS:

Permission Management

Access Validation

Policy Enforcement

Tool Gating

Audit Logging



Brand Naming Agent

Naming Specialist

Autonomy 2/10

Creates memorable brand names with linguistic analysis, cultural screening, and trademark viability assessment.

SKILLS:

Name Creation

Linguistic Analysis

Cultural Screening

Trademark Research

Domain Check



Brand Strategy Agent

Brand Strategist

Autonomy 3/10

Develops brand strategy including positioning, architecture, voice, and long-term brand evolution plans.

SKILLS:

Brand Strategy

Brand Positioning

Brand Architecture

Voice Development

Brand Evolution



Campaign Creative Agent

Campaign Creative

Autonomy 3/10

Develops creative concepts for marketing campaigns across channels with consistent messaging and visual themes.

SKILLS:

Campaign Concepts

Multi-Channel Creative

Theme Development

Visual Systems

Message Testing



Content Design Agent

Content Designer

Autonomy 3/10

Designs content strategy and information architecture for products, websites, and digital experiences.

SKILLS:

Content Strategy

Information Architecture

UX Writing

Content Systems

Taxonomy



Copywriting Agent

Copywriter

Autonomy 3/10

Writes compelling copy for all channels including web, ads, email, social, and product interfaces.

SKILLS:

Web Copy

Ad Copy

Email Copy

Social Copy

Product Copy



Creative Direction Agent

Creative Director

Autonomy 4/10

Sets creative vision across all brand touchpoints, ensures consistency, and drives innovative creative execution.

SKILLS:

Creative Vision

Art Direction

Campaign Concepts

Brand Consistency

Innovation



Design System Agent

Design System Manager

Autonomy 3/10

Maintains and evolves the design system including components, tokens, patterns, and usage guidelines.

SKILLS:

Component Library

Design Tokens

Pattern Library

Usage Guidelines

System Governance



Graphic Design Agent

Graphic Designer

Autonomy 3/10

Creates graphic design assets including illustrations, infographics, social media graphics, and print materials.

SKILLS:

Illustration

Infographics

Social Graphics

Print Design

Visual Storytelling



Interaction Design Agent

Interaction Designer

Autonomy 3/10

Designs micro-interactions, animations, transitions, and user flow patterns for engaging experiences.

SKILLS:

Micro-Interactions

Animation Design

Flow Patterns

Prototyping

Motion Design



Landing Page Creative Agent

Landing Page Designer

Autonomy 3/10

Designs high-converting landing pages with optimized layouts, compelling copy, and clear CTAs.

SKILLS:

Landing Page Design

Conversion Optimization

CTA Design

A/B Layouts

Copy-Visual Alignment



Motion Design Agent

Motion Designer

Autonomy 3/10

Creates motion graphics, animated content, and kinetic typography for digital platforms.

SKILLS:

Motion Graphics

Animation

Kinetic Typography

Video Effects

Animated Assets



Narrative Strategy Agent

Narrative Strategist

Autonomy 3/10

Develops brand narratives, storytelling frameworks, and content strategies that build emotional connections.

SKILLS:

Brand Narrative

Storytelling

Content Strategy

Emotional Design

Story Frameworks



Presentation Design Agent

Presentation Designer

Autonomy 2/10

Designs executive presentations, pitch decks, and visual reports with professional layouts and data visualization.

SKILLS:

Presentation Design

Pitch Decks

Data Visualization

Executive Layouts

Visual Reports



UI Design Agent

UI Designer

Autonomy 3/10

Designs user interfaces with focus on usability, visual hierarchy, interaction patterns, and accessibility.

SKILLS:

UI Design

Visual Hierarchy

Interaction Patterns

Accessibility

Responsive Design



UX Research Agent

UX Researcher

Autonomy 3/10

Conducts user research including interviews, usability testing, surveys, and behavioral analysis.

SKILLS:

User Interviews

Usability Testing

Survey Design

Behavioral Analysis

Persona Creation



Video Production Agent

Video Producer

Autonomy 3/10

Produces video content including concept development, scripting, storyboarding, and post-production direction.

SKILLS:

Video Concept

Scripting

Storyboarding

Production Planning

Post-Production



Visual Identity Agent

Visual Identity Designer

Autonomy 3/10

Creates visual identity systems including logos, color palettes, typography, and brand assets.

SKILLS:

Logo Design

Color Systems

Typography

Brand Assets

Identity Guidelines



Category Creation Analysis Agent

Category Analyst

Autonomy 3/10

Identifies opportunities to create new market categories and designs category entry strategies.

SKILLS:

- Category Analysis
- Market Creation
- Positioning Strategy
- Category Design
- Entry Planning



Competitive Intelligence Agent

Competitive Analyst

Autonomy 3/10

Tracks and analyzes competitors including strategy, product moves, pricing, and market positioning.

SKILLS:

- Competitor Tracking
- Strategy Analysis
- Product Intel
- Pricing Watch
- SWOT Analysis



Customer Research Agent

Customer Researcher

Autonomy 3/10

Researches customer segments, behaviors, preferences, and needs through qualitative and quantitative methods.

SKILLS:

- Customer Segmentation
- Behavioral Analysis
- Preference Research
- Needs Assessment
- Insight Synthesis



Demand Mapping Agent

Demand Mapper

Autonomy 3/10

Maps demand across markets, segments, and geographies to identify growth opportunities.

SKILLS:

- Demand Mapping
- Geographic Analysis
- Segment Sizing
- Growth Identification
- Demand Forecasting



Emerging Signals Agent

Emerging Signals Analyst

Autonomy 3/10

Monitors weak signals of emerging opportunities, threats, and market shifts before they become obvious.

SKILLS:

- Signal Detection
- Early Warning
- Pattern Analysis
- Horizon Scanning
- Signal Validation



Industry Mapping Agent

Industry Mapper

Autonomy 3/10

Maps industry structures, value chains, key players, and ecosystem dynamics across sectors.

SKILLS:

- Industry Mapping
- Value Chain Analysis
- Player Identification
- Ecosystem Dynamics
- Structure Analysis



Macro Trend Agent

Macro Trend Analyst

Autonomy 3/10

Analyzes macroeconomic trends, geopolitical shifts, and global forces that impact business strategy.

SKILLS:

- Macro Analysis
- Geopolitical Assessment
- Economic Forecasting
- Global Trends
- Impact Modeling



Market Research Agent

Market Researcher

Autonomy 3/10

Conducts comprehensive market research including size estimation, segmentation, and trend analysis.

SKILLS:

- Market Sizing
- Segmentation
- Trend Analysis
- Consumer Research
- Industry Analysis



Opportunity Discovery Agent

Opportunity Discoverer

Autonomy 3/10

Discovers new business opportunities through market analysis, trend intersection, and need gap identification.

SKILLS:

- Opportunity Discovery
- Market Analysis
- Trend Intersection
- Need Gap
- Idea Generation



Pricing Intelligence Agent

Pricing Intelligence Analyst

Autonomy 3/10

Gathers and analyzes competitive pricing data, market rate benchmarks, and pricing strategy intelligence.

SKILLS:

- Price Monitoring
- Competitive Pricing
- Benchmark Analysis
- Strategy Intel
- Market Rates



Regulatory Research Agent

Regulatory Researcher

Autonomy 3/10

Researches regulatory landscape, upcoming changes, and compliance requirements across jurisdictions.

SKILLS:

- Regulatory Research
- Change Tracking
- Compliance Analysis
- Jurisdiction Mapping
- Impact Assessment



Research Synthesis Agent

Research Synthesizer

Autonomy 3/10

Synthesizes research from multiple sources and agents into coherent, actionable intelligence briefs.

SKILLS:

- Research Synthesis
- Multi-Source Integration
- Brief Creation
- Insight Extraction
- Knowledge Distillation



Scenario Intelligence Agent

Scenario Intelligence Analyst

Autonomy 3/10

Develops intelligence-driven scenarios for strategic planning and decision support.

SKILLS:

Scenario Development Intelligence Synthesis Strategic Scenarios Decision Support Likelihood Assessment



Signal Weakening/Strengthening Agent

Signal Analyst

Autonomy 3/10

Tracks signal strength changes over time, identifies strengthening or weakening trends.

SKILLS:

Signal Tracking Strength Analysis Trend Direction Momentum Assessment Alert Generation



Technology Scouting Agent

Technology Scout

Autonomy 3/10

Scouts emerging technologies, evaluates maturity, and assesses applicability to business objectives.

SKILLS:

Tech Scouting Maturity Assessment Applicability Analysis Innovation Tracking Adoption Readiness



Thesis Validation Agent

Thesis Validator

Autonomy 3/10

Validates strategic and investment theses against market evidence, data, and expert analysis.

SKILLS:

Thesis Validation Evidence Assessment Data Verification Expert Analysis Confidence Scoring



Trend Detection Agent

Trend Analyst

Autonomy 3/10

Detects emerging trends across industries using data analysis, signal detection, and pattern recognition.

SKILLS:

Trend Detection Signal Analysis Pattern Recognition Early Warning Trend Validation



Feature Prioritization Agent

Feature Priority Manager

Autonomy 3/10

Prioritizes features using RICE, MoSCoW, and weighted scoring frameworks based on user impact and business value.

SKILLS:

Feature Prioritization RICE Scoring Impact Analysis Value Assessment Backlog Management



Product Analytics Agent

Product Analyst

Autonomy 3/10

Analyzes product usage data, identifies patterns, generates insights, and recommends data-driven improvements.

SKILLS:

Usage Analysis Pattern Detection Insight Generation A/B Analysis Cohort Tracking



Product Architect Agent

Product Architect

Autonomy 3/10

Designs product architecture including system boundaries, integration points, and scalability patterns.

SKILLS:

Product Architecture System Design Integration Planning Scalability Patterns Technical Strategy



Product Instrumentation Agent

Instrumentation Engineer

Autonomy 3/10

Designs and implements product analytics instrumentation including event tracking, funnel definitions, and metric schemas.

SKILLS:

Event Tracking Funnel Design Metric Schema Analytics Setup Data Quality



Product Iteration Agent

Iteration Manager

Autonomy 3/10

Manages rapid product iteration cycles including sprint planning, velocity tracking, and continuous improvement.

SKILLS:

Sprint Planning Velocity Tracking Iteration Management Continuous Improvement Retrospectives



Product Launch Readiness Agent

Launch Readiness Manager

Autonomy 3/10

Assesses launch readiness across engineering, marketing, support, and operations with go/no-go checklists.

SKILLS:

Launch Readiness Go/No-Go Assessment Cross-Team Coordination Risk Evaluation Checklist Management



Product Packaging Agent

Product Packaging Specialist

Autonomy 3/10

Designs product packaging including tier structure, feature bundling, and pricing presentation.

SKILLS:

Tier Design Feature Bundling Pricing Presentation Package Optimization SKU Management



Product Portfolio Agent

Product Portfolio Manager

Autonomy 3/10

Manages the product portfolio, balances investment across products, and identifies gaps or redundancies.

SKILLS:

Portfolio Management Investment Balancing Gap Analysis Product Rationalization Lifecycle Planning



Product Positioning Agent

Product Positioning Specialist

Autonomy 3/10

Defines product positioning, messaging frameworks, and competitive differentiation for target segments.

SKILLS:

Product Positioning Messaging Framework Competitive Differentiation Segment Targeting Value Communication



Product Quality Review Agent

Quality Reviewer

Autonomy 3/10

Reviews product quality across dimensions including UX, performance, reliability, and user satisfaction.

SKILLS:

Quality Review UX Assessment Performance Review Reliability Analysis Satisfaction Tracking



Product Requirements Agent

Requirements Engineer

Autonomy 3/10

Elicits, documents, and manages product requirements including functional specs, acceptance criteria, and edge cases.

SKILLS:

Requirements Engineering Spec Writing Acceptance Criteria Edge Case Analysis Stakeholder Alignment



Product Retention Agent

Retention Specialist

Autonomy 3/10

Analyzes and optimizes user retention through engagement loops, re-activation campaigns, and churn prediction.

SKILLS:

Retention Analysis Churn Prediction Engagement Loops Re-activation Cohort Analysis



Product Strategy Agent

Product Strategist

Autonomy 4/10

Defines product strategy aligned with business objectives, market positioning, and competitive differentiation.

SKILLS:

Product Strategy

Market Positioning

Competitive Differentiation

Strategic Roadmapping

Portfolio Strategy



Product Sunset Agent

Product Sunset Manager

Autonomy 3/10

Plans and executes product deprecation including migration paths, customer communication, and data handling.

SKILLS:

Deprecation Planning

Migration Design

Customer Communication

Data Handling

Timeline Management



Product Testing Agent

Product Test Manager

Autonomy 3/10

Plans and coordinates product testing including QA strategy, test plans, UAT, and regression testing.

SKILLS:

Test Strategy

QA Planning

UAT Coordination

Regression Testing

Bug Triage



Roadmap Planning Agent

Roadmap Planner

Autonomy 3/10

Creates and maintains product roadmaps with themes, milestones, dependencies, and capacity constraints.

SKILLS:

Roadmap Planning

Theme Organization

Milestone Tracking

Dependency Management

Capacity Planning



User Needs Analysis Agent

User Needs Analyst

Autonomy 3/10

Analyzes user needs through research, behavioral data, and feedback to inform product decisions.

SKILLS:

User Research

Needs Analysis

Behavioral Data

Feedback Synthesis

Persona Development



Advertising Compliance Agent

Ad Compliance Officer

Autonomy 3/10

Ensures advertising and marketing materials comply with regulations including FTC, ASA, and industry rules.

SKILLS:

Ad Compliance FTC Rules Claims Verification Disclosure Requirements Industry Standards



Approval Policy Agent

Approval Policy Manager

Autonomy 3/10

Manages approval policies including thresholds, escalation paths, and delegation of authority.

SKILLS:

Approval Policies Threshold Management Escalation Design Authority Delegation Policy Compliance



Audit Readiness Agent

Audit Readiness Manager

Autonomy 3/10

Prepares organization for audits including documentation, evidence gathering, and compliance verification.

SKILLS:

Audit Preparation Documentation Review Evidence Gathering Compliance Verification Gap Remediation



Compliance Monitoring Agent

Compliance Monitor

Autonomy 3/10

Monitors ongoing compliance across operations, detects violations, and tracks remediation.

SKILLS:

Compliance Monitoring Violation Detection Remediation Tracking Audit Preparation Gap Analysis



Contract Review Agent

Contract Reviewer

Autonomy 2/10

Reviews contracts for risk, compliance, and favorable terms. Identifies problematic clauses.

SKILLS:

Contract Review Risk Identification Clause Analysis Term Negotiation Template Management



Data Privacy Agent

Data Privacy Officer

Autonomy 3/10

Manages data privacy compliance including GDPR, CCPA, data mapping, and privacy impact assessments.

SKILLS:

Privacy Compliance GDPR/CCPA Data Mapping PIA/DPIA Consent Management



Entity Structure Review Agent

Entity Structure Reviewer

Autonomy 2/10

Reviews and recommends corporate entity structures for tax efficiency, liability protection, and operational needs.

SKILLS:

Entity Review Structure Optimization Tax Efficiency Liability Analysis Formation Advice



Ethics Review Agent

Ethics Reviewer

Autonomy 3/10

Reviews decisions and actions for ethical implications including bias, fairness, and social responsibility.

SKILLS:

Ethics Review Bias Assessment Fairness Analysis Social Responsibility Ethical Guidelines



Governance Policy Agent

Governance Policy Manager

Autonomy 3/10

Develops and maintains governance policies including board governance, decision authority, and oversight frameworks.

SKILLS:

Governance Policy Board Governance Decision Authority Oversight Frameworks Policy Updates



IP Risk Agent

IP Risk Analyst

Autonomy 2/10

Assesses intellectual property risks including patent, trademark, copyright, and trade secret issues.

SKILLS:

IP Risk Assessment Patent Analysis Trademark Review Copyright Compliance Trade Secrets



Jurisdiction Analysis Agent

Jurisdiction Analyst

Autonomy 2/10

Analyzes jurisdictional requirements for operations, products, and services across regions.

SKILLS:

Jurisdiction Analysis Cross-Border Rules Regional Requirements Compliance Mapping Risk Assessment



Legal Research Agent

Legal Researcher

Autonomy 2/10

Conducts legal research across jurisdictions, case law, regulations, and legal precedents.

SKILLS:

Legal Research Case Law Regulatory Analysis Precedent Search Jurisdiction Review



Policy Interpretation Agent

Policy Interpreter

Autonomy 2/10

Interprets internal and external policies, provides guidance on compliance, and resolves policy ambiguities.

SKILLS:

Policy Interpretation

Compliance Guidance

Ambiguity Resolution

Policy Application

Advisory



Regulatory Change Detection Agent

Regulatory Change Monitor

Autonomy 3/10

Monitors regulatory changes and upcoming legislation that may impact business operations.

SKILLS:

Regulatory Monitoring

Change Detection

Impact Assessment

Legislation Tracking

Preparation Planning



Sector-Specific Compliance Agent

Sector Compliance Specialist

Autonomy 3/10

Manages sector-specific compliance requirements including healthcare, finance, and regulated industries.

SKILLS:

Sector Compliance

HIPAA/SOX

Industry Regulations

Specialized Requirements

Compliance Programs



Vendor Compliance Agent

Vendor Compliance Manager

Autonomy 3/10

Ensures vendor compliance with contractual obligations, security requirements, and regulatory standards.

SKILLS:

Vendor Compliance

Contract Enforcement

Security Assessment

Regulatory Adherence

Audit Coordination

 **Account Expansion Agent** Account Expansion Manager Autonomy 3/10

Identifies and executes expansion opportunities within existing accounts through upselling and cross-selling.

SKILLS:

Account ExpansionUpsellingCross-SellingOpportunity IdentificationRelationship Building

 **CRM Strategy Agent** CRM Strategist Autonomy 3/10

Develops CRM strategy including data architecture, automation rules, and customer journey tracking.

SKILLS:

CRM StrategyData ArchitectureAutomation RulesJourney TrackingIntegration Planning

 **Churn Prevention Agent** Churn Prevention Specialist Autonomy 3/10

Predicts and prevents customer churn through early warning signals, intervention strategies, and win-back campaigns.

SKILLS:

Churn PredictionEarly WarningIntervention DesignWin-Back CampaignsRisk Scoring

 **Deal Desk Agent** Deal Desk Manager Autonomy 2/10

Manages deal structuring, pricing approvals, and non-standard deal terms with proper governance.

SKILLS:

Deal StructuringPricing ApprovalsNon-Standard TermsRisk AssessmentGovernance

 **Lead Qualification Agent** Lead Qualifier Autonomy 3/10

Qualifies leads using scoring models, intent signals, and ICP fit assessment.

SKILLS:

Lead ScoringICP AssessmentIntent AnalysisQualification CriteriaPipeline Quality

 **Offer Design Agent** Offer Designer Autonomy 3/10

Designs compelling offers including pricing, packaging, incentives, and value propositions for deals.

SKILLS:

Offer DesignPricing StrategyIncentive StructureValue PropositionDeal Packaging

 **Pipeline Management Agent** Pipeline Manager Autonomy 3/10

Manages sales pipeline including stage progression, forecast accuracy, and deal health monitoring.

SKILLS:

Pipeline ManagementStage TrackingForecast AccuracyDeal HealthRisk Identification

 **Proposal Generation Agent** Proposal Writer Autonomy 3/10

Generates professional proposals including scope, pricing, timelines, and terms tailored to prospect needs.

SKILLS:

Proposal WritingScope DefinitionPricing TablesTimeline PlanningCustom Tailoring

 **Prospect Research Agent** Prospect Researcher Autonomy 3/10

Researches prospects including company intelligence, decision-maker mapping, and engagement history.

SKILLS:

Company ResearchDecision-Maker MappingIntelligence GatheringEngagement HistoryOpportunity Sizing

 **Retention Revenue Agent** Retention Revenue Manager Autonomy 3/10

Protects and grows revenue from existing customers through retention strategies and renewal optimization.

SKILLS:

Retention StrategyRenewal OptimizationChurn PreventionCustomer Value GrowthRisk Mitigation

 **Revenue Forecast Coordination Agent** Revenue Forecaster Autonomy 3/10

Coordinates revenue forecasting across teams, models pipeline conversion, and tracks forecast accuracy.

SKILLS:

Revenue ForecastingPipeline ModelingForecast AccuracyVariance AnalysisPrediction Models

 **Revenue Operations Agent** RevOps Manager Autonomy 3/10

Aligns sales, marketing, and customer success operations to maximize revenue efficiency.

SKILLS:

RevOps StrategyProcess AlignmentRevenue EfficiencySystem IntegrationReporting



Sales Enablement Agent

Sales Enablement Manager

Autonomy 3/10

Creates sales enablement content, training materials, battle cards, and competitive intelligence for sales teams.

SKILLS:

- Sales Content
- Training Materials
- Battle Cards
- Competitive Intel
- Playbook Design



Sales Strategy Agent

Sales Strategist

Autonomy 3/10

Develops sales strategies including territory planning, quota setting, and sales process optimization.

SKILLS:

- Sales Strategy
- Territory Planning
- Quota Setting
- Process Optimization
- Win Rate Analysis



Acquisition Evaluation Agent

Acquisition Evaluator

Autonomy 2/10

Evaluates acquisition targets including strategic fit, valuation, integration complexity, and synergy potential.

SKILLS:

Target Evaluation Valuation Analysis Integration Assessment Synergy Potential Risk Factors



Capital Reallocation Agent

Capital Reallocator

Autonomy 2/10

Recommends capital reallocation between ventures based on performance, opportunity cost, and strategic fit.

SKILLS:

Reallocation Analysis Performance Comparison Opportunity Cost Strategic Fit Transition Planning



Concentration Risk Agent

Concentration Risk Analyst

Autonomy 3/10

Monitors and manages concentration risk across the portfolio including sector, stage, and geography.

SKILLS:

Concentration Analysis Risk Monitoring Diversification Strategy Geographic Risk Sector Risk



Decline Review Agent

Decline Reviewer

Autonomy 3/10

Reviews declining ventures for turnaround potential, sunset timing, and asset recovery.

SKILLS:

Decline Analysis Turnaround Assessment Sunset Planning Asset Recovery Wind-Down



Exit Strategy Agent

Exit Strategist

Autonomy 3/10

Develops exit strategies including IPO readiness, M&A positioning, and secondary sale preparation.

SKILLS:

Exit Strategy IPO Readiness M&A Positioning Secondary Sales Timing Optimization



Incubation Portfolio Agent

Incubation Manager

Autonomy 3/10

Manages early-stage incubation portfolio with milestone tracking, mentor matching, and graduation criteria.

SKILLS:

Incubation Management Milestone Tracking Mentor Matching Graduation Criteria Resource Allocation



Kill/Scale Decision Agent

Kill/Scale Analyst

Autonomy 3/10

Evaluates ventures for kill/scale decisions using performance data, market signals, and economic analysis.

SKILLS:

Kill/Scale Analysis Performance Evaluation Market Signals Economic Assessment Decision Framework



Mature Portfolio Agent

Mature Portfolio Manager

Autonomy 3/10

Manages mature portfolio companies focused on optimization, dividend policy, and long-term value creation.

SKILLS:

Portfolio Optimization Dividend Policy Value Creation Growth Sustaining Efficiency



Opportunity Screening Agent

Opportunity Screener

Autonomy 3/10

Screens investment opportunities using defined criteria, market analysis, and risk filters.

SKILLS:

Opportunity Screening Criteria Evaluation Market Filtering Risk Assessment Pipeline Ranking



Portfolio Balancing Agent

Portfolio Balancer

Autonomy 3/10

Balances portfolio allocation across stages, sectors, and risk profiles for optimal returns.

SKILLS:

Portfolio Balancing Allocation Optimization Risk Diversification Sector Balancing Stage Mix



Portfolio Reporting Agent

Portfolio Reporter

Autonomy 2/10

Creates comprehensive portfolio reports including performance metrics, valuations, and investor updates.

SKILLS:

Portfolio Reporting Performance Metrics Valuation Updates Investor Communication Benchmarking



Reinvestment Strategy Agent

Reinvestment Strategist

Autonomy 3/10

Plans reinvestment of returns including follow-on investments, new opportunities, and reserve management.

SKILLS:

Reinvestment Planning Follow-On Strategy Reserve Management Return Recycling Allocation



Synergy Mapping Agent

Synergy Analyst

Autonomy 3/10

Maps synergies between portfolio companies including shared resources, cross-selling, and technology reuse.

SKILLS:

- Synergy Mapping
- Cross-Selling
- Resource Sharing
- Technology Reuse
- Value Creation



Venture Scoring Agent

Venture Scorer

Autonomy 3/10

Scores ventures using multi-factor models including market, team, technology, and financial metrics.

SKILLS:

- Venture Scoring
- Multi-Factor Analysis
- Comparative Ranking
- Score Calibration
- Benchmark



Capacity Planning Agent

Capacity Planner

Autonomy 3/10

Plans and manages capacity across infrastructure, team, and systems to meet demand.

SKILLS:

- Capacity Planning
- Demand Modeling
- Resource Forecasting
- Scaling Strategy
- Constraint Analysis



Documentation Operations Agent

Documentation Ops Manager

Autonomy 2/10

Manages operational documentation including procedures, runbooks, and knowledge base maintenance.

SKILLS:

- Documentation Management
- Runbooks
- Procedure Maintenance
- Version Control
- Access Management



Internal Service Agent

Internal Service Manager

Autonomy 3/10

Manages internal service delivery including IT support, facilities, and shared services.

SKILLS:

- Service Delivery
- IT Support
- Facilities Management
- Shared Services
- Service Catalog



Inventory Coordination Agent

Inventory Coordinator

Autonomy 3/10

Manages inventory levels, demand forecasting, reorder points, and stock optimization.

SKILLS:

- Inventory Management
- Demand Forecasting
- Reorder Optimization
- Stock Control
- Waste Reduction



Logistics Coordination Agent

Logistics Coordinator

Autonomy 3/10

Coordinates logistics operations including shipping, warehousing, distribution, and delivery optimization.

SKILLS:

- Logistics Planning
- Shipping Optimization
- Distribution Management
- Delivery Tracking
- Cost Control



Operational Risk Agent

Operational Risk Analyst

Autonomy 3/10

Identifies and assesses operational risks, develops mitigation strategies, and monitors risk indicators.

SKILLS:

- Risk Identification
- Risk Assessment
- Mitigation Strategy
- Risk Monitoring
- Incident Analysis



Operations Strategy Agent

Operations Strategist

Autonomy 4/10

Develops operational strategy aligned with business objectives, identifies efficiency opportunities.

SKILLS:

- Operations Strategy
- Efficiency Analysis
- Process Reengineering
- Operational Excellence
- KPI Design



Process Design Agent

Process Designer

Autonomy 3/10

Designs new business processes with documentation, SOP creation, and change management support.

SKILLS:

- Process Design
- SOP Creation
- Documentation
- Change Management
- Training Design



Procurement Analysis Agent

Procurement Analyst

Autonomy 3/10

Analyzes procurement needs, evaluates suppliers, negotiates terms, and optimizes purchasing decisions.

SKILLS:

- Procurement Analysis
- Supplier Evaluation
- Negotiation Support
- Cost Optimization
- Spend Analysis



Program Operations Agent

Program Operations Manager

Autonomy 3/10

Manages operational aspects of programs including resource coordination, milestone tracking, and risk management.

SKILLS:

- Program Management
- Resource Coordination
- Milestone Tracking
- Risk Management
- Status Reporting



SLA Monitoring Agent

SLA Monitor

Autonomy 3/10

Monitors SLA compliance across all services, tracks performance metrics, and alerts on potential breaches.

SKILLS:

- SLA Monitoring
- Compliance Tracking
- Performance Metrics
- Breach Alerts
- Remediation



Vendor Management Agent

Vendor Manager

Autonomy 3/10

Manages vendor relationships including selection, performance monitoring, contract management, and risk assessment.

SKILLS:

- Vendor Selection
- Performance Monitoring
- Contract Management
- Risk Assessment
- Relationship Management



Workflow Optimization Agent

Workflow Optimizer

Autonomy 3/10

Analyzes and optimizes business workflows using lean principles, automation, and process mining.

SKILLS:

Workflow Analysis

Lean Optimization

Process Mining

Automation Design

Bottleneck Removal



Workforce Planning Agent

Workforce Planner

Autonomy 3/10

Plans workforce needs including capacity modeling, skill gap analysis, and resource allocation optimization.

SKILLS:

Capacity Modeling

Skill Gap Analysis

Resource Allocation

Hiring Planning

Team Structure



Community Management Agent

Community Manager

Autonomy 3/10

Builds and manages customer communities including forums, events, and user groups.

SKILLS:

Community Building

Forum Management

Event Planning

User Groups

Engagement Programs



Complaint Analysis Agent

Complaint Analyst

Autonomy 3/10

Analyzes customer complaints to identify systemic issues, patterns, and root causes for improvement.

SKILLS:

Complaint Analysis

Pattern Detection

Root Cause Analysis

Systemic Issues

Improvement Recommendations



Customer Health Agent

Customer Health Monitor

Autonomy 3/10

Monitors customer health scores, identifies at-risk accounts, and triggers proactive interventions.

SKILLS:

Health Scoring

Risk Detection

Proactive Outreach

Usage Monitoring

Engagement Tracking



Customer Support Agent

Customer Support Specialist

Autonomy 3/10

Provides frontline customer support with issue resolution, product guidance, and escalation when needed.

SKILLS:

Issue Resolution

Product Guidance

Ticket Management

Escalation

Customer Communication



Education and Enablement Agent

Education Manager

Autonomy 3/10

Creates educational content, training programs, and enablement resources for customer success.

SKILLS:

Training Programs

Educational Content

Webinar Design

Certification

Learning Paths



Escalation Handling Agent

Escalation Handler

Autonomy 3/10

Manages customer escalations with priority routing, resolution tracking, and executive escalation procedures.

SKILLS:

Escalation Management

Priority Routing

Resolution Tracking

Executive Escalation

Root Cause



Knowledge Base Agent

Knowledge Base Manager

Autonomy 3/10

Maintains customer-facing knowledge base with articles, FAQs, troubleshooting guides, and search optimization.

SKILLS:

KB Management

Article Writing

FAQ Maintenance

Search Optimization

Content Gaps



Onboarding Agent

Onboarding Specialist

Autonomy 3/10

Designs and delivers customer onboarding experiences that drive activation and time-to-value.

SKILLS:

Onboarding Design

Activation Tracking

Time-to-Value

Success Milestones

Guided Tours



Renewal Readiness Agent

Renewal Manager

Autonomy 3/10

Prepares for and manages subscription renewals including health assessment, value documentation, and negotiation.

SKILLS:

Renewal Preparation

Health Assessment

Value Documentation

Negotiation Support

Risk Mitigation



Satisfaction Measurement Agent

Satisfaction Analyst

Autonomy 3/10

Measures and tracks customer satisfaction through NPS, CSAT, CES surveys and analysis.

SKILLS:

NPS Tracking

CSAT Analysis

CES Measurement

Survey Design

Trend Analysis



Success Planning Agent

Success Planner

Autonomy 3/10

Creates customer success plans with goals, milestones, and action items aligned to customer outcomes.

SKILLS:

Success Planning

Goal Setting

Milestone Tracking

Outcome Alignment

QBR Preparation



Support QA Agent

Support Quality Analyst

Autonomy 3/10

Monitors and improves support quality through ticket audits, agent coaching, and process improvement.

SKILLS:

Quality Audits

Agent Coaching

Process Improvement

Metric Tracking

Best Practices



Support Workflow Agent

Support Workflow Manager

Autonomy 3/10

Designs and optimizes support workflows including routing, SLA management, and automation.

SKILLS:

Workflow Design Routing Rules SLA Management Automation Efficiency Metrics



Voice of Customer Agent

VoC Analyst

Autonomy 3/10

Collects and analyzes customer feedback, identifies themes, and translates insights into product improvements.

SKILLS:

Feedback Collection Theme Analysis Insight Translation Sentiment Tracking NPS Analysis



Access Control Agent

Access Control Manager

Autonomy 3/10

Manages role-based and attribute-based access control policies across all systems.

SKILLS:

RBAC/ABAC

Policy Management

Access Reviews

Least Privilege

Audit



Anomaly Detection Agent

Anomaly Detector

Autonomy 3/10

Detects anomalies in system behavior, user activity, and data patterns that may indicate security issues.

SKILLS:

Anomaly Detection

Behavioral Analysis

Pattern Recognition

Alert Generation

False Positive Management



Authentication Integrity Agent

Authentication Integrity Monitor

Autonomy 3/10

Monitors authentication system integrity including credential compromise detection and session management.

SKILLS:

Auth Monitoring

Credential Compromise

Session Management

Token Integrity

Breach Detection



Data Security Agent

Data Security Officer

Autonomy 3/10

Protects data at rest and in transit with encryption, access controls, and data loss prevention.

SKILLS:

Data Encryption

DLP

Access Controls

Data Classification

Breach Prevention



Deployment Security Agent

Deployment Security Manager

Autonomy 3/10

Ensures security of deployment processes including signing, verification, and rollback capabilities.

SKILLS:

Deploy Security

Code Signing

Verification

Rollback

Pipeline Security



Identity Security Agent

Identity Security Manager

Autonomy 3/10

Manages identity security including authentication, authorization, and identity lifecycle management.

SKILLS:

Identity Management

Authentication

Authorization

Lifecycle Management

SSO/MFA



Incident Response Agent

Security Incident Responder

Autonomy 4/10

Responds to security incidents with containment, eradication, recovery, and post-incident analysis.

SKILLS:

Incident Response

Containment

Eradication

Recovery

Post-Incident Analysis



Policy Hardening Agent

Security Policy Hardener

Autonomy 3/10

Continuously hardens security policies based on threat landscape changes and incident learnings.

SKILLS:

Policy Hardening

Threat Adaptation

Baseline Improvement

Control Strengthening

Best Practices



Recovery Security Agent

Security Recovery Manager

Autonomy 3/10

Manages security aspects of system recovery including integrity verification and secure restoration.

SKILLS:

Secure Recovery

Integrity Verification

Restoration

Evidence Preservation

Clean State



Secrets Management Agent

Secrets Manager

Autonomy 3/10

Manages secrets lifecycle including key generation, rotation, distribution, and revocation.

SKILLS:

Secrets Management

Key Rotation

Distribution

Revocation

Vault Management



Security Posture Agent

Security Posture Analyst

Autonomy 3/10

Assesses and reports on overall security posture with risk scoring and improvement recommendations.

SKILLS:

Posture Assessment

Risk Scoring

Improvement Plans

Benchmark Comparison

Compliance Status



Threat Monitoring Agent

Threat Monitor

Autonomy 3/10

Monitors for security threats including intrusion attempts, malware, and suspicious activity.

SKILLS:

Threat Monitoring

Intrusion Detection

Malware Detection

Suspicious Activity

Alert Triage



Tool Abuse Detection Agent

Tool Abuse Monitor

Autonomy 3/10

Detects misuse or abuse of system tools and capabilities by agents or external actors.

SKILLS:

- Abuse Detection
- Usage Monitoring
- Policy Violation
- Rate Limiting
- Quarantine



API Reliability Agent API Reliability Manager Autonomy 3/10

Monitors and ensures API reliability including uptime, latency SLAs, and error budgets.

SKILLS:

- API Monitoring
- SLA Tracking
- Error Budgets
- Latency Management
- Reliability Reports



Capability Discovery Agent Capability Discoverer Autonomy 3/10

Discovers available capabilities across the system to prevent duplication and enable reuse.

SKILLS:

- Capability Discovery
- Reuse Assessment
- Duplication Prevention
- Service Mapping
- Gap Analysis



Capability Marketplace Agent Marketplace Manager Autonomy 3/10

Manages the internal capability marketplace for tool and service discovery.

SKILLS:

- Marketplace Management
- Service Catalog
- Discovery UX
- Rating System
- Usage Analytics



Capability Reuse Agent Reuse Promoter Autonomy 3/10

Promotes capability reuse by matching needs to existing tools and services.

SKILLS:

- Reuse Promotion
- Need Matching
- Tool Recommendation
- Integration Guidance
- Best Practices



Connector Qualification Agent Connector Qualifier Autonomy 3/10

Qualifies new connectors and integrations for security, reliability, and performance standards.

SKILLS:

- Connector Qualification
- Security Review
- Reliability Testing
- Performance Benchmark
- Approval



Cost-Aware Tool Selection Agent Cost-Aware Selector Autonomy 3/10

Selects tools based on cost-effectiveness considering quality, speed, and reliability tradeoffs.

SKILLS:

- Cost Analysis
- Tool Comparison
- Tradeoff Analysis
- Budget Optimization
- ROI Assessment



Fallback Tool Routing Agent Fallback Router Autonomy 3/10

Routes to fallback tools when primary tools fail, managing graceful degradation.

SKILLS:

- Fallback Routing
- Graceful Degradation
- Alternative Selection
- Recovery Paths
- Priority Fallback



Sandbox Tool Agent Sandbox Manager Autonomy 3/10

Manages sandbox environments for tool testing and new capability evaluation.

SKILLS:

- Sandbox Management
- Tool Testing
- Isolation
- Evaluation
- Promotion Readiness



Tool Audit Agent Tool Auditor Autonomy 3/10

Audits tool usage including compliance, cost, performance, and access pattern analysis.

SKILLS:

- Usage Auditing
- Compliance Check
- Cost Analysis
- Performance Review
- Access Patterns



Tool Health Agent Tool Health Monitor Autonomy 3/10

Monitors tool health including availability, latency, error rates, and reliability scoring.

SKILLS:

- Health Monitoring
- Availability Tracking
- Latency Analysis
- Error Rate
- Reliability Scoring



Tool Registry Agent Tool Registry Manager Autonomy 3/10

Manages the central tool registry including registration, metadata, and version control.

SKILLS:

- Tool Registration
- Metadata Management
- Version Control
- Discovery
- Deprecation



Tool Schema Validation Agent

Schema Validator

Autonomy 3/10

Validates tool input/output schemas, detects schema drift, and ensures compatibility.

SKILLS:

Schema Validation

Drift Detection

Compatibility Check

Contract Testing

Version Alignment



Alert Routing Agent

Alert Router

Autonomy 3/10

Routes alerts to appropriate teams and agents based on severity, type, and ownership.

SKILLS:

Alert Routing

Severity Classification

Team Assignment

Deduplication

Noise Reduction



Cost Monitoring Agent

Cost Monitor

Autonomy 3/10

Monitors system costs in real-time including API spend, compute, and operational costs.

SKILLS:

Cost Tracking

Spend Alerts

Budget Monitoring

Provider Costs

Optimization Recommendations



Dashboard Generation Agent

Dashboard Generator

Autonomy 3/10

Generates and maintains operational dashboards for system visibility and decision support.

SKILLS:

Dashboard Creation

Widget Design

Data Binding

Real-Time Updates

Custom Views



Escalation Trend Agent

Escalation Trend Analyst

Autonomy 3/10

Tracks escalation trends to identify systemic problems and improve resolution processes.

SKILLS:

Escalation Tracking

Trend Analysis

Process Improvement

Root Cause Patterns

Resolution Metrics



Failure Detection Agent

Failure Detector

Autonomy 3/10

Detects system failures including service outages, error spikes, and degraded performance.

SKILLS:

Failure Detection

Error Spikes

Service Outages

Performance Degradation

Alert Routing



Incident Review Agent

Incident Reviewer

Autonomy 3/10

Reviews incidents for root cause, impact assessment, and improvement recommendations.

SKILLS:

Incident Review

Root Cause Analysis

Impact Assessment

Improvement Recommendations

Lessons Learned



Metrics Consolidation Agent

Metrics Consolidator

Autonomy 3/10

Consolidates metrics from across the system into unified dashboards and reports.

SKILLS:

Metric Collection

Consolidation

Dashboard Updates

Anomaly Detection

Trend Reporting



Postmortem Agent

Postmortem Manager

Autonomy 3/10

Facilitates blameless postmortems and ensures action items are tracked and completed.

SKILLS:

Postmortem Facilitation

Action Items

Blameless Culture

Documentation

Follow-Up Tracking



Retry Pattern Analysis Agent

Retry Analyst

Autonomy 3/10

Analyzes retry patterns to identify systemic issues and optimize retry strategies.

SKILLS:

Retry Analysis

Pattern Detection

Strategy Optimization

Root Cause Linking

Circuit Breaking



Root Cause Analysis Agent

RCA Specialist

Autonomy 3/10

Conducts deep root cause analysis using systematic investigation methods.

SKILLS:

Root Cause Analysis

5 Whys

Fishbone Diagrams

Fault Trees

Contributing Factors



SLO/SLA Agent

SLO/SLA Manager

Autonomy 3/10

Defines and monitors SLOs/SLAs with error budget tracking and compliance reporting.

SKILLS:

SLO Definition

SLA Monitoring

Error Budget

Compliance Reporting

Threshold Management



Uptime Analysis Agent

Uptime Analyst

Autonomy 3/10

Analyzes system uptime, availability patterns, and reliability trends.

SKILLS:

Uptime Analysis

Availability Tracking

SLA Reporting

Trend Analysis

Improvement Planning



Archive Curation Agent

Archive Curator

Autonomy 2/10

Curates archived knowledge with retention policies, access controls, and retrieval optimization.

SKILLS:

Archive Management

Retention Policies

Access Controls

Retrieval Design

Storage Optimization



Context Packing Agent

Context Packer

Autonomy 3/10

Optimizes context assembly for agent tasks by selecting and packing the most relevant knowledge.

SKILLS:

Context Assembly

Relevance Selection

Token Optimization

Priority Packing

Context Windows



Entity Resolution Agent

Entity Resolver

Autonomy 3/10

Resolves entity references across knowledge sources, deduplicates, and maintains canonical entities.

SKILLS:

Entity Resolution

Deduplication

Canonical Entities

Cross-Reference

Merge Operations



Episodic Memory Agent

Episodic Memory Manager

Autonomy 3/10

Manages episodic memories including task histories, decision records, and event sequences.

SKILLS:

Episode Recording

History Management

Decision Records

Event Sequencing

Recall Optimization



Freshness Reconciliation Agent

Freshness Manager

Autonomy 3/10

Monitors knowledge freshness, flags stale information, and triggers refresh workflows.

SKILLS:

Freshness Monitoring

Stale Detection

Refresh Triggers

Age Tracking

Update Prioritization



Knowledge Graph Maintenance Agent

Knowledge Graph Manager

Autonomy 3/10

Maintains the knowledge graph including entity creation, relationship updates, and consistency checks.

SKILLS:

Graph Maintenance

Entity Management

Relationship Updates

Consistency Checks

Query Optimization



Knowledge Quality Agent

Knowledge Quality Manager

Autonomy 3/10

Ensures knowledge quality through validation, accuracy checks, and freshness monitoring.

SKILLS:

Quality Validation

Accuracy Checks

Freshness Monitoring

Contradiction Detection

Quality Scoring



Memory Compaction Agent

Memory Compactor

Autonomy 3/10

Compacts memory by merging related entries, removing redundancy, and optimizing storage.

SKILLS:

Memory Compaction

Redundancy Removal

Entry Merging

Storage Optimization

Quality Preservation



Memory Retrieval Agent

Memory Retrieval Specialist

Autonomy 3/10

Optimizes memory retrieval with relevance ranking, context-aware search, and multi-layer queries.

SKILLS:

Relevance Ranking

Context Search

Multi-Layer Query

Recall Optimization

Result Ranking



Memory Summarization Agent

Memory Summarizer

Autonomy 3/10

Summarizes and compresses memory entries while preserving key information and context.

SKILLS:

Summarization

Compression

Key Extraction

Context Preservation

Deduplication



Provenance Tracking Agent

Provenance Tracker

Autonomy 2/10

Tracks provenance of knowledge including sources, creation context, and modification history.

SKILLS:

Source Tracking

Creation Context

Modification History

Trust Assessment

Lineage



Semantic Memory Agent

Semantic Memory Manager

Autonomy 3/10

Manages semantic knowledge including concepts, relationships, and factual knowledge.

SKILLS:

Concept Management

Relationship Mapping

Fact Storage

Knowledge Organization

Semantic Search



Citation Integrity Agent

Citation Verifier

Autonomy 3/10

Verifies citation accuracy, source existence, and claim-to-source alignment.

SKILLS:

Citation Verification

Source Existence

Claim Alignment

URL Validation

Reference Accuracy



Code Verification Agent

Code Verifier

Autonomy 3/10

Reviews and verifies code quality, security, and correctness before deployment.

SKILLS:

Code Review

Security Scan

Correctness Check

Test Coverage

Best Practices



Compliance Verification Agent

Compliance Verifier

Autonomy 3/10

Verifies compliance of outputs and actions against regulatory and policy requirements.

SKILLS:

Compliance Check

Regulatory Verification

Policy Adherence

Audit Trail

Gap Detection



Experiment Verification Agent

Experiment Verifier

Autonomy 3/10

Verifies experiment design, execution integrity, and statistical validity of results.

SKILLS:

Experiment Audit

Design Verification

Execution Integrity

Statistical Validity

Bias Check



Fact Verification Agent

Fact Checker

Autonomy 3/10

Verifies factual claims against reliable sources with confidence scoring and citation.

SKILLS:

Fact Checking

Source Verification

Confidence Scoring

Citation

Claim Analysis



Financial Verification Agent

Financial Verifier

Autonomy 3/10

Verifies financial models, projections, and calculations for accuracy and reasonableness.

SKILLS:

Model Verification

Projection Audit

Calculation Check

Assumption Validation

Reasonableness



Hallucination Risk Agent

Hallucination Detector

Autonomy 3/10

Detects potential AI hallucinations in outputs by cross-referencing with known facts and sources.

SKILLS:

Hallucination Detection

Cross-Reference

Confidence Assessment

Source Grounding

Risk Flagging



Launch Readiness Verification Agent

Launch Verifier

Autonomy 3/10

Verifies launch readiness across all dimensions including technical, marketing, support, and legal.

SKILLS:

Launch Verification

Readiness Assessment

Cross-Team Check

Risk Review

Go/No-Go



Multi-Agent Crosscheck Agent

Crosscheck Coordinator

Autonomy 3/10

Coordinates verification across multiple agents to ensure consensus and catch disagreements.

SKILLS:

Cross-Agent Verification

Consensus Building

Disagreement Detection

Resolution

Confidence Aggregation



Output Consistency Agent

Consistency Checker

Autonomy 3/10

Checks consistency across agent outputs, ensuring aligned facts, numbers, and recommendations.

SKILLS:

Consistency Check

Cross-Output Validation

Alignment Verification

Contradiction Detection

Reconciliation



Quantitative Verification Agent

Quantitative Verifier

Autonomy 3/10

Verifies quantitative claims, calculations, and statistical analyses for accuracy.

SKILLS:

Calculation Verification

Statistical Review

Accuracy Check

Methodology Audit

Data Validation



Source Reliability Agent

Source Reliability Analyst

Autonomy 3/10

Assesses reliability of information sources using trust scoring, bias detection, and history analysis.

SKILLS:

- Source Assessment
- Trust Scoring
- Bias Detection
- History Analysis
- Credibility Rating



Corporate Strategy Agent

Corporate Strategist

Autonomy 4/10

Manages corporate-level strategy including competitive positioning, market expansion, M&A evaluation, and strategic partnerships.

SKILLS:

Corporate Strategy M&A Evaluation Market Positioning Partnership Strategy Competitive Analysis



Crisis Response Strategy Agent

Crisis Strategist

Autonomy 5/10

Develops crisis response strategies, coordinates rapid response across departments, and manages strategic communication during crises.

SKILLS:

Crisis Strategy Rapid Response Stakeholder Communication Damage Control Recovery Planning



Ecosystem Design Agent

Ecosystem Architect

Autonomy 4/10

Designs multi-product ecosystems, identifies synergies between ventures, and architects integration strategies.

SKILLS:

Ecosystem Design Synergy Identification Integration Strategy Platform Architecture Network Effects



Executive Summary Agent

Executive Summarizer

Autonomy 3/10

Creates concise, decision-ready executive summaries from complex data, reports, and multi-agent outputs.

SKILLS:

Executive Summaries Decision Briefs Data Synthesis Key Insight Extraction Board Reports



Long-Horizon Forecasting Agent

Long-Horizon Forecaster

Autonomy 4/10

Models long-term trends, creates multi-year projections, and identifies emerging opportunities and threats across industries.

SKILLS:

Trend Modeling Multi-Year Projections Opportunity Detection Threat Analysis Scenario Planning



Multi-Venture Coordination Agent

Venture Coordinator

Autonomy 4/10

Coordinates activities across multiple concurrent ventures, manages shared resources, and identifies inter-venture synergies.

SKILLS:

Venture Coordination Resource Sharing Synergy Detection Cross-Venture Planning Portfolio Optimization



Objective Prioritization Agent

Objective Priority Manager

Autonomy 4/10

Prioritizes organizational objectives using evidence-based frameworks, manages OKR cascades, and resolves priority conflicts.

SKILLS:

OKR Management Priority Frameworks Conflict Resolution Resource Prioritization Impact Assessment



Opportunity Portfolio Agent

Opportunity Portfolio Manager

Autonomy 3/10

Manages the pipeline of strategic opportunities, scores and prioritizes them, and tracks opportunity lifecycle.

SKILLS:

Opportunity Scoring Pipeline Management Priority Ranking Lifecycle Tracking Decision Support



Program Synthesis Agent

Program Synthesizer

Autonomy 4/10

Synthesizes complex multi-domain programs from strategic objectives, coordinating cross-functional execution plans.

SKILLS:

Program Design Cross-Domain Coordination Execution Planning Dependency Mapping Resource Estimation



Strategic Planning Agent

Strategic Planner

Autonomy 4/10

Develops long-range strategic plans, translates vision into executable multi-phase programs with measurable milestones.

SKILLS:

Strategic Planning Vision Translation Milestone Design Program Architecture Long-Range Forecasting



Strategic Review Agent

Strategic Reviewer

Autonomy 4/10

Conducts periodic strategic reviews, evaluates program effectiveness, and recommends course corrections.

SKILLS:

Strategy Review Performance Evaluation Course Correction Benchmark Analysis Gap Assessment



Strategic Thesis Agent

Strategic Thesis Developer

Autonomy 4/10

Develops and validates strategic theses for new markets, products, and ventures using evidence-based frameworks.

SKILLS:

Thesis Development

Market Validation

Evidence Synthesis

Hypothesis Testing

Strategic Framing



Counterfactual Analysis Agent Counterfactual Analyst Autonomy 3/10

Analyzes counterfactual scenarios to understand what would have happened under different decisions.

SKILLS:

- Counterfactual Analysis
- Alternative History
- Decision Review
- Outcome Attribution
- Learning



Ecosystem Dynamics Agent Ecosystem Dynamics Modeler Autonomy 3/10

Models ecosystem dynamics including competitive interactions, network effects, and market evolution.

SKILLS:

- Ecosystem Modeling
- Network Effects
- Competitive Dynamics
- Market Evolution
- System Dynamics



Failure Mode Forecast Agent Failure Mode Analyst Autonomy 3/10

Forecasts potential failure modes across systems, products, and operations with prevention strategies.

SKILLS:

- FMEA
- Failure Prediction
- Prevention Strategy
- Risk Prioritization
- Mitigation Planning



Long-Term Foresight Agent Long-Term Forecaster Autonomy 3/10

Projects long-term trends and their implications for strategy, products, and organizational design.

SKILLS:

- Long-Term Forecasting
- Trend Projection
- Implication Analysis
- Strategic Foresight
- Futures Analysis



Market Shock Analysis Agent Market Shock Analyst Autonomy 3/10

Analyzes potential market shocks and their impacts on portfolio, products, and operations.

SKILLS:

- Shock Analysis
- Impact Modeling
- Stress Testing
- Recovery Planning
- Resilience Assessment



Monte Carlo Agent Monte Carlo Analyst Autonomy 3/10

Runs Monte Carlo simulations for financial projections, risk assessment, and outcome distributions.

SKILLS:

- Monte Carlo Simulation
- Risk Distribution
- Financial Projection
- Confidence Intervals
- Sensitivity



Product Demand Forecast Agent Demand Forecaster Autonomy 3/10

Forecasts product demand using historical data, market signals, and seasonal patterns.

SKILLS:

- Demand Forecasting
- Historical Analysis
- Seasonal Patterns
- Market Signals
- Accuracy Tracking



Resource Stress Testing Agent Stress Tester Autonomy 3/10

Stress tests resource allocation under extreme conditions to identify breaking points and reserves needed.

SKILLS:

- Stress Testing
- Breaking Point Analysis
- Reserve Planning
- Extreme Scenarios
- Resilience



Risk Simulation Agent Risk Simulator Autonomy 3/10

Simulates risk scenarios including cascading failures, market shocks, and operational disruptions.

SKILLS:

- Risk Simulation
- Cascade Analysis
- Shock Modeling
- Disruption Planning
- Impact Assessment



Scenario Modeling Agent Scenario Modeler Autonomy 3/10

Creates detailed scenario models for strategic planning with probability-weighted outcomes.

SKILLS:

- Scenario Modeling
- Probability Weighting
- Outcome Analysis
- Branching Scenarios
- Strategic Planning



Sensitivity Analysis Agent Sensitivity Analyst Autonomy 3/10

Conducts sensitivity analysis on key variables to understand impact ranges and critical thresholds.

SKILLS:

- Sensitivity Analysis
- Variable Impact
- Threshold Detection
- Range Modeling
- Key Driver



Strategic Option Comparison Agent

Option Comparator

Autonomy 3/10

Compares strategic options using multi-criteria decision analysis with quantified tradeoffs.

SKILLS:

Option Comparison

Multi-Criteria Analysis

Tradeoff Quantification

Decision Matrix

Ranking



A/B Test Agent

A/B Test Manager

Autonomy 3/10

Designs and manages A/B tests including sample sizing, randomization, and result interpretation.

SKILLS:

A/B Testing

Sample Sizing

Randomization

Result Interpretation

Statistical Significance



Channel Testing Agent

Channel Tester

Autonomy 3/10

Tests different acquisition and distribution channels for effectiveness and scalability.

SKILLS:

Channel Testing

Effectiveness Measurement

Scalability Assessment

Cost Comparison

ROI Analysis



Creative Experiment Agent

Creative Experimenter

Autonomy 3/10

Tests creative variations including messaging, visuals, and formats across channels.

SKILLS:

Creative Testing

Message Variation

Visual Testing

Format Testing

Performance Ranking



Experiment Registry Agent

Experiment Registry Manager

Autonomy 2/10

Maintains the experiment registry tracking all active, completed, and planned experiments.

SKILLS:

Experiment Registry

Status Tracking

Result Archive

Learning Capture

Conflict Detection



Funnel Experiment Agent

Funnel Experimenter

Autonomy 3/10

Runs experiments on conversion funnels to identify and improve bottleneck stages.

SKILLS:

Funnel Experiments

Bottleneck Testing

Stage Optimization

Variant Design

Impact Measurement



Hypothesis Design Agent

Hypothesis Designer

Autonomy 3/10

Designs testable hypotheses with clear metrics, success criteria, and expected outcomes.

SKILLS:

Hypothesis Design

Metric Selection

Success Criteria

Expected Outcomes

Test Planning



Landing Page Experiment Agent

Landing Page Experimenter

Autonomy 3/10

Runs landing page experiments testing layouts, copy, CTAs, and conversion elements.

SKILLS:

Landing Page Testing

Layout Experiments

CTA Testing

Copy Variants

Conversion Optimization



Pricing Experiment Agent

Pricing Experimenter

Autonomy 3/10

Designs and runs pricing experiments to optimize revenue and customer willingness to pay.

SKILLS:

Price Testing

Willingness to Pay

Revenue Optimization

Elastic Analysis

Segment Pricing



Product-Market-Fit Experiment Agent

PMF Experimenter

Autonomy 3/10

Designs experiments to validate product-market fit including Sean Ellis test and usage metrics.

SKILLS:

PMF Testing

Sean Ellis Test

Usage Metrics

Retention Analysis

Fit Scoring



Statistical Confidence Agent

Statistical Analyst

Autonomy 3/10

Calculates and validates statistical confidence levels for experiment results and decisions.

SKILLS:

Confidence Calculation

Power Analysis

Sample Validation

Bias Detection

P-Value Analysis



Variant Analysis Agent

Variant Analyst

Autonomy 3/10

Analyzes experiment variants for statistical significance, effect size, and practical importance.

SKILLS:

Variant Analysis

Statistical Significance

Effect Size

Practical Importance

Recommendation



Action Execution Agent Action Executor Autonomy 3/10

Executes approved actions through the execution gateway with proper logging and verification.

SKILLS:

- Action Execution
- Gateway Integration
- Logging
- Verification
- Status Tracking



Approval-Aware Execution Agent Approval-Aware Executor Autonomy 3/10

Executes actions that require approval, managing the approval workflow before and after execution.

SKILLS:

- Approval Workflow
- Pre-Execution Check
- Post-Execution Logging
- Threshold Management
- Escalation



Billing Execution Agent Billing Executor Autonomy 2/10

Executes billing operations including invoicing, payment processing, and subscription management.

SKILLS:

- Invoice Generation
- Payment Processing
- Subscription Management
- Refund Processing
- Revenue Recording



Campaign Execution Agent Campaign Executor Autonomy 3/10

Executes marketing campaigns across channels with scheduling, tracking, and optimization.

SKILLS:

- Campaign Launch
- Channel Execution
- Scheduling
- Performance Tracking
- Optimization



Communications Execution Agent Communications Executor Autonomy 3/10

Executes communications including emails, notifications, and stakeholder updates.

SKILLS:

- Email Execution
- Notifications
- Stakeholder Updates
- Template Management
- Delivery Tracking



Deployment Execution Agent Deployment Executor Autonomy 3/10

Executes deployment operations with staged rollout, monitoring, and rollback capability.

SKILLS:

- Deployment
- Staged Rollout
- Monitoring
- Rollback
- Health Checks



Environment-Specific Execution Agent Environment Executor Autonomy 3/10

Manages environment-specific execution rules for sandbox, staging, and production.

SKILLS:

- Environment Rules
- Sandbox Execution
- Production Controls
- Staging Validation
- Environment Selection



Post-Execution Logging Agent Post-Execution Logger Autonomy 2/10

Logs all execution outcomes including success, failure, side effects, and performance metrics.

SKILLS:

- Execution Logging
- Outcome Recording
- Side Effects
- Performance Metrics
- Audit Trail



Rollback Trigger Agent Rollback Trigger Manager Autonomy 4/10

Detects conditions requiring rollback and triggers rollback procedures automatically.

SKILLS:

- Rollback Detection
- Auto-Trigger
- State Restoration
- Impact Assessment
- Communication



Safe Rollout Agent Safe Rollout Manager Autonomy 3/10

Manages safe rollouts with canary deployments, feature flags, and progressive delivery.

SKILLS:

- Canary Deployment
- Feature Flags
- Progressive Delivery
- Health Monitoring
- Auto-Rollback



Workflow Trigger Agent Workflow Trigger Manager Autonomy 3/10

Triggers automated workflows based on events, schedules, and conditional rules.

SKILLS:

- Workflow Triggers
- Event Processing
- Scheduling
- Conditional Logic
- Chain Execution



Alerting Agent Alert Communications Manager Autonomy 3/10

Manages alert communications including severity-based routing, formatting, and delivery tracking.

SKILLS:

- Alert Communication
- Severity Routing
- Format Templates
- Delivery Tracking
- Escalation Alerts



Board Memo Agent Board Memo Writer Autonomy 2/10

Drafts board memos, meeting materials, and governance documentation.

SKILLS:

- Board Memos
- Meeting Materials
- Governance Docs
- Resolution Drafts
- Compliance Reports



Briefing Agent Briefing Specialist Autonomy 3/10

Creates briefings for stakeholders including situation summaries, action items, and context.

SKILLS:

- Situation Briefs
- Action Summaries
- Context Packaging
- Stakeholder Tailoring
- Quick Updates



Executive Dashboard Agent Executive Dashboard Manager Autonomy 3/10

Maintains executive dashboards with real-time KPIs, alerts, and strategic metrics.

SKILLS:

- Dashboard Management
- KPI Updates
- Alert Integration
- Strategic Metrics
- Custom Views



Executive Reporting Agent Executive Reporter Autonomy 3/10

Creates executive-level reports with key metrics, strategic insights, and decision recommendations.

SKILLS:

- Executive Reports
- KPI Summaries
- Strategic Insights
- Decision Support
- Visual Reports



Internal Communications Agent Internal Comms Manager Autonomy 3/10

Manages internal communications including announcements, newsletters, and organizational updates.

SKILLS:

- Internal Comms
- Announcements
- Newsletters
- Org Updates
- Change Communication



Investor Relations Agent IR Communications Manager Autonomy 2/10

Manages investor communications including updates, reports, and relationship maintenance.

SKILLS:

- Investor Updates
- IR Reports
- Relationship Management
- Q&A Preparation
- Disclosure



Knowledge Distribution Agent Knowledge Distributor Autonomy 3/10

Distributes knowledge and insights to relevant stakeholders and teams.

SKILLS:

- Knowledge Distribution
- Insight Sharing
- Team Updates
- Learning Broadcast
- Targeted Delivery



Narrative Synthesis Agent Narrative Synthesizer Autonomy 3/10

Synthesizes complex information into coherent narratives for different audiences.

SKILLS:

- Narrative Synthesis
- Audience Adaptation
- Complexity Reduction
- Story Structure
- Message Clarity



Stakeholder Update Agent Stakeholder Update Manager Autonomy 3/10

Creates regular stakeholder updates with progress, risks, and action items.

SKILLS:

- Progress Updates
- Risk Communication
- Action Items
- Timeline Updates
- Stakeholder Mapping



Arbitration Agent

Arbitrator

Autonomy 4/10

Arbitrates disputes between agents or departments with evidence-based decisions.

SKILLS:

Arbitration

Evidence Review

Decision Making

Fairness Assessment

Binding Resolution



Budget Conflict Agent

Budget Conflict Resolver

Autonomy 3/10

Resolves budget conflicts between competing priorities with allocation recommendations.

SKILLS:

Budget Arbitration

Priority Balancing

Allocation Recommendation

Impact Analysis

Compromise Design



Conflict Detection Agent

Conflict Detector

Autonomy 3/10

Detects conflicts between agents, goals, resources, and policies across the system.

SKILLS:

Conflict Detection

Early Warning

Pattern Recognition

Scope Analysis

Priority Assessment



Decision Rationale Agent

Rationale Documenter

Autonomy 2/10

Documents decision rationale for conflict resolutions including evidence, reasoning, and precedent.

SKILLS:

Rationale Documentation

Evidence Linking

Reasoning Chain

Precedent Recording

Learning Capture



Evidence Consolidation Agent

Evidence Consolidator

Autonomy 3/10

Consolidates evidence from multiple sources for conflict resolution and decision support.

SKILLS:

Evidence Gathering

Source Consolidation

Credibility Assessment

Summary Creation

Gap Identification



Final Recommendation Agent

Final Recommender

Autonomy 4/10

Provides final resolution recommendations after all evidence and analysis is gathered.

SKILLS:

Final Recommendation

Decision Synthesis

Rationale Documentation

Implementation Plan

Follow-Up



Policy Conflict Agent

Policy Conflict Resolver

Autonomy 3/10

Resolves conflicts between competing policies with precedence analysis and recommendations.

SKILLS:

Policy Analysis

Precedence Rules

Conflict Resolution

Policy Updates

Guidance



Scope Conflict Agent

Scope Conflict Resolver

Autonomy 3/10

Resolves scope conflicts between projects and teams with boundary clarification.

SKILLS:

Scope Analysis

Boundary Clarification

Overlap Resolution

Responsibility Assignment

Agreement



Timeline Conflict Agent

Timeline Conflict Resolver

Autonomy 3/10

Resolves timeline conflicts between competing deadlines and resource constraints.

SKILLS:

Timeline Analysis

Deadline Negotiation

Resource Balancing

Priority Sequencing

Compromise



Tradeoff Analysis Agent

Tradeoff Analyst

Autonomy 3/10

Analyzes tradeoffs in conflicting decisions with quantified impact and recommendation.

SKILLS:

Tradeoff Analysis

Impact Quantification

Option Comparison

Recommendation

Sensitivity



Data Restoration Agent

Data Restorer

Autonomy 3/10

Restores data from backups with validation, integrity checks, and selective recovery.

SKILLS:

- Data Restoration
- Backup Recovery
- Integrity Validation
- Selective Recovery
- Point-in-Time



Disaster Recovery Agent

Disaster Recovery Manager

Autonomy 4/10

Manages disaster recovery planning and execution including cross-region failover.

SKILLS:

- DR Planning
- Failover
- Cross-Region
- RTO/RPO Management
- DR Testing



Failure Recovery Agent

Recovery Manager

Autonomy 4/10

Manages recovery from system failures with automated and manual recovery procedures.

SKILLS:

- Recovery Management
- Automated Recovery
- Manual Procedures
- State Restoration
- Validation



Partial Rebuild Agent

Partial Rebuild Manager

Autonomy 3/10

Manages partial system rebuilds when full rollback is not necessary or feasible.

SKILLS:

- Partial Rebuild
- Component Recovery
- Selective Restoration
- Dependency Handling
- Validation



Quarantine Agent

Quarantine Manager

Autonomy 4/10

Quarantines faulty components, agents, or workflows to prevent cascading failures.

SKILLS:

- Quarantine Management
- Isolation
- Impact Containment
- Analysis Support
- Release Planning



Resume Planning Agent

Resume Planner

Autonomy 3/10

Plans system resumption after incidents including dependency ordering and validation.

SKILLS:

- Resume Planning
- Dependency Ordering
- Validation Steps
- Progressive Restart
- Health Checks



Retry Strategy Agent

Retry Strategist

Autonomy 3/10

Designs retry strategies with exponential backoff, jitter, and circuit breaking.

SKILLS:

- Retry Strategy
- Backoff Design
- Jitter
- Circuit Breaking
- Timeout Management



Rollback Agent

Rollback Manager

Autonomy 4/10

Executes rollback procedures to restore system to a known-good state.

SKILLS:

- Rollback Execution
- State Restoration
- Data Recovery
- Version Revert
- Verification



Safe Shutdown Agent

Safe Shutdown Manager

Autonomy 4/10

Manages graceful system shutdowns with proper state preservation and resource cleanup.

SKILLS:

- Graceful Shutdown
- State Preservation
- Resource Cleanup
- Connection Draining
- Notification



State Reconciliation Agent

State Reconciler

Autonomy 3/10

Reconciles system state after recovery to ensure consistency across all components.

SKILLS:

- State Reconciliation
- Consistency Check
- Data Sync
- Conflict Resolution
- Integrity Verification



37 Core Systems & Capabilities

Everything the platform can do, explained simply

These are the 37 major tools and systems built into MAARS Command. Each one does something different, and they all work together. Think of them like apps on your phone -- each has a specific purpose, but they can share information with each other.

Autonomous Orchestration Engine -- The Boss Brain That Runs Everything

Think of this as the main brain of MAARS Command. When you tell it what you want to accomplish (like "launch a new product" or "write a marketing plan"), it breaks your big goal into smaller steps, figures out which AI agent is best at each step, and assigns the work automatically. It's like having a super-smart project manager who never sleeps, never forgets, and can manage hundreds of tasks at the same time.

HOW IT WORKS STEP BY STEP:

- You type in a goal in plain English (like "Create a social media campaign for our new app")
- The system scores your goal on how clear it is and how hard it will be to complete
- It creates a step-by-step plan with milestones (checkpoints) and individual tasks
- Each task gets automatically assigned to the AI agent who is best at that type of work
- Multiple tasks can run at the same time (not one-by-one), so things get done faster
- If something goes wrong with a task, the system automatically tries again with a different AI model
- It can also generate images (using GPT Image 1) and videos (using Sora 2) as part of the plan
- When tasks need real-world actions (like sending an email), it hands off to the Personal Secretary agent

Custom Brain Profiles -- Personalizing How Each Agent Thinks

Every AI agent has its own "brain profile" -- a set of settings that control how it thinks, talks, and works. You can change these settings anytime. It's like adjusting the personality and skills of each team member. Want one agent to be super creative? Turn up its creativity. Want another to be very precise and careful? Turn down its creativity and pick a more accurate AI model.

HOW IT WORKS STEP BY STEP:

- Each agent can use a different AI model (like GPT-5 for hard tasks, or a faster model for simple ones)
- Autonomy Level (1 to 10): Level 1 means the agent always asks you before doing anything. Level 10 means it does everything on its own
- Communication Style: You can set each agent to talk professionally, casually, technically, or creatively
- Tool Permissions: You decide what tools each agent can use -- web browsing, making images, running code, etc.
- Creativity Temperature (0.0 to 1.0): 0.0 means the agent gives the same answer every time (very predictable). 1.0 means it gets very creative and surprising
- Max Tokens: This controls how long the agent's responses can be (more tokens = longer answers)

Quality Control & Failure Recovery -- Making Sure Every Answer Is Good Enough

Every time an AI agent finishes a task, another AI (called the Critic) checks the work and gives it a score from 1 to 10. If the score is too low, the system tries again with a different AI model. If it still can't get a good result after several tries, it tells you so you can handle it yourself. This way, you never get bad work without knowing about it.

HOW IT WORKS STEP BY STEP:

- The Critic Module (powered by GPT-4o) automatically reviews every single task output
- It scores work on: completeness (did it finish everything?), accuracy (is it correct?), actionability (can you use it?), and professionalism (does it look good?)
- Each score is from 1 to 10. Higher is better
- If a task fails, the system automatically tries again using different AI models in this order: GPT-5 first, then Claude Sonnet, then Gemini Pro, then GPT-4o
- If ALL retries fail, the task gets flagged and you receive a notification so you can review it yourself
- The Quality Dashboard shows you pass rates (what percentage of tasks succeed), average scores, and how often recovery was needed



Smart Universal AI Router -- 33 Providers, 600+ Models — Routing the Right Brain for Every Job

MAARS Command has access to 600+ different AI models from 33 companies via the MAARS Universal API Gateway. The Smart Router classifies your prompt's task type (code, math, reasoning, creative, translation, research, etc.) and your current credit balance, then automatically picks the ideal model. Easy chats go to blazing-fast Groq or Cerebras. Complex reasoning goes to O4-mini or DeepSeek R1. Research goes to Perplexity Deep Research first. Developers get a maars-sk-* key and can call any model through a single OpenAI-compatible endpoint at /developer.

HOW IT WORKS STEP BY STEP:

- Premium Tier (hard tasks): O4, O4-mini, Claude Opus 4.6, Claude Sonnet 4.6, GPT-5, Grok 3 — smartest models for complex reasoning, legal, and research
- Standard Tier (medium tasks): GPT-4.1, Gemini 2.5 Pro, DeepSeek R1, Mistral Large, Cohere Command A — balanced quality and cost
- Economy Tier (easy tasks): Groq Llama 4 Scout, Cerebras Llama 3.3, Gemini Flash, DeepSeek Chat, Mistral Nemo — ultra-fast at rock-bottom cost
- Task classification: the router scores your prompt against 11 keyword banks (code, math, reasoning, legal, creative, translation, summary, research, data, vision, chat) to find the best match
- Credit-budget awareness: if you have ≤ 3 credits (micro), it routes to economy only. 11-50 credits = full routing. 51+ = best model, no restrictions
- Per-task preferred models: code → GPT-4.1 or Claude Sonnet 4.6 first; math → O4-mini or DeepSeek R1; research → Perplexity Deep Research; translation → DeepSeek or Qwen
- Manual override: pick any specific model in chat or set a default in Settings. Quality (economy/standard/premium) and task hint (code/math/etc.) can be saved as preferences
- All 33 providers are accessed directly via their own API keys — the MAARS Universal Key also covers OpenAI, Anthropic, and Gemini as a fallback



Autonomous Collaboration Engine -- Agents Automatically Working Together

When one agent finishes a task, the system automatically figures out if other agents should know about it. For example, if the Marketing agent creates a campaign plan, the system automatically tells the Creative agent (to make visuals) and the Technical agent (to build the landing page). This happens without you doing anything -- the agents coordinate by themselves.

HOW IT WORKS STEP BY STEP:

- All 458+ agents are organized into 9 work domains: executive, product, technical, creative, marketing, operations, finance, governance, and intelligence
- The system has built-in rules for which domains affect each other. Example: when marketing creates something, it triggers creative + technical
- There are 4 types of collaboration: information sharing (FYI), review request (please check this), data handoff (here's what you need), and coordination (let's work together)
- Each task creates up to 3 automatic collaborations to keep things focused and not overwhelm agents with notifications
- You can also manually start a collaboration between any two agents using the Collaboration Engine page



Universal Reference Intelligence -- Learning Your Style From Examples

Give this system any example of writing or design you like, and it will study it to understand the style. Upload a marketing email you love? It extracts the writing tone, word choices, and messaging patterns. Upload a brand image? It identifies colors, design patterns, and visual style. It saves all of this as a "Style Blueprint" that other tools can use to create content that matches your brand perfectly.

HOW IT WORKS STEP BY STEP:

- Text Analysis: paste in any text (an email, article, ad copy) and it identifies the writing style, tone, target audience, and messaging patterns
- Image Analysis: upload any image and it detects brand elements, visual style, emotional tone, color palette, and design patterns
- The output is a "Style Blueprint" -- a saved profile that captures everything about that style
- You can reuse Style Blueprints in the Content Generator to create new content that matches the same style
- History tracking keeps all your previous analyses so you can go back to them anytime



Content Generator -- Creating On-Brand Content Instantly

Need to write a blog post, social media caption, email, or ad? Just tell this tool what you need, optionally select a Style Blueprint to match a specific tone, and it generates professional content for you. It's like having a writing team that can produce any type of content in seconds, and it always matches your brand voice.

HOW IT WORKS STEP BY STEP:

- 8 content types you can generate: marketing copy, social media posts, emails, blog articles, ad copy, press releases, brand guidelines, and custom (anything else)
- Style Blueprint integration: pick a blueprint you created in Reference Intelligence, and the content will match that exact tone and style
- Length control: choose short (100-200 words), medium (300-500 words), or long (800-1500 words)
- Optional tone override: even with a blueprint, you can adjust the tone for specific pieces
- Full history: every piece of content you generate is saved. You can copy it, reuse it, or delete it



Vibe Coding App Builder -- Build Apps Just By Talking

Describe an app you want to build using plain English, and this tool creates a complete working web application for you. No coding knowledge needed. You can say things like "make me a todo list app with a dark theme" and it builds it instantly. You can keep talking to modify it -- "add a calendar view" or "change the color to blue" -- and it updates in real time.

HOW IT WORKS STEP BY STEP:

- Chat-based interface: just describe what you want in normal English
- Creates complete HTML/CSS/JavaScript applications with modern design (Tailwind CSS)
- Live preview: see your app running in real-time right next to the chat
- Iterative modifications: keep chatting to change, add, or remove features
- Code view: switch to see the actual code if you want to learn or customize further
- Download: export your finished app as a file you can host anywhere



Agent Activity Monitor -- Watching All Agents Work in Real-Time

This is a live dashboard that shows you what every AI agent is doing right now. It updates every 8 seconds using a real-time connection (WebSocket). You can see which agents are active, what they're working on, who they're talking to, and how fast they're completing tasks. Think of it like a security camera system, but for your AI workforce.

HOW IT WORKS STEP BY STEP:

- Real-time updates every 8 seconds using WebSocket (a fast, always-on internet connection)
- Shows agent status: which agents are active, idle, or working on tasks
- Communication flows: see messages being sent between agents (who sent what to whom)
- Task dependency graph: a visual map showing how tasks relate to each other and their current status
- Tool execution tracking: see when agents use tools (like web search or image generation) and how long each tool takes
- Pause/Resume: you can pause the live feed to examine a specific moment, then resume
- Connection indicator: shows if you're connected via WebSocket (fast), REST polling (slower backup), or Paused

Simulation vs. Execution Mode -- A Safety Switch for Real-World Actions

MAARS Command has a master safety switch. In Simulation Mode (the default), agents can plan and think, but they CAN'T do anything in the real world -- no real emails are sent, no real calendar events are created, no real social media posts go live. Everything is pretend. When you're ready to go live, you flip the switch to Execution Mode, and agents can perform real actions. This keeps you safe while testing.

HOW IT WORKS STEP BY STEP:

- System-wide toggle: one switch controls ALL agents at once (found on the KPI Dashboard)
- Simulation Mode (default): all real-world actions return fake, safe responses. Nothing actually happens in the real world
- Execution Mode: agents send real emails via Gmail, create real Google Calendar events, and interact with real APIs
- Visual indicators: the UI clearly shows which mode you're in, so you always know if actions are real or simulated
- You can switch between modes at any time -- no restart needed

Real-World Action Layer -- Letting Agents Do Things in the Real World

This is the system that actually connects agents to real-world services. Right now it supports Google Suite (Gmail for sending emails and Google Calendar for creating events). When an agent needs to send an email or schedule a meeting, this layer handles the actual sending. All actions are controlled by the Simulation/Execution toggle -- nothing real happens unless you explicitly allow it.

HOW IT WORKS STEP BY STEP:

- Google Gmail: agents can compose and send real emails on your behalf
- Google Calendar: agents can create, update, and manage calendar events
- OAuth 2.0 authentication: you securely connect your Google account (your password is never shared with MAARS)
- You connect and disconnect from the Settings page
- Integration status dashboard shows which services are connected and working
- The Personal Secretary agent is the one who actually performs these actions (other agents request through it)

MAARS Universal API Gateway -- 600+ Models · 33 Providers · One OpenAI-Compatible Key

MAARS is a self-hosted universal AI gateway — on steroids. A single maars-sk-* API key lets you call any of 175,000+ models across 33 providers using the same OpenAI SDK you already use. MAARS smart routing aliases (maars/auto, maars/code, maars/fast...) automatically pick the best model for your task and budget. The full Developer Portal is at /developer.

HOW IT WORKS STEP BY STEP:

- OpenAI: GPT-5, GPT-4.1/mini/nano (1M ctx), GPT-4o, O4 & O4-mini & O3 & O3-pro (reasoning/math)
- Anthropic: Claude Opus 4.6 & 4.5 (deepest reasoning), Claude Sonnet 4.6 & 3.7 (creative + code), Claude Haiku 4.5 (fast + cheap)
- Google Gemini: Gemini 2.5 Pro (2M ctx research), Gemini 2.5 Flash (speed), Gemma 3 open models (1B–27B)
- xAI Grok: Grok-4 & Grok-4 Fast Reasoning (256K ctx), Grok-3, Grok-2 Vision — made by Elon Musk's xAI
- DeepSeek: DeepSeek V3-0324 (code + data, ultra cheap), DeepSeek R1-0528 (math + reasoning, open-source champ)
- Mistral AI: Mistral Large, Magistral Medium (reasoning), Codestral (code-specialist, 256K ctx), Pixtral Large (vision)
- Perplexity: Sonar Deep Research (live web research), Sonar Pro, Sonar Reasoning Pro — web-grounded answers
- Cohere: Command A (enterprise docs), Command R+ — optimized for business document retrieval
- Groq + Cerebras: blazing-fast open-source inference — Llama 4, Qwen QwQ-32B at sub-second latency
- Together AI (39 models), Fireworks AI (20), SambaNova (14), NVIDIA NIM (12), Novita AI (12), Lepton AI (8), Lambda Labs (6): open-source hosting
- AI21 Jamba (256K ctx), Moonshot Kimi (128K), Qwen/Alibaba (multilingual), Minimax AI, Inception AI (Mercury), Arcee AI, Amazon Bedrock (Nova), Yi, Zhipu, Doubao, Hyperbolic, Upstage, LLaMA, Writer, HuggingFace
- 10 MAARS smart routing aliases: maars/auto, maars/smart, maars/economy, maars/premium, maars/code, maars/vision, maars/reasoning, maars/search, maars/fast, maars/standard
- Live model validation: before every call, the gateway checks the provider's live /models API and auto-fallbacks if a model is not yet live
- Webhooks: register HTTPS endpoints for budget alerts (75%/90%/100%), rate limit events, and request completion — signed with HMAC-SHA256



Voice Command Interface -- Talk to MAARS Instead of Typing

Instead of typing, you can just talk. Click the microphone button, say what you want (like "show me my agents" or "open the dashboard"), and MAARS Command understands you and takes action. Your voice is converted to text using OpenAI's Whisper technology, which understands 50+ languages.

HOW IT WORKS STEP BY STEP:

- Click the microphone button in the Command Palette (or press / then click the mic icon)
- Speak naturally in any of 50+ supported languages
- Your speech is recorded using your browser's microphone
- OpenAI Whisper converts your speech to text in real-time
- The text automatically fills the search box, so you can navigate, search agents, or trigger actions
- Visual feedback: the mic pulses red while recording, shows a spinner while transcribing



Admin Code Explorer -- Browse the Entire Codebase

For admin users who want to look under the hood, the Code Explorer lets you browse every file in the MAARS Command codebase. You can navigate through folders, read any file with color-coded syntax highlighting, search for specific files, and even download the entire codebase as a ZIP file. It's like having a built-in file manager for the whole system.

HOW IT WORKS STEP BY STEP:

- Full file tree: browse all backend and frontend directories with expand/collapse folders
- Code viewer: read any file with line numbers, automatic language detection, and a copy button
- File search: find files by name across all directories
- ZIP export: download the complete codebase (backend + frontend) as a single ZIP file
- Language-colored icons: Python files are blue, JavaScript files are yellow, JSON is green, etc.
- Security: only admin users can access this feature, and path traversal attacks are blocked



Memory Governance -- The AI's Long-Term Memory System

Just like humans remember important things, MAARS Command has a memory system. Memories are facts, decisions, preferences, and knowledge that agents learn over time. You can create, edit, and delete memories manually. Old or unimportant memories automatically fade over time (like forgetting what you had for lunch last month), keeping the system focused on what matters most.

HOW IT WORKS STEP BY STEP:

- Full control: create, read, update, and delete memory entries
- Category tagging: organize memories by type (facts, decisions, preferences, instructions, context)
- Version history: every change to a memory is tracked with timestamps and reasons for the change
- Time-decay scoring: memories become less relevant over time (30-day half-life). Frequently accessed memories stay stronger
- Relevance = (time decay score) + (how often it's accessed) + (importance rating you give it)
- Auto-pruning: you can preview and remove low-relevance memories to keep the system clean
- Usage statistics: see total entries, average relevance score, and breakdown by category
- 500-entry limit per user with a percentage indicator showing how much you've used

Agent Memory Auto-Learning -- Agents Get Smarter After Every Task

Every time an agent completes a task, the system automatically extracts the most important lessons learned and saves them as memories. Over time, this creates a growing knowledge base that makes every agent smarter. It's like how you get better at your job the more experience you have -- except the AI remembers everything perfectly.

HOW IT WORKS STEP BY STEP:

- After every completed task, GPT-4o-mini automatically extracts 1-3 key learnings
- Learnings are saved as: facts, decisions, preferences, instructions, or context
- Each learning gets an importance score based on how significant the task was and what was learned
- Near-duplicate detection: if the system already knows something similar, it won't save a duplicate
- Auto-learned entries are marked with a special cyan badge in the Memory Governance page so you can tell them apart from manually created memories
- This runs in the background and never slows down task completion

Command Palette -- Instant Search & Navigation (Like Spotlight)

Press the / key (or Cmd+K on Mac) and a search box pops up. Start typing anything -- a page name, an agent name, or an action - and it instantly shows matching results. Select one and you're taken there immediately. It's just like Spotlight on Mac or the search bar in VS Code. It searches across all 50+ pages, all 458+ agents, and all quick actions.

HOW IT WORKS STEP BY STEP:

- Open it with / key or Cmd+K (Ctrl+K on Windows)
- Searches across: all pages (Dashboard, Agents, Chat, etc.), all 458+ agents, and quick actions
- Keyboard navigation: use arrow keys to move up/down, Enter to select, Escape to close
- Recent searches: your last searches are remembered so you can quickly redo them
- Results are grouped into categories: Recent, Pages, Agents, Quick Actions
- Voice command: click the microphone icon to search by speaking instead of typing
- Also accessible from the search button in the sidebar

Agent Team Builder -- Create Custom Teams for Projects

Pick agents from the full 458+ agent workforce and group them into custom teams. Give each team a name and a mission (like "Q1 Marketing Team" or "Product Launch Crew"). This makes it easy to organize agents for specific projects, see which skills your team covers, and quickly deploy the right group of agents for any job.

HOW IT WORKS STEP BY STEP:

- Search and filter agents by name, role, network category, or specific capability
- Create as many teams as you want with custom names and mission descriptions
- Select agents from the full 458+ workforce pool -- no limits on team size
- Team cards show: member count, network coverage (how many different specialties), and avatar previews
- Edit or delete teams anytime
- Teams are saved to your account so they persist between sessions

🔄 Trust Analytics & Scoring -- Rating How Reliable Each Agent Is

Every agent gets a trust score from 0 to 100 based on its track record. Did it complete tasks on time? Were the results good? Did it cause any errors? The Trust Analytics dashboard shows you which agents are your most reliable performers, which ones need improvement, and how trust levels are trending over time. It's like a performance review for your AI workforce.

HOW IT WORKS STEP BY STEP:

- Trust scores range from 0 (never trust) to 100 (completely reliable)
- Scores are calculated from: success rate (did tasks complete?), latency (how fast?), quality (how good?), and consistency (same quality every time?)
- Trust Distribution chart: see how many agents fall into Excellent (90-100), Good (70-89), Fair (50-69), and Poor (below 50) ranges
- Top Performers: a ranking of your best agents
- Needs Improvement: a ranking of agents that need attention
- Anomaly detection: if an agent's performance suddenly drops, the system flags it
- Historical trend graphs: see how trust scores change over time

⚡ Workflow Builder -- Create Automated Step-by-Step Processes

A visual drag-and-drop tool for building automated workflows. Think of it like building a flowchart: connect different steps together, and MAARS Command runs them automatically. For example: "When a new lead comes in, have the Research agent look them up, then have the Sales agent write a personalized email, then have the Secretary agent send it."

HOW IT WORKS STEP BY STEP:

- Visual canvas: drag and drop nodes (steps) onto a canvas and connect them with lines
- Node types: Agent Task (assign work to an agent), Decision (if/then branching), Parallel Split (do multiple things at once), Merge (wait for all parallel tasks to finish), Trigger (what starts the workflow)
- Conditional branching: set rules like "if the quality score is above 7, proceed; otherwise, redo"
- Workflow templates: pre-built workflows for common business processes to get you started
- Execution history: see every time a workflow ran, how long it took, and if anything failed
- Triggers: start workflows on a schedule (daily, weekly) or when specific events happen

> MAARS Developer API Gateway -- One Key · 600+ Models · 33 Providers · OpenAI-Compatible

MAARS is a self-hosted universal AI gateway — on steroids. Every registered user gets a `maars-sk-*` API key. With one key and a single `base_url` change, developers can call any of 175,000+ AI models across 33 providers using their existing OpenAI SDK code. The gateway validates models live against provider APIs, automatically falls back when a model is unavailable, fires signed webhooks on budget thresholds, and lets you run side-by-side model comparisons in one request.

HOW IT WORKS STEP BY STEP:

- Developer Portal at `/developer`: 4 tabs — My API Key (masked/reveal/copy), Model Browser (175,000+ models, searchable, filterable), Playground (test any model in browser), Quickstart (copy-paste code for Python/JS/cURL)
- OpenAI-compatible: `POST /v1/chat/completions` — change `base_url` from `api.openai.com` to your domain. No other code changes needed
- 609+ real models: OpenAI (35), Groq (35), Together AI (77), Novita AI (56), Fireworks AI (44), HuggingFace (38), Qwen/Alibaba (32), Mistral (31), Lepton AI (28), NVIDIA NIM (27), SambaNova (20), Lambda Labs (20), Cerebras (18), Google/Gemini (18), Hyperbolic (14), Anthropic (13), DeepSeek (12), xAI (11), Zhipu/GLM (11), ByteDance/Doubao (10), Cohere (9), Moonshot (8), Meta Llama API (6), Perplexity (6), 01.AI/Yi (5), Writer (5), AI21/Jamba (4), Amazon Bedrock (4), Minimax (4), Upstage (3), Arcee AI (3), Inception AI (2), ElevenLabs (voice)
- 10 MAARS smart routing aliases: `maars/auto`, `maars/smart`, `maars/economy`, `maars/standard`, `maars/premium`, `maars/code`, `maars/vision`, `maars/reasoning`, `maars/search`, `maars/fast`
- Live model validation: before each call, checks provider's `/models` API (cached 1 hour). If model not found, skips to fallback and includes `model_warning` in response
- Auto-fallback chain: tries up to 4 providers on failure — GPT-4.1 → Claude Sonnet → Gemini Flash → DeepSeek
- Webhooks: register HTTPS endpoints for `budget.75`, `budget.90`, `budget.100`, `rate_limit.exceeded`, `request.completed` — all signed with HMAC-SHA256
- Model comparison: `POST /v1/models/compare` runs the same prompt on up to 4 models in parallel and returns latency, cost, tokens, and content per model side-by-side
- Real token billing from provider responses, not estimates. Full parameter passthrough: `temperature`, `top_p`, `max_tokens`, `tools`, `vision`, `seed`, `response_format`



Campaign Builder -- Plan Multi-Channel Marketing Campaigns

Plan and run marketing campaigns across multiple channels (email, social media, web, ads) all in one place. Define your campaign goals, pick which channels to use, and let AI agents create content variations for each channel. Track how each campaign performs and let the AI optimize your budget allocation.

HOW IT WORKS STEP BY STEP:

- Multi-channel support: email, social media, website, and paid advertising
- AI-generated content: agents create different versions of your message for each channel automatically
- A/B testing: create multiple variations and see which performs better
- Performance tracking: see opens, clicks, conversions, and engagement for each campaign
- Budget allocation: set your budget and let the AI suggest how to split it across channels
- Campaign templates and scheduling: plan campaigns in advance and have them launch automatically
- Works with Content Generator and Reference Intelligence for on-brand content



Social Media Command Center -- 10 Platforms · AI-Powered Posting, Messaging & Geo-Targeted Boosting

MAARS agents can now operate across 10 major social media and messaging platforms simultaneously. From composing geo-targeted posts in the right language, to sending cold DMs, running multi-region ad boosts, broadcasting on WhatsApp and Telegram, and making cold calls via Twilio — all from a single command center. This isn't traditional social media management. It's an AI-driven omnichannel growth engine.

HOW IT WORKS STEP BY STEP:

- 10 platforms connected: Facebook, Instagram, Twitter/X, TikTok, WhatsApp Business, Viber, LINE, LinkedIn, YouTube, Telegram
- Post & Publish: agents auto-generate and publish content across all connected platforms in one click — with platform-specific formatting
- Geographic Targeting: select target regions (North America, Southeast Asia, Middle East, etc.) and content is auto-published with the right language for each region
- Non-Traditional Boosting: beyond basic boosts — A/B test ad creatives per region, TikTok Spark Ads, LinkedIn Lead Gen forms, Facebook dynamic creatives, YouTube TrueView campaigns
- Auto-Translation: agents translate your content to Japanese (Japan), Arabic (Middle East), Thai (Southeast Asia), Hindi (South Asia), etc. — automatically matched to your target region
- Cold Email Outreach: send personalized cold emails via SendGrid or Resend with geo-aware subject lines and body copy
- Cold Calling: initiate Twilio voice calls with AI-generated scripts delivered in the target region's language (English, Spanish, French, Arabic, Hindi, and more)
- Direct Messaging: send DMs on WhatsApp, Instagram, Viber, LINE, and Telegram — including WhatsApp Business templates and LINE rich messages
- Reply to Everything: agents can reply to comments, mentions, and DMs across all platforms — in the right language
- Scheduling: queue posts, calls, meetings, and campaigns for future dates with full calendar view
- Multi-Platform Campaigns: one campaign name, content in, agents adapt and post across all selected platforms simultaneously
- Omnichannel Analytics: track reach, impressions, clicks, and conversions per platform and per geographic region



Task Graph Execution Kernel -- DAG-Based Task Orchestration with Dependency Resolution

The Kernel turns any goal into a directed-acyclic task graph (DAG) and executes it. Instead of running steps one-by-one, it figures out which tasks depend on which, runs independent branches in parallel, and holds dependent tasks until their inputs are ready. Built-in retry logic, budget enforcement per branch, and checkpointed resumption let a workflow survive provider outages, human pauses, and multi-hour runs without losing state.

HOW IT WORKS STEP BY STEP:

- 131 dedicated Infinity endpoints under `/api/infinity` (task-graphs, goals, scheduler, budget, policies, runtime, recovery, workers, environments)
- Task graph compiler: accepts a natural-language goal → emits a DAG with nodes (agent/tool calls), edges (data dependencies), checkpoints, and rollback markers
- Parallel branch execution: independent branches run concurrently, respecting per-branch credit budgets and per-agent rate limits
- Deterministic resumption: every step emits a checkpoint so a halted graph can resume mid-run without re-executing completed nodes
- Scheduler + worker pool: `/kernel/scheduler.py` dispatches ready nodes to worker processes; job queue tracked at `/api/infinity/workers/*`
- Budget controller enforces per-user, per-agent, per-environment credit caps — hits a ceiling, the graph pauses instead of overspending

10-Tier Autonomy Framework -- From Full Human Approval (T1) to Unsupervised Autonomy (T10)

Every agent runs at one of ten autonomy tiers. Tier 1 means every action needs explicit human approval before it leaves the sandbox. Tier 10 means the agent operates unsupervised within its budget and policy envelope. Tiers 2-9 interpolate — some tools are pre-approved, some actions auto-execute if trust score > threshold, some require dual-agent cross-verification. You dial autonomy per agent, and the kernel enforces the tier on every tool call.

HOW IT WORKS STEP BY STEP:

- Tier 1 — Supervised: every action queues for human approval in the Approvals page
- Tier 2-3 — Assisted: non-destructive reads auto-approve; writes and external actions require approval
- Tier 4-6 — Delegated: pre-approved tools run freely; high-cost or high-impact actions gate on approval
- Tier 7-9 — Empowered: all allow-listed actions run; only novel tool calls or policy edge cases escalate
- Tier 10 — Autonomous: full autonomy within budget + policy; post-facto audit only
- Per-agent override: set any agent's tier from the Brain Profiles page; changes logged immutably
- Trust-gated elevation: an agent's effective tier is $\min(\text{assigned_tier}, \text{ceil}(\text{trust_score} / 10))$ — low trust can't bypass review
- Tier enforcement lives in `backend/governance/autonomy.py` and is checked on every tool invocation before execution

Embedded Browser Runtime -- Self-Hosted Chromium · Multi-Tab · Three-Way Shared Control · Autonomous Vision Agent

MAARS ships with a full Chromium runtime — not a call out to a remote browser service, not a thin iframe. The binary lives inside the MAARS install under `backend/browser_data/`, so when you container-ize or ship MAARS the browser comes with it. Agents, the orchestrator, and the human user all drive the same session: when the autonomous BrowserAgent hits a 2FA screen it hands the driver lock to you, you complete the one-time code, and the agent picks back up where it left off — cookies preserved, no re-login. Every one of 29 integrations has a click-to-connect button that opens the right OAuth URL in a tagged tab and optionally invokes the agent to walk through the flow.

HOW IT WORKS STEP BY STEP:

- Chromium runs on a dedicated OS thread with a `ProactorEventLoop` — sidesteps the `uvicorn --reload SelectorEventLoop` subprocess bug on Windows
- Multi-tab per session: Chrome-style tab strip, click to switch, × to close, + for blank — context shared across tabs so cookies/localStorage are unified
- Three-way shared control: driver $\in \{\text{shared}, \text{user}, \text{agent}, \text{system}\}$; take-control / release-control via REST, WS message, or UI buttons
- Autonomous BrowserAgent (`services/browser_agent.py`): perceive → decide (vision) → act loop with 2FA / CAPTCHA / payment hand-off to the human
- Vision fallback chain: OpenAI GPT-4o → Anthropic Claude. Returns strict JSON `{ kind, x, y, text, url, reason, confidence }` — callable directly via `browser_see_and_act`
- WebSocket live view: event-driven or continuous streams, PNG (lossless) or JPEG (~10× smaller), 1–24 fps, user input events forwarded back to Playwright
- Integration Connect: `POST /api/browser/integrations/{id}/connect` — opens the right login URL in a tagged tab, optional `run_agent=true` drives OAuth automatically
- Governance: per-env domain allow/deny policy (fail-closed), per-user browser-minutes daily budget (default 60 min), 429 on exhaust
- 11 agent tools: `browser_open`, `browser_navigate`, `browser_click`, `browser_fill`, `browser_type`, `browser_extract`, `browser_screenshot`, `browser_evaluate`, `browser_close`, `browser_see_and_act`, `browser_run_goal`
- 32 routes under `/api/browser/*` incl. sessions, tabs, take/release-control, screenshot, extract, evaluate, agent/run (SSE), policy/domains, policy/budget, usage, `integrations/{id}/connect`
- Frontend at `/browser`: session sidebar, integration grid (29 chips), tab strip, URL bar, driver status strip, agent goal runner with SSE event log, fullscreen mode, stream config

Multi-Environment Segregation -- Production · Staging · Sandbox · Simulation — Isolated Policies & Budgets

MAARS ships with four isolated execution environments. Agents, workflows, and data in Simulation never touch Production. Each environment has its own budget multiplier, policy set, data retention rule, and audit trail. Promote from Sandbox → Staging → Production only after Verification Civilization signs off. Simulation runs at zero real cost (all external calls mocked) so you can stress-test autonomous behavior without risk.

HOW IT WORKS STEP BY STEP:

- Simulation: cost multiplier = 0, all integrations mocked, full logging, 7-day retention — safe for dry-runs and chaos testing

- Sandbox: real calls to test-mode APIs (Stripe test keys, Gmail test inbox), 30-day retention, relaxed policies
- Staging: production-like credentials, strict policies, 90-day retention, required for release promotion
- Production: real money, real users, immutable audit, 7-year retention, all circuit breakers enabled
- Environment-scoped budgets: cap spend per env independently — e.g., \$10/day sandbox, \$2,000/day production
- Env-scoped trust scores: an agent's trust in sandbox doesn't inherit to production until it's promoted
- Managed via `backend/governance/environments.py` and `/api/infinity/environments/*` endpoints



Verification Civilization -- Multi-Verifier Fact-Checking with 6-Dimensional Consensus Scoring

Every factual claim an agent makes gets cross-verified by 2-4 independent verifiers (different models, different prompts, different retrieval sources). A claim only passes if consensus is reached across six dimensions: factual accuracy, source grounding, logical coherence, completeness, bias detection, and temporal freshness. Disagreements escalate to a human reviewer or a higher-tier verifier. The whole pipeline is inspired by Popperian epistemology — agents don't just "answer," they make claims that survive attempted refutation.

HOW IT WORKS STEP BY STEP:

- 6-dimensional scoring: factual accuracy, source grounding, logical coherence, completeness, bias, temporal freshness (0-100 each)
- Multi-model crosscheck: the same claim gets verified by e.g. Claude Opus + Gemini 2.5 Pro + DeepSeek R1 — disagreement = flag
- Source citations mandatory: every fact cites a retrievable source from the knowledge_chunks collection or live web search
- Hallucination detector: detects confabulation patterns (unverifiable specifics, made-up URLs, invented quotes) and quarantines the response
- Escalation routes: low-confidence → retry with higher-tier model; persistent disagreement → queue for human review
- Verification metrics tracked at /api/infinity/verification/stats — see pass rate, top failure modes, per-agent accuracy
- Engine lives in backend/verification/engine.py; exposed via /api/infinity/verification/* (4 routes)



Hierarchical Memory System -- Working · Episodic · Semantic · Knowledge Graph — 4 Tiers, Automatic Decay

Human-inspired memory architecture with four tiers. Working memory holds the current task's context (capped at 50 items). Episodic memory logs every event chronologically for later recall. Semantic memory stores concept embeddings for similarity search. The Knowledge Graph captures entity-relationship edges so agents can traverse "who knows what, connected to which, via which relationship." Old items automatically decay unless reinforced by usage, keeping the system focused without manual pruning.

HOW IT WORKS STEP BY STEP:

- Working memory (backend/memory_system/working.py): per-task scratchpad, 50-item cap, purged on task completion
- Episodic memory (episodic.py): time-ordered event log with time-decay scoring (30-day half-life), reinforces on recall
- Semantic memory (semantic.py): vector embeddings, cosine-similarity retrieval, concept clustering
- Knowledge Graph (knowledge_graph.py): entity + relationship store, traversable, Neo4j-style queries
- 16 REST endpoints under /api/infinity/memory/* for CRUD, search, promotion between tiers, and decay tuning
- Auto-promotion: frequently-accessed episodic items get promoted to semantic; heavily-referenced semantic clusters get nodes in the graph
- 500-entry soft cap per user with auto-pruning of lowest-relevance items when full
- Full UI at frontend/src/pages/MemoryHierarchy.jsx with tier breakdowns, decay charts, and per-entry relevance scoring



Recovery & Incident Response -- Quarantine · Rollback · Retry — Automatic Blast-Radius Containment

When something breaks (agent misbehavior, provider outage, runaway spend, repeated verification failure), the Recovery system isolates the problem instead of letting it cascade. Misbehaving agents get quarantined (tier forced to 1, tools stripped, flagged for review). Broken task graphs roll back to the last clean checkpoint. Transient failures retry with exponential backoff across alternate providers. Every incident gets a root-cause record with timeline, blast radius, and remediation steps.

HOW IT WORKS STEP BY STEP:

- Quarantine: an agent flagged for review loses all tools and drops to tier 1 until a human clears it
- Checkpointed rollback: any task graph can roll back to its last successful checkpoint — partial progress preserved
- Smart retry: failed LLM calls retry with a different model (GPT-5 → Claude Sonnet → Gemini Pro → GPT-4o) before escalating
- Incident ledger: every fault creates an incident record with timeline, severity, affected users/agents, and linked audit entries
- Circuit breakers (governance/circuit_breaker.py): 5 default breakers monitor spend spikes, error rate, retry storms, provider downtime, policy violations
- Recovery endpoints (/api/infinity/recovery/*): 5 routes for quarantine, rollback, retry, status, clear
- Post-incident reports auto-generated with root cause, blast radius, remediation, and preventive action taken



Approval Workflow Gates -- Human-in-Loop Execution · Approve · Reject · Edit · Escalate

When an agent hits a step that requires explicit approval (per its autonomy tier, per a policy, or per a workflow gate), the action queues in the Approvals page. A human reviews the proposed action, the agent's reasoning, the expected cost/impact, and signs off — or rejects, edits the payload, or escalates to a higher role. Every decision is logged immutably with approver identity, timestamp, and rationale. The agent resumes from the gate with the approved (or modified) payload.

HOW IT WORKS STEP BY STEP:

- Gate types: tool-call approval, budget exceedance, policy exception, cross-environment promotion, high-impact-action review
- Approver experience: rich preview of proposed action (full prompt, parameters, cost estimate, similar past actions)
- Decision options: Approve · Reject with reason · Edit payload and approve · Escalate to role
- Multi-signature gates: high-risk actions can require 2+ approvers before proceeding
- SLA tracking: gates that sit unapproved past a deadline auto-escalate or auto-timeout per policy
- Every decision written to immutable audit log — who approved what, when, why, with what edits
- Backend: /kernel/approval_controller.py + 7 endpoints under /api/approvals and /api/infinity/approvals/*



Knowledge Base & RAG Engine -- Document Uploads · Per-Agent Knowledge · Cited Retrieval

Upload PDFs, DOCX, Markdown, or plain text and attach them to any agent. The Knowledge Base chunks and embeds each document, stores the chunks in MongoDB's knowledge_chunks collection, and retrieves the top-K relevant chunks at chat time. Every agent answer cites the exact source chunks it drew from. Combined with the 58k skill library and live web search, every agent grounds its replies in three retrieval sources: skills (curated), uploads (yours), and web (live).

HOW IT WORKS STEP BY STEP:

- Supported upload formats: PDF (via PyMuPDF), DOCX (python-docx), Markdown, plain text; up to 25 MB per file
- Automatic chunking: semantic paragraph chunks with 200-token overlap for retrieval quality
- Per-agent scoping: attach a document to one agent or a team — retrieval only pulls from that agent's documents
- Cited retrieval: every chat response shows which chunks (with doc title + page/section) grounded each claim
- Three-layer RAG: agent's skill library (from .claude/skills/) + its knowledge base (your uploads) + live web search (Perplexity)
- GET /agents/{id}/brain exposes the full chunk ledger for any agent — transparent debugging
- UI: KnowledgeBaseTab with upload, preview, chunk inspector, and re-embed controls
- Powered by backend/services/rag_service.py + 5 routes under /api/knowledge-base



Integration Hub -- 29 External Providers · OAuth, API Keys, Webhooks — One Control Plane

Every external service MAARS talks to is configured through one Integration Hub. OAuth providers (Google Suite, GitHub), API-key providers (Stripe, Twilio, SendGrid, ElevenLabs, all 33 LLM providers), webhook-based integrations (Slack, custom), and the self-hosted Embedded Browser all share the same credential vault, status dashboard, rate-limit tracker, and health checker. Flip an integration on/off globally; rotate keys without code changes; see per-integration call counts and error rates at a glance.

HOW IT WORKS STEP BY STEP:

- 29 integrations: Stripe, Gmail, Google Calendar, Twilio, SendGrid, ElevenLabs, Whisper, GitHub, Salesforce, HubSpot, Shopify, Airtable, Slack, Zapier-style webhooks, Embedded Browser (self-hosted Chromium), plus all 33 LLM providers
- Categorized: Payments · Email · Calendar · SMS/Voice · CRM · E-commerce · Data · Dev Tools · Social Media · Search · Voice/Audio · LLM
- OAuth flows for Google Suite and GitHub; API-key config for the rest; signed-webhook receivers for inbound events
- Per-integration health dashboard: call count, success rate, p95 latency, last error, circuit-breaker state
- Integration-scoped audit: every external call logs to immutable audit trail with masked credentials
- Hot-swap credentials: rotate a key in Settings, live traffic fails over without downtime
- Backend: /routes/admin.py /admin/integrations/* + services/integration_service.py

Point MAARS at a product URL, a company, or a market segment and the Product Scanner builds a living intelligence file. It pulls pricing tiers, feature lists, recent announcements, social sentiment, funding rounds, review aggregates, and positioning language — then keeps refreshing on a schedule. Agents across the Growth, Sales, Research, and Strategic networks use these files to plan campaigns, write competitive battlecards, and flag market shifts.

HOW IT WORKS STEP BY STEP:

- Scans: product pages, pricing pages, About/company pages, Crunchbase/LinkedIn public data, review sites, social mentions
- Extracts: pricing tiers, feature matrix, recent product updates, team size, funding stage, Glassdoor ratings, trust signals
- Living files: schedule a re-scan daily/weekly/monthly; diff detector surfaces material changes
- Competitive battlecards: one click turns a product file into a battlecard (their claims → our response + proof points)
- 11 endpoints under `/api/products` for scan, catalog CRUD, subscribe-to-changes, export
- Integrated with CampaignBuilder and ContentGenerator so competitive insights flow directly into generated content
- Backend: `backend/services/product_scanner.py`

MAARS Command can run as the OS for a whole portfolio of ventures — each with its own agent team, budget, KPIs, stage gate, and runway. The Enterprise layer tracks revenue/burn/retention/NPS per venture, enforces stage gates (idea → MVP → PMF → scale → profit), and surfaces cross-venture benchmarks. If a venture misses two consecutive stage gates, the portfolio manager escalates a decision: double down, pivot, or sunset.

HOW IT WORKS STEP BY STEP:

- Per-venture: agent team, budget, KPIs, OKRs, stage, runway, founder, cap table stub
- Stage gates: 5 stages (Idea · MVP · PMF · Scale · Profit) with exit criteria per stage — enforced automatically
- KPI tracking: revenue, burn, runway, retention, NPS, CAC, LTV — pulled from connected integrations
- Cross-venture benchmarks: anonymous comparisons against portfolio peers (e.g., "your NPS is in the bottom quartile")
- Quality metrics framework: 22 endpoints under `/api/enterprise` covering KPIs, collaborations, ventures, quality, approvals
- Decision intelligence: automated quarterly reviews produce double-down / pivot / sunset recommendations with evidence
- UI: `VenturePortfolio.jsx` with portfolio grid, stage gates, KPI charts, and decision inbox

MAARS ships with a massive library of 58,422 Claude Code skills at `.claude/skills/`. Each skill is a self-contained Markdown document with a YAML front-matter (name + description) and optional reference files (scripts, docs, evals). When you create a new agent, a keyword-matcher instantly picks the skills whose topics align with that agent's role, description, and capabilities, then chunks the Markdown and writes the chunks into the MongoDB `knowledge_chunks` collection so the agent can retrieve them during chat (RAG). You can also bulk-ingest the entire library for all agents at once via the ingest script. The original unfixed copies are kept at `.claude/skills_backup/` as a safety net.

HOW IT WORKS STEP BY STEP:

- 58,422 SKILL.md files, each with valid YAML front-matter — covers A/B testing, accessibility, ad creative, aluminum rolling, animal handling, Axum, AWS, booking flows, carbon-lang, creative frontends, data arch, delivery, delon, DevOps docs, drift check, frugal rerouter, healthcare, hyper-casual games, Kubernetes, LinkedIn ads, MkDocs, NestJS, Next.js, Norway roads, PostgreSQL, React, reviewing code, Salesforce, Vue, and tens of thousands more
- Automatic matching on agent creation: `services/skills_service.py` runs `match_skills_for_agent(agent)` against the agent's role, description, capabilities, network, and tags using `SKILL_KEYWORDS`, then calls `ensure_agent_skills()` to chunk and write the matched skills to `db.knowledge_chunks`
- Provider knowledge pack: `match_providers_for_agent()` also ships LLM-provider cheat sheets (from `data/provider_skills.py`) so agents know the strengths of every model they can call
- Zero-broken-YAML guarantee: every front-matter parses cleanly under `yaml.safe_load`; no stale "hooks" fields (hooks belong in `settings.json`, not skill frontmatter); description capped at 1024 chars and collapsed to one line
- Bulk ingest on demand: `cd backend && python scripts/ingest_skills.py` re-processes all agents — use `--force` to overwrite, `--agents-only` or `--providers-only` to narrow scope
- Knowledge chunks stored with `source="skill" + skill_name + doc_id` so you can audit exactly which skill-chunks each agent has in its brain
- RAG-ready: `routes/chats.py` retrieves matching chunks during chat so agents ground their replies in the skill library instead of hallucinating

- Per-agent brain viewer: GET `/agents/{agent_id}/brain` returns every knowledge chunk currently loaded for that agent — full transparency into what any agent "knows"
- Extensible: drop a new folder under `.claude/skills/<name>/` with a `SKILL.md`, and the next agent creation or ingest run picks it up automatically



The MAARS Universal Key is now fully pay-as-you-go. Top up any amount — no subscriptions required. 65% of every payment goes directly to AI model costs (OpenAI, Anthropic, Google, and 30 other providers). The remaining 35% covers platform infrastructure, routing, failover, monitoring, and support. Credits are issued instantly. Each credit $\approx$ one LLM API call, priced by the model used.

HOW IT WORKS STEP BY STEP:

- Pay any amount from \$5 to \$10,000 USD (or BDT equivalent) — no monthly commitment
- 65% allocation: goes to your AI model budget — funds actual API calls to OpenAI, Anthropic, Google, DeepSeek, Groq, etc.
- 35% allocation: covers MAARS infrastructure — routing layer, failover logic, monitoring dashboards, logging, support, and platform development
- Credits issued instantly: your AI budget is converted to MAARS credits at ~ 16.67 credits per \$1 AI budget
- Real-time tracking: watch your credit balance and USD AI budget update in real-time after every LLM call
- One key, every provider: your maars-sk-* key unlocks access to all 33 providers and 600+ models — smart routing picks the best model for each task
- Budget transparency: every top-up shows the exact split — how much goes to AI costs vs. platform. No hidden fees
- Transaction history: full log of every top-up with amount paid, AI budget allocated, platform portion, and credits issued
- Usage stats: see breakdowns by provider, model, task type, and quality tier — understand exactly how your credits are spent
- Top-up anytime: credits never expire during your subscription period. Top up \$10 today, \$500 next month — fully flexible



How MAARS Command Is Built

The technology behind the scenes

MAARS Command is built with three main layers that work together, like a three-layer cake. Each layer has a specific job and they communicate constantly to deliver a seamless experience.



Backend (The Engine)

This is the part that runs on the server and does all the heavy work behind the scenes. It processes every request, manages data, and coordinates all 458+ agents.

FastAPI (Python) -- the programming framework that handles all requests at high speed with async processing

MongoDB -- the database that stores all agents, teams, memories, and user data across 27+ collections

212+ API endpoints -- connection points that the frontend uses to get and send data for every feature

WebSocket -- a real-time connection for live updates like the Agent Activity Monitor and collaboration feeds

MAARS Universal API Gateway: 175,000+ models across 33 direct providers, OpenAI-compatible /v1/chat/completions, live model validation, auto-fallback routing, webhooks, and model comparison endpoint

Stripe Integration -- handles credit card payments, subscription billing, and cost tracking

Background Task Queue -- manages long-running operations like batch agent deployments and report generation

Circuit Breaker System -- automatically detects and isolates failing services to prevent cascade failures

Skill ingestion pipeline (services/skills_service.py) -- reads .claude/skills/{name}/SKILL.md files, chunks the Markdown, and upserts into db.knowledge_chunks per agent using keyword-based matching



Frontend (What You See)

This is the part that runs in your web browser -- all the pages, buttons, charts, and visuals that you interact with every day.

React 18 -- the framework that builds the interactive user interface with real-time state management

Tailwind CSS -- makes everything look beautiful with a modern, consistent dark-themed design system

Shadcn/UI -- pre-built, accessible components like buttons, cards, menus, dialogs, and data tables

50+ pages & dashboards -- every feature has its own dedicated page with custom layouts and interactions

Progressive Web App (PWA) -- can be installed on your phone or desktop like a native app with offline support

Command Palette + Voice -- search and navigate instantly by typing keyboard shortcuts or speaking commands

Real-Time Charts & Visualizations -- live data rendered using Recharts for trust analytics and cost tracking

Print-Ready Documentation -- the entire system documentation can be exported as a formatted document



AI Layer (The Brains)

This connects MAARS Command to all the AI providers that power the agents. It routes tasks to the best model automatically.

OpenAI -- GPT-5, GPT-4.1/mini/nano (1M ctx), O4/O3 reasoning, Whisper STT, GPT Image 1, Sora 2

Anthropic -- Claude Opus 4.6 & 4.5, Sonnet 4.6 & 3.7, Haiku 4.5 — 200K context, extended thinking

Google -- Gemini 2.5 Pro (2M ctx), Gemini 2.5 Flash, Gemma 3 open models (1B–27B)

xAI -- Grok-4 & Grok-4 Fast Reasoning (256K ctx), Grok-3 (1M ctx), Grok-2 Vision

Groq, Cerebras -- sub-second inference on Llama 4, Qwen QwQ-32B

Together AI (39), Fireworks AI (20), SambaNova (14), NVIDIA NIM (12), Novita AI (12), Lepton AI (8), Lambda Labs (6), Minimax AI (4), Inception AI (2), Arcee AI (3), Amazon Bedrock (4) -- direct provider hosting

DeepSeek, Mistral, Perplexity, Cohere, AI21, Moonshot, Qwen, Yi, Zhipu, Doubao, Hyperbolic, Upstage, Writer, HuggingFace, LLaMA API

Live model validation -- checks provider /models API before every call, auto-fallbacks on stale IDs



All 33 AI Providers & 609+ Models

Every AI brain available via the MAARS Universal API Gateway

MAARS connects to 33 different AI companies, giving you access to 175,609+ models via a single `maars-sk-*` API key. Each model has different strengths — some are super fast but less powerful, some are incredibly smart but cost more. The Smart Router picks the best one automatically, or call any model directly via the Developer Portal at [/developer](#).

🌟 01.AI / Yi

- Yi Lightning Economy
- Yi Large FC Standard
- Yi Medium 200K Long-Context

🌟 AI Media & Voice

- Nano Banana 2 Image Gen
- GPT Image 1 Image Gen
- DALL-E 3 Image Gen
- Sora 2 Video Gen
- ElevenLabs Voice
- Whisper STT

🌟 AI21 (Jamba)

- Jamba Large 1.7 Flagship
- Jamba Mini 1.7 Economy

🌟 Amazon Bedrock

- Nova Pro Flagship
- Nova Lite Fast
- Nova Micro Economy

🌟 Anthropic

- Claude Sonnet 4.6 Flagship
- Claude Opus 4.6 Premium
- Claude Haiku 4.5 Economy

🌟 Arcee AI

- Arcee Maestro Flagship
- Arcee Blaze Fast
- Arcee Spark Economy

🌟 ByteDance Doubao

- Doubao Pro 128K Flagship
- Doubao Pro 32K Standard
- Doubao Lite 128K Economy

🌟 Cerebras

- Llama 3.3 70B Fast
- Llama 3.1 70B Fast
- Qwen 3-32B Economy

🌟 Cohere

- Command R+ Flagship
- Command R Economy

🌟 DeepSeek

- DeepSeek Chat Economy
- DeepSeek Reasoner Reasoning

🌟 Fireworks AI

- Llama 4 Scout Economy
- Llama 4 Maverick Fast
- DeepSeek V3 Fast

🌟 Google Gemini

- Gemini 2.5 Flash Fast
- Gemini 2.5 Pro Flagship
- Gemini 3 Flash / Pro Fast

🌟 Groq (Llama 4)

- Llama 4 Scout Economy
- Llama 4 Maverick Fast
- Llama 3.3 70B Fast

🌟 HuggingFace Inference

- Microsoft Phi-4 Compact Reasoning
- Gemma 2 9B Economy
- Mistral 7B Economy

🌟 Hyperbolic

- Llama 3.1 405B Flagship
- DeepSeek-R1 Reasoning
- Llama 3.3 70B Standard

🌟 Inception AI

- Mercury Coder Small Code
- Mercury Coder Large Code

🌟 Lambda Labs

- Hermes 3 405B Flagship
- Llama 4 Scout Economy

🌟 Lepton AI

- Llama 4 Maverick Fast
- DeepSeek R1-0528 Reasoning

🌟 Meta Llama API

- Llama 4 Scout Standard
- Llama 4 Maverick Flagship
- Llama 3.3 70B Economy

🌟 Minimax AI

- MiniMax Text-01 Flagship
- MiniMax VL-01 Fast

🌟 Mistral AI

- Mistral Large Flagship
- Mistral Medium Fast

🌟 Moonshot AI (Kimi)

- Kimi Auto Fast
- Kimi 128K Standard

- Mistral Small **Economy**

- Kimi 32K **Economy**

Novita AI

- Llama 4 Maverick **Fast**
- Qwen 3 235B **Flagship**
- DeepSeek R1 **Reasoning**

Nvidia NIM

- Nemotron Ultra 253B **Premium**
- Nemotron Super 49B **Standard**
- Llama 3.3 70B **Fast**

OpenAI

- GPT-5 **Flagship**
- GPT-4.1 / Mini / Nano **Fast**
- GPT-4o **Fast**
- O4 / O4-mini **Reasoning**
- O3 / O3-mini **Reasoning**

Perplexity

- Sonar **Search**
- Sonar Pro **Research**

Qwen / Alibaba

- Qwen Max **Flagship**
- Qwen Plus **Standard**
- QwQ-32B **Reasoning**
- Qwen Turbo **Economy**

SambaNova

- Llama 4 Maverick **Fast**
- DeepSeek R1-0528 **Reasoning**
- Qwen 2.5 72B **Fast**

Together AI

- Llama 4 Maverick FP8 **Fast**
- Llama 3.3 70B Turbo **Fast**
- DeepSeek R1 **Reasoning**

Upstage Solar

- Solar Pro **Enterprise**
- Solar Mini **Economy**

Writer Palmyra

- Palmyra X 004 **Enterprise**
- Palmyra Med **Medical**
- Palmyra Fin **Finance**

xAI (Grok)

- Grok-4 **Flagship**
- Grok-4 Fast / Reasoning **Fast**
- Grok-3 **Premium**
- Grok-3 Mini **Economy**
- Grok-2 Vision **Fast**

Zhipu AI (GLM)

- GLM-4-Plus **Flagship**
- GLM-4-Air **Economy**
- GLM-Z1-Air **Reasoning**



Security & Governance

How MAARS Command keeps everything safe and controlled

Running 458+ AI agents requires strong safety controls. Without proper governance, agents could overspend, leak data, or take unintended actions. Here are the six security systems that keep everything safe, fair, and auditable at all times.



Role-Based Access Control (RBAC)

Not everyone should have access to everything. RBAC lets you set different permission levels for different users. Admins can do anything -- manage all agents, change settings, and view billing. Managers can manage agents and view reports but cannot change system settings. Regular users can chat with agents and use tools but cannot modify configurations. Every page and every API endpoint checks permissions before allowing access. This prevents unauthorized users from accidentally (or intentionally) changing critical settings or accessing sensitive data.



Circuit Breakers

If an AI model or service starts failing repeatedly, the circuit breaker automatically stops sending requests to it. This prevents one broken service from crashing the entire system -- like a fuse box in your house that trips to protect your appliances. The system monitors error rates in real-time. When failures exceed the threshold, the circuit 'opens' and redirects traffic to backup models. Once the original service recovers and passes health checks, the circuit breaker automatically re-enables it. All transitions are logged for review.



Cost Governance

AI models cost money every time they're used -- some models cost just \$0.10 per million tokens, while others cost \$75 per million. Cost Governance tracks every single API call and its exact cost in real-time with per-agent and per-user breakdowns. You can set daily, weekly, or monthly spending limits for individual users and for the entire organization. When spending approaches the limit (80%), the system sends warning alerts. When limits are reached (100%), it automatically stops sending requests to paid models. This prevents surprise bills and ensures your AI budget is always under control.



Audit Logging

Every action in MAARS Command is recorded in an immutable audit log -- who did what, when they did it, and what the result was. This includes agent task completions, user logins, configuration changes, permission modifications, and API calls to external services. These logs cannot be edited or deleted after creation, making them tamper-proof and suitable for compliance audits, legal requirements, and internal investigations. Logs can be searched, filtered, and exported for reporting. Retention policies ensure logs are kept for the required duration.



Trust Scoring

Every agent gets a trust score from 0 to 100 based on its complete performance history. The score factors in task completion rate, response quality, speed, error frequency, and user feedback. Agents that consistently deliver accurate, fast results earn higher scores over time. Agents that fail often, produce low-quality work, or generate errors get lower scores. Low-trust agents can be automatically restricted to simulation-only mode or flagged for human review. High-trust agents get priority assignment for critical tasks. The scoring algorithm is transparent -- you can see exactly why an agent has its current score.



Simulation Mode

A system-wide safety switch that controls whether agents can take real-world actions. When Simulation Mode is ON (the default), all external actions -- sending emails, creating calendar events, making API calls, posting content -- return realistic fake responses instead of actually executing. This lets you safely test and review exactly what agents would do before going live. You can review simulated results, adjust agent configurations, and only flip to Execution Mode when you're confident everything works correctly. Individual agents can also be toggled independently between simulation and execution modes.



Why MAARS Command Is Different

What separates this from every other AI platform

Most AI platforms are either (a) one giant chatbot, (b) a skinny wrapper around OpenAI, or (c) a single-purpose agent (customer support, coding copilot, etc.). MAARS Command is built as an **enterprise operating system for autonomous AI** — the way a cloud platform is built for workloads, or an ERP is built for a business. Below are the nine things you get in MAARS that you cannot get by stitching together ChatGPT, a Zapier workflow, and a custom GPT.

458+ Agents in 27 Networks — not one chatbot

You don't prompt a generalist and hope it specializes. You dispatch work to the right specialist. A Commander coordinates. A CFO agent models the P&L. A Legal agent reads contracts. A Growth agent runs experiments. Each agent has its own brain profile, tool set, trust score, and memory — and they talk to each other through a real collaboration protocol, not fake multi-agent prompting.

Universal Gateway — 175,609+ models, one key

A single maars-sk-* API key reaches 33 providers (OpenAI, Anthropic, Google, xAI, DeepSeek, Mistral, Perplexity, Groq, Cerebras, plus 175,000+ HuggingFace models) through an OpenAI-compatible /v1/chat/completions endpoint. Smart routing, live model validation, auto-fallback, real token billing, webhooks. The other way to get this is to integrate 33 SDKs yourself.

58,422-Skill Knowledge Pack pre-built

Every agent is grounded in curated skill documents the moment it's created. You don't write system prompts from scratch. You don't hunt prompt libraries on Reddit. The keyword matcher picks the right skills per agent; the RAG engine retrieves them per turn; the brain endpoint shows you exactly what each agent "knows."

10-Tier Autonomy — not a binary switch

"Agent frameworks" tend to let agents run wild or demand approval for everything. MAARS has 10 granular tiers: pre-approved tools vs. gated tools, trust-scored elevation, dual-signature gates, policy-aware bypass. You can run one agent at Tier 3 (conservative) and another at Tier 9 (autonomous) in the same workflow.

Verification Civilization instead of "hope it's right"

Most AI output is trusted because the user has no other choice. MAARS cross-verifies claims across 2-4 independent models on 6 dimensions (factual, grounded, coherent, complete, bias, fresh) before the answer reaches you. Disagreements escalate. Hallucinations get quarantined. The agent's job isn't to sound confident — it's to survive attempted refutation.

4 Environments · Immutable Audit · Simulation mode

Simulation, Sandbox, Staging, Production — each with its own budget, policies, credentials, and retention. Run autonomous experiments at zero real cost in Simulation. Flip to Production only after the Verification engine signs off. Every action — approvals, tool calls, model responses, policy violations — written to an immutable audit trail. Enterprise-ready out of the box.

Circuit Breakers, Recovery, Incident Response built-in

Spend spike? Breaker trips, agents pause, you're paged. Agent misbehaves? It's quarantined — tools stripped, tier dropped to 1, flagged for review. Task graph fails halfway? Rollback to last checkpoint. Provider outage? Auto-fallback across 33 alternates. These are not bolt-ons — they are core kernel functionality.

Task Graph Kernel — real DAG execution

Give MAARS a goal in English; it emits a task graph (DAG) with parallel branches, dependency resolution, retry, checkpoints, and per-branch budgets. Kill it at hour 4 — resume at hour 5 from exactly where it stopped. This is how long-horizon autonomous work actually ships, not by stuffing everything into one mega-prompt and hoping.

Pay-as-you-go with 65%/35% transparency

No subscription lock-in. Top up any amount. Exactly 65% of every dollar goes to AI model costs (shown to you in real time), 35% covers the platform (routing, failover, governance, support). Credits never expire during your active period. Compare that to the "your model quota resets every month at 11:59 pm" world.

\$

Value Evaluation — What MAARS Saves You

Rough ROI per feature so you know where the money goes

Building the same capability stack by hand takes roughly 18-24 months of engineering plus ongoing integration/ops overhead. Below is a concrete per-feature breakdown of what each MAARS subsystem replaces and the rough monthly cost/effort avoided. Numbers are conservative mid-market estimates; enterprise stacks are typically 2-5× higher.

Capability	What It Replaces	Typical Monthly Cost (DIY)	Effort Avoided
Universal Gateway (33 providers)	33 separate LLM SDK integrations + routing layer	\$800–\$2,500	6–9 months eng
458+ Agent Workforce	Hiring 10-20 specialists + onboarding	\$80,000–\$200,000 (salary)	Ongoing
58,422 Skill Library	Custom prompt engineering per use case	\$3,000–\$15,000 (agency)	2–4 months
Task Graph Kernel (DAG)	Custom workflow engine (Temporal/Airflow + glue)	\$1,500–\$5,000 (infra + eng)	4–8 months eng
Verification Civilization	Manual QA review cycles + fact-check contractors	\$5,000–\$20,000	Ongoing
Hierarchical Memory (4 tiers)	Redis + Pinecone + Neo4j + custom orchestration	\$1,200–\$4,000 (infra)	3–6 months eng
10-Tier Autonomy + Approvals	Custom approval workflow + audit UI	\$2,000–\$8,000 (eng + tooling)	3–6 months eng
Multi-Environment Segregation	Separate AWS/GCP accounts + IAM + data isolation	\$1,000–\$5,000 (infra + compliance)	2–4 months eng
Circuit Breakers + Recovery	Custom SRE playbooks + monitoring + runbooks	\$3,000–\$10,000 (SRE ops)	Ongoing
Integration Hub (29 providers)	29 separate OAuth/API integrations + embedded browser + credential vault	\$2,000–\$6,000	6–12 months eng
Knowledge Base + RAG	Vector DB + chunking pipeline + UI + eval harness	\$800–\$3,000	3–5 months eng
Trust Scoring + Analytics	Custom scoring engine + dashboards	\$1,500–\$5,000	2–3 months eng
Product Scanner	Competitive intel tooling (Crayon, Klue, Kompyte)	\$1,500–\$6,000 (SaaS)	Ongoing
Content Generator + Reference Intelligence	Copy agency + brand voice consultancy	\$5,000–\$20,000	Ongoing
Social Media Command (10 platforms)	Hootsuite/Sprout + SMM agency	\$3,000–\$15,000	Ongoing
Vibe Coding App Builder	Junior dev or low-code platform (Retool, Bubble)	\$2,000–\$10,000	Ongoing
Voice Command + STT/TTS	ElevenLabs + Whisper + custom UI wiring	\$400–\$1,500	1–2 months eng
Observability + Metrics + Alerts	Datadog + PagerDuty + custom dashboards	\$2,000–\$8,000	Ongoing
Developer Portal + OpenAI-compatible API	Building an LLM reseller gateway	\$3,000–\$12,000 (eng + infra)	8–12 months eng
Venture Portfolio + KPIs	Portfolio ops tooling (Visible, Carta, manual)	\$1,000–\$4,000	Ongoing

Aggregate estimate: a comparable DIY stack runs **\$110,000–\$355,000/month** in tooling + salary, plus **18–24 months** of engineering to integrate. MAARS Command collapses that into one subscription with one API key.

✨

Total System Value — What MAARS Command Is Worth

Adding up every replaced subsystem, every avoided hire, every month of engineering collapsed into a turnkey platform:

\$3.5M–\$8.5M

Build-From-Scratch Cost

8–15 engineers × 18–24 months

\$1.3M–\$4.3M

Annual Operating Cost

Tooling, infra, SaaS, salaries

\$7.4M–\$21.4M

3-Year TCO (DIY)

Build + 3 yrs ops + integration debt

18–24 months

Time-To-Market Saved

From zero to enterprise-ready

Conservative replacement value: \$7.4M over 3 years. This is the TCO to rebuild MAARS Command from scratch with a mid-market engineering team — architecture + integrations + governance + UI + testing + ongoing SRE — *before* the model-call bill itself.

Enterprise-grade replacement value: \$21.4M+ over 3 years. Organizations that need SOC2 compliance, multi-region deployment, per-tenant isolation, 24/7 on-call, and domain-specific verifier tuning typically land at the top of this band — some well above it.

Real strategic value » replacement cost. Time-to-market is the hidden multiplier: shipping an autonomous-AI platform 18–24 months faster than a competitor is worth more than any line-item on this page. For a venture studio running 3+ products, MAARS Command replaces an entire AI platform team plus the tool stack — effectively a **\$10M–\$30M+ internal program** substituted by one account.

Ranges are conservative mid-market estimates. Enterprise deployments (SOC2, multi-region, custom verifiers, dedicated infra) routinely land 2–5× higher. Replacement cost ≠ list price — MAARS Command sells for a small fraction of TCO because it spreads platform R&D across every customer.



What MAARS Command Is Worth

The whole-system valuation, seen through four lenses

["Value"] depends on who's asking. A CFO asks replacement cost. A buyer asks acquisition comp. An investor asks pre-money. A founder asks strategic worth. Here is MAARS Command valued through all four lenses, each with its own math and comparables — so whichever seat you're in, the number is grounded.

HEADLINE VALUATION

\$40M – \$180M

As-built worth of the MAARS Command platform today — centered on the fair-market band for a pre-revenue enterprise AI platform with 484 endpoints, 458 agents, 58,422 skills, 33 LLM providers, full governance + verification, and an OpenAI-compatible gateway. The floor is salvage/acqui-hire. The ceiling is strategic acquisition by an incumbent.

Replacement Cost

CFO / rebuild-from-scratch view

\$7.4M – \$21.4M

Cost to rebuild MAARS Command from zero: \$3.5M–\$8.5M to build (8–15 engineers × 18–24 months) plus \$1.3M–\$4.3M/yr operating. 3-year TCO = \$7.4M–\$21.4M before a single customer API call. Enterprise-grade (SOC2, multi-region, 24/7 on-call) lands at the top of the band or above.

Fair Market Value

Independent appraisal / asset transfer

\$15M – \$45M

Arms-length sale of the codebase + IP + agent library + skill pack + gateway, assuming a knowledgeable buyer and no special synergies. Premium over pure replacement because the bugs are already found, the integrations work, and the design decisions are made — a year of de-risking baked in.

Strategic Acquisition Value

Incumbent buyer (OpenAI, Anthropic, Salesforce, MSFT)

\$60M – \$180M

What a strategic buyer pays to acquire the platform plus the team. Comparable AI-platform acqui-hires and strategic deals in 2024–2025 range \$30M–\$250M at this feature depth (governance + verification + gateway is rare). A vertical player buying to add enterprise orchestration would pay top-of-band.

Pre-Money Venture Valuation

Seed / Series A investor

\$40M – \$120M pre

What a VC underwrites based on platform depth + TAM + team. AI-infrastructure rounds in 2024–2025 closed at \$40M–\$150M pre-money for platforms with this feature surface even pre-revenue. Early revenue traction (>\$1M ARR) pushes the top to \$250M+. Vertical SKU launches (MAARS-for-Law etc.) stack on.

Bottom line: MAARS Command is worth between **\$40M at the low end** (pre-revenue venture valuation or strategic floor) and **\$180M at the high end** (strategic acquisition by an incumbent who needs the platform + team). The numbers stack across lenses: rebuilding costs \$7M+, fair market sale clears \$15M+, a strategic buyer pays \$60M+, and a VC underwrites \$40M+ pre-money. Add early revenue (>\$1M ARR), a SOC2 cert, and one vertical SKU shipped — all reachable within 12 months — and the whole band shifts to **\$150M–\$500M+**.

All ranges are mid-2026 benchmarks. Private-market comps: Harvey AI (\$3B), Glean (\$4.6B), Writer (\$1.9B), Cresta (\$1.6B), Decagon (\$1.5B), Sierra (\$4.5B) — all AI-platform categories MAARS competes in. MAARS is earlier stage and pre-revenue, which is why the band starts at \$40M, not \$1B+.



MAARS Command — Pitch Deck

14 slides · the investor story, at a glance

Fundraising-ready deck summarizing the business case in 14 slides. Each slide is a self-contained card — screenshot, paste into a Google Slides template, or use as speaker notes. All numbers are grounded in the platform you just read about.

01

Cover
Autonomous AI Enterprise Operating System

MAARS Command

The operating system for autonomous AI workforces. 458 agents across 28 networks, 58,422 curated skills, 33 LLM providers, one API key.

MAARS Global Corporation · Est. 2026 · Seeking Seed / Series A

02

The Problem
Building enterprise AI is 18–24 months of integration hell

- Every org wants AI agents. Nobody wants to build the plumbing.
- 33 LLM providers, 29 integrations, 4 memory tiers, governance, verification, billing — an 18–24 month build.
- DIY stacks cost **\$110k–\$355k/month** in tooling + salary.
- Meanwhile: chatbot wrappers hallucinate, no audit trail, no rollback, no cost controls, no trust.

03

The Solution
One platform. One key. Every capability enterprises need.

- **Task Graph Kernel** — DAG execution with checkpoints + rollback
- **Verification Civilization** — multi-model cross-check on 6 dimensions
- **10-Tier Autonomy** — human-approval → unsupervised, per agent
- **58,422 Skill Library** — grounded RAG, no prompt engineering
- **Universal Gateway** — OpenAI-compatible, 175,609+ models, one key

04

How It Works
From English goal to executed workflow

- ① You describe a goal in English →
- ② Kernel compiles a task graph (DAG) →
- ③ Router picks the best model per node →
- ④ Agents execute in parallel, grounded in skills + KB →
- ⑤ Verification cross-checks every claim →
- ⑥ Approval gates where human sign-off required →
- ⑦ Audit log records every decision, forever.

05

Market Size
Enterprise AI orchestration — emerging \$100B+ category

TAM — Global enterprise AI ops	\$150B by 2030
SAM — AI agent platforms	\$40B by 2028
SOM — Mid-market + venture studios	\$2.5B 3-yr reachable

Gartner, IDC, McKinsey 2024–2025 enterprise AI spend forecasts.

06

Product Depth

What's already built and shipping

458

AGENTS

27

NETWORKS

58,422

SKILLS

33

LLM PROVIDERS

175,609+

MODELS

484

API ENDPOINTS

62

FRONTEND PAGES

28

INTEGRATIONS

07

Why We Win

Moats that compound over time

- **Depth moat** — 35 core systems vs. competitors' 3–5
- **Data flywheel** — every verification pass is labeled data
- **Platform lock-in** — custom skills + KB + brain profiles create switching cost
- **Provider neutrality** — 33 providers = we win when models commoditize
- **Governance trust** — SOC2 roadmap unlocks enterprise contracts competitors can't touch

08

Business Model

Hybrid: subscription + PAYG + marketplace

- **PAYG credits** — 65% to AI costs, 35% platform take
- **Subscription tiers** — \$49/\$299/\$1,499/month user tiers
- **Enterprise** — \$50k–\$500k+ ARR (SOC2, dedicated env, SLA)
- **Marketplace (roadmap)** — 30% take on skill + agent + workflow sales
- **Gateway API** — developer revenue on 175k+ model calls

09

Traction

Platform shipped; commercial phase opening

- Platform code-complete — 522 endpoints, 62 pages, 29 integrations live (incl. self-hosted Embedded Browser)
- 58,422-skill library ingested + matched — largest known skill catalog
- Universal Gateway operational — OpenAI-compatible, live model validation
- Governance + Verification engines in production flow
- Early user cohort onboarding — feedback loop active
- *Seeking seed to accelerate GTM, compliance, and vertical SKUs.*

10

Competitive Landscape

Owning the orchestration layer above every model

LangChain / LlamaIndex

Dev library · no governance, no UI

CrewAI / AutoGen

Agent framework · no platform

Zapier / Make

No-code · no autonomy, no agents

Glean / Writer / Harvey

Vertical SaaS · single use case

Salesforce Agentforce

Bolt-on · locked in SF ecosystem

MAARS = platform depth of all five combined, provider-neutral.

11

Go-to-Market

Three wedges, sequenced for compounding distribution

- Wedge 1 — Developers — Universal Gateway with maars-sk-* key; bottom-up adoption via OpenAI compatibility
- Wedge 2 — Mid-market — \$299/\$1,499 tiers; Campaign Builder + Content Generator + Social Command as acquisition magnets
- Wedge 3 — Enterprise — SOC2 + vertical SKUs (Law, Finance, Healthcare); \$50k–\$500k ARR deals
- Marketplace unlocks in year 2 — 3rd-party creators drive network effects

12

Financial Projections

3-year ARR trajectory (base case)

Year 1 — Dev + mid-market

\$1.2M ARR

Year 2 — SOC2 + first vertical

\$8M ARR

Year 3 — Enterprise + marketplace

\$35M ARR

Gross margin target

65–72%

Comp benchmarks: Harvey \$75M→\$250M ARR Y2→Y3, Writer similar trajectory.

13

Valuation

As-built worth · with traction uplift

Today (pre-revenue, as-built)

\$40M–\$180M

+12mo (ARR + SOC2 + vertical)

\$150M–\$500M

Category comps (2024–25)

\$1B–\$4.6B

Category: Harvey (\$3B), Glean (\$4.6B), Writer (\$1.9B), Sierra (\$4.5B), Decagon (\$1.5B).

14

The Ask

Raising to accelerate GTM and compliance

\$8M Seed / Series A

18-month runway · use of funds:

- 40% — GTM + sales team (mid-market + enterprise wedges)
- 25% — SOC2 + HIPAA + first vertical SKU
- 20% — Engineering — marketplace, SDKs, mobile
- 10% — Brand + content + category creation
- 5% — Compliance + ops + legal

Target milestone: \$8M ARR + SOC2 + 50 enterprise logos in 18 months.



A-Z Feature Index

Every feature, alphabetized — use Ctrl+F to jump

Quick reference of every user-facing feature in MAARS Command, alphabetized. Each entry is a short description so you know what it does at a glance.

Admin Dashboard	Central admin hub — users, pricing, integrations, branding, analytics
Agent Activity Monitor	Live WebSocket feed of every agent's state, messages, and tool calls
Agent Brain Viewer	GET /agents/{id}/brain — see every knowledge chunk loaded for any agent
Agent Catalog	Browse all 458+ agents with network, role, and capability filters
Agent Networks	28 domain categories (Strategic, Engineering, Creative, Growth, Core Team, etc.)
Agent Team Builder	Group any subset of agents into custom named teams
Analytics Dashboard	Usage, cost, trust, performance — per user, agent, network, env
API Keys	Issue/rotate maars-sk-* keys for developer access to the Universal Gateway
Approvals	Human-in-loop queue for gated actions — approve, reject, edit, escalate
Audit Log	Immutable record of every action — tamper-proof, queryable, exportable
Autonomy Tiers	10 levels from full-approval (T1) to unsupervised (T10) per agent
Avatar Generator	Admin batch-generates agent avatars (unique per agent)
Brain Profiles	Per-agent model, tone, temperature, max-tokens, tool permissions
Branding	Admin configures logo, domain, theme — multi-tenant ready
Browser — Embedded	Self-hosted Chromium under backend/browser_data/ — lives inside MAARS, never at OS user level
Browser — Multi-Tab	Chrome-style tab strip; agents, system, and user share the same tabs with three-way handoff
Browser — Fullscreen	One-click fullscreen (Fullscreen API) — canvas scales to full viewport
Browser — Stream Modes	PNG/JPEG × event-driven or continuous 1-24 fps; persists per user
Browser Agent (Autonomous)	Natural-language goal → perceive (vision) → act → repeat. 2FA / CAPTCHA hand-off
Browser Vision	GPT-4o / Claude multi-modal decides next action from screenshot + goal
Browser Integration Connect	Click any of 29 integrations — opens login/OAuth in tagged tab, agent drives optionally
Browser Domain Policy	Per-env allow/deny hostname patterns; default blocks localhost/cloud-metadata
Browser Minutes Budget	Per-user daily cap (default 60 min); admin-overridable, 429 on exhaust
Budget Controller	Per-user, per-agent, per-env credit caps with soft/hard limits
Campaign Builder	Multi-channel campaign planning + AI-generated variants + A/B
Catalog Manager	458+ agent registry with network lookup and hot-swap
Chat History	Persistent chat threads per agent with search and PDF export
Circuit Breakers	Auto-trip on spend spike, error rate, retry storm, policy violation
Code Explorer (Admin)	Browse every backend/frontend file — view, search, ZIP export
Collaboration Engine	Agent-to-agent info share, review request, handoff, coordination
Command Palette	Press / or Ctrl+K — search pages, agents, actions with voice input
Content Generator	Marketing copy, social posts, emails, ads, press releases, blueprints
Cost Governance	Real-time per-call cost tracking with alert thresholds (80/100%)
Credits	Pay-as-you-go credit balance with instant top-up and transaction log
Custom Brain	Per-user custom instructions attached to agent knowledge

Developer Portal	/developer — API key, model browser, playground, quickstart
Environments	Simulation · Sandbox · Staging · Production — isolated policies/budgets
Execution Mode	System-wide toggle: real external actions ON vs. mocked (Simulation)
Gmail Integration	OAuth-based real email sending through connected Gmail account
Google Calendar	OAuth-based event creation, updates, and scheduling
Image Generation	GPT Image 1, DALL-E 3, Nano Banana (Gemini 3.1 Flash image)
Incident Ledger	Every fault logged with timeline, severity, blast radius, remediation
Ingest Script	backend/scripts/ingest_skills.py — bulk re-ingest skills to all agents
Insights Page	Personal analytics — your usage, best agents, cost breakdown
Integration Hub	29 providers: payments, email, CRM, ecommerce, social, dev, LLM, embedded browser
Kernel Dashboard	Task graphs, goals, scheduler, budget, policies — one control plane
Knowledge Base	Upload PDFs/DOCX/MD per agent; chunked, embedded, cited in replies
Knowledge Graph	Entity-relationship memory tier — traversable, queryable
KPI Dashboard	Venture/org-level KPIs: revenue, burn, retention, NPS, CAC, LTV
LLM Router	Classifies task type + budget → picks best model across 33 providers
Media Upload	Files, audio, video uploads with STT transcription (Whisper)
Memory Governance	CRUD memories with category tags, importance, decay, pruning
Memory Hierarchy	4 tiers: working · episodic · semantic · knowledge graph
Messenger Chat	Floating assistant widget on every page
Model Comparison	POST /v1/models/compare — run same prompt on up to 4 models
Model Router Dashboard	Live routing decisions, classification accuracy, per-model latency
Notifications	Real-time alerts for task completion, approvals, incidents, budget
OAuth	Google, GitHub OAuth flows with secure credential vault
Observability Dashboard	Live metrics, alerts, trace logs, per-service health
Operator Control Panel	Low-level kernel controls — halt, resume, reassign, escalate
Organization	Org-level settings, teams, roles, permissions
Payments	Stripe checkout, subscriptions, PAYG credit top-ups, invoices
Personal Secretary	Dedicated agent that executes real-world actions on your behalf
Policy Engine	Governance rules, compliance, autonomy enforcement per env
Preview Mode	Preview a workflow before executing in Production
Pricing Admin	Plan editor — tiers, credits, overage rules, per-provider markup
Product Scanner	Real-time competitive intelligence — pricing, features, sentiment
Project Catalog	Browse all active projects with phase, owner, timeline
Projects	Multi-task project lifecycle with phases and active summary
Quality Service	Pass rates, escalation tracking, auto-review of agent outputs
RBAC	Role-based access — admin · manager · user with page/route gates
Real-World Actions	Gmail, Calendar, Twilio, SendGrid — real external calls when enabled
Recovery	Quarantine · rollback · retry — automatic blast-radius containment
Reference Intelligence	Paste text/image → Style Blueprint for on-brand content generation

RAG Engine	Retrieval with chunk-level citations across skills, KB, web
Settings	User preferences, integrations, notifications, default model/tier
Simulation Mode	All external actions mocked — zero cost, full logging, chaos-safe
Skill Library	58,422 SKILL.md files auto-matched to agents at creation
SMTP Config	Admin-configurable SMTP for transactional email
Social Media Command	10 platforms: post, boost, DM, cold-email, cold-call, translate
Stats (Admin)	Platform-wide analytics: MRR, DAU, cost, top agents, top errors
STT / Whisper	Speech-to-text in 50+ languages via OpenAI Whisper
Subscriptions	Monthly plans, PAYG, hybrid billing with fallover
Task Graph	DAG execution with parallel branches, checkpoints, resumable
Tasks	Task creation, assignment, status, project linking
Team Builder	Group agents into named teams with missions
Teams	User team management with invites and roles
Test Harness	Scenario testing for autonomous workflows in Simulation
Tool Registry	Tool discovery, validation, sandboxing across 70+ built-in tools
Trust Analytics	0-100 trust score per agent — success, latency, quality, consistency
TTS / ElevenLabs	Text-to-speech in multiple languages including Bangla + English
Universal Gateway Key	One maars-sk-* key for all 175,609+ models, OpenAI-compatible
Universal Reference Intelligence	Extract style from any text/image → reusable Blueprint
Usage Stats	Per-user, per-agent, per-provider, per-task-type cost breakdowns
Venture Portfolio	Multi-product tracking with stage gates and cross-venture benchmarks
Verification Civilization	Multi-verifier fact-checking with 6-dimensional consensus
Vibe Coding	Build apps by talking — HTML/CSS/JS generated live with preview
Video Generation	Sora 2 — 4-12s AI video from text descriptions
Voice Commands	Microphone-based navigation and search via Whisper
Web Search	Perplexity-grounded live web search embedded in every chat
WebSocket Streams	Live activity feed + Infinity execution streaming
Webhooks	Signed HMAC-SHA256 webhooks for budget, rate limit, completion events
Workers	Job queue with per-worker stats, retry logic, priority levels
Workflow Builder	Visual drag-drop flowchart for automated multi-step workflows
Workspace	Per-user brain, artifacts, tool logs, KPIs, execution gateway

37 core systems · 522 API endpoints · 62 pages · 84+ UI components · 29 external integrations (incl. Embedded Browser) · 33 LLM providers · 175,609+ models · 58,422 skills · 458+ agents across 28 networks.



Roadmap — What's Shipped, In-Flight, and Next

12-month execution plan across product, compliance, and GTM

Where MAARS Command is today and where it's going next. Shipped items are live in production. In-flight items are in active development. Next items are committed but not yet started.

Shipped

✓ Live

- Task Graph Execution Kernel (131 endpoints)
- 458+ agents across 28 networks
- 58,422-skill library + RAG ingestion
- 33 LLM providers · OpenAI-compat gateway
- 10-tier autonomy + approval workflow
- Verification Civilization (6-dim scoring)
- Hierarchical memory (4 tiers)
- Multi-environment segregation (4 envs)
- Circuit breakers + recovery + incidents
- 29 external integrations (OAuth + API)
- Embedded Browser runtime (multi-tab, fullscreen, agent-driven)
- Autonomous BrowserAgent with 2FA hand-off + vision (GPT-4o / Claude)
- Vibe Coding · Content Gen · Campaign Builder
- Social Media Command (10 platforms)
- PAYG + subscription billing
- Developer Portal + maars-sk-* keys
- 522 API endpoints · 62 frontend pages

In-Flight (Q1-Q2)

In Progress

- SOC2 Type I audit — observation period active
- Skill Marketplace (beta) — creator revenue share
- Mobile PWA optimization for tablet/phone
- Per-tenant fine-tuned verifiers
- Advanced cost optimization (route-by-price)
- Enterprise SSO (Okta, Azure AD, Google Workspace)
- Dedicated VPC deployment option
- Webhook subscription library
- Expanded Slack + Teams native integration
- Voice agent (phone + ElevenLabs live calls)
- Real-time collaboration on task graphs
- Benchmark leaderboard (public)

Next (Q3-Q4)

Planned

- SOC2 Type II certification
- HIPAA compliance + healthcare SKU
- Vertical SKUs: MAARS-for-Law, MAARS-for-Finance
- Agent Marketplace (third-party certified agents)
- Workflow Marketplace (template commerce)
- Native iOS + Android apps
- EU data residency (Frankfurt region)
- GDPR DPA templates + CCPA toolkit
- GitHub App — MAARS as a PR reviewer
- ISO 27001 kickoff
- White-label / OEM partner program
- Verification-as-a-Service (/v1/verify public)
- Trust-Score-as-a-Service API

- Academic research partnerships + eval dataset publishing

Roadmap items are directional commitments; exact timing depends on enterprise customer feedback, compliance auditor scheduling, and ecosystem readiness.



Security & Compliance — Current Status

Which certifications and controls are live, in-progress, or planned

Enterprise buyers need evidence, not promises. Here's exactly where MAARS stands on every major compliance framework and security control today.

SOC 2 Type I Observation period active · target audit completion Q2	In Progress
SOC 2 Type II Q4 after Type I; 6-month continuous-controls window	Planned
GDPR DPA template available · data-export + right-to-erasure live	Ready
CCPA California resident opt-out flow implemented	Ready
HIPAA BAA-ready architecture · full compliance with healthcare SKU	Planned
ISO 27001 Kickoff planned post-SOC2 Type II	Planned
PCI-DSS All card data handled by Stripe (PCI-DSS Level 1)	Delegated
Data Encryption TLS 1.3 in transit · AES-256 at rest · hardware-backed keys	Live
Audit Logging Immutable · 7-yr retention (prod) · full who/what/when	Live
RBAC Admin · Manager · User · Custom roles · route + action gated	Live
Credential Vault Encrypted secrets store · hot-rotate without downtime	Live
Penetration Testing Annual third-party pentest · next cycle pre-SOC2 audit	Scheduled
Incident Response Playbooks · on-call rotation · incident ledger · RCA process	Live
SSO (Okta/Azure AD) Enterprise SSO beta Q2; SAML + OIDC	In Progress
Data Residency US-only today · EU (Frankfurt) region Q4	Planned
DDoS Protection	Live

Cloud-provider WAF · rate-limit per API key · circuit breakers

Backup & DR

Live

Hourly snapshots · multi-AZ · tested quarterly

Zero-Trust Network

Live

Per-service mTLS · no flat-network access · JIT elevation

Live = production, auditable today. In Progress = active work toward completion. Planned = committed, scheduled, not started. Delegated = handled by a certified third party.



FAQ — Common Questions

What buyers, investors, and developers ask most often

The questions we get every week. If you have one that isn't here, email support.maars@marsgc.net and we'll add it.

How is this different from LangChain, CrewAI, or building on OpenAI directly?

LangChain and CrewAI are libraries — you still build the platform around them: governance, verification, billing, UI, audit, recovery. Building on OpenAI directly means you're locked to one provider and you build everything else yourself. MAARS gives you the whole enterprise platform (governance, verification, multi-provider routing, UI, audit, recovery, skill library, 458 agents) as one product. Days to value, not quarters.

Can I bring my own API keys and avoid your gateway markup?

Yes. Every LLM provider can be configured with your own API key in Settings → Integrations. You pay the provider directly and MAARS just routes through. The Universal Gateway with `maars-sk-*` key is a convenience for teams who don't want to manage 33 provider relationships — but it's optional.

What happens when a model provider goes down?

Circuit breakers detect elevated error rates or latency and trip. The router automatically falls back through a configured chain (e.g. GPT-5 → Claude Sonnet → Gemini Pro → GPT-4o). If all fallbacks fail, the task graph pauses at the affected node and can resume when a provider recovers — no work lost.

How do you handle hallucinations?

Every factual claim goes through the Verification Civilization: 2-4 independent verifier models cross-check the same claim on 6 dimensions (factual, grounded, coherent, complete, unbiased, time-fresh). Disagreement flags the claim. Unverifiable specifics get quarantined. The agent's answer ships with the verification score and source citations. No blind trust.

Where is data stored? Is it used to train models?

Production data lives in the MongoDB cluster in the customer's region (US today; EU region Q4). Nothing is used to train third-party models — every provider is called in data-processor mode with zero-retention headers. Chat history is retained per your settings and fully exportable or deletable on demand (GDPR-compliant).

Can I self-host or run on my own cloud?

Docker Compose deployment is live today (docker-compose up — full stack in one command). Dedicated VPC deployment (your AWS/GCP/Azure account with MAARS-managed plane) is in-flight for Q2. True air-gapped on-prem is available for enterprise contracts.

How much does it cost for a realistic workload?

Pay-as-you-go pricing: 65% of every dollar goes to AI model costs (shown live), 35% covers platform. A typical mid-market team with 3-5 active workflows and moderate chat use lands at \$500-\$2,500/month. Enterprise contracts (SOC2, dedicated env, SLA) start at \$50k ARR. Developer gateway usage is metered per call at near-pass-through pricing.

Can agents take real-world actions (send emails, make charges, post to social)?

Yes — through the Real-World Action Layer. But only when you flip Simulation Mode off. By default every external action returns a realistic mock response so you can safely test a workflow end-to-end before going live. Per-agent autonomy tiers gate what each agent can do, and high-impact actions require human approval.

What happens if an agent misbehaves?

The Recovery system quarantines it: tools stripped, autonomy forced to Tier 1, flagged for human review. The incident ledger logs the full timeline. You decide whether to reinstate, retrain, or retire the agent. Your workflows auto-reroute through healthy agents in the meantime.

How do I add a custom skill or custom agent?

Custom skills: drop a new folder at `.claude/skills/<name>/` with a SKILL.md file — it's ingested automatically on agent creation or via the bulk ingest script. Custom agents: POST /agents with role, description, capabilities, and optional brain profile. The skill matcher runs automatically and populates the new agent's brain.

Is there an OpenAI-compatible API I can point existing code at?

Yes. Your maars-sk-* key works at POST /v1/chat/completions with the exact OpenAI SDK — only change the base_url. You get access to all 175,609+ models (609 curated + 175k HuggingFace pass-through) through the same SDK you already use. Model comparison endpoint (/v1/models/compare) runs the same prompt on 4 models in parallel.

What's the team's AI background?

MAARS was built by MAARS Global Corporation with a team combining full-stack engineering, AI research, and enterprise operations experience. The platform has been in active development since 2026 with focus on governance, verification, and multi-provider orchestration. Team details available under NDA for qualified investors.



Get Started

Book a demo, join the developer waitlist, or talk to enterprise sales

Book a Demo

30-minute walkthrough of MAARS Command with a product specialist. Best for evaluating fit for your team.

[Email to schedule](#) >

Developer Access

Get a maars-sk-* key and start calling 175,609+ models through the OpenAI-compatible gateway today.

[Open Developer Portal](#) >

Enterprise Sales

SOC2 roadmap, dedicated VPC, BAA, SLA, custom training, and volume pricing. Deals starting at \$50k ARR.

[Contact enterprise](#) >

 support.maars@marsgc.net  enterprise.maars@marsgc.net  MAARS Global Corporation  Est. 2026

MAARS Command v1.0 -- Autonomous AI Enterprise Operating System
499+ Agents | 28 Networks | 37 Core Systems | 33 Providers | 600+ Models | 58,422 Skills | 522+ Endpoints
Built by MAARS Global Corporation | Est. 2026
Contact: support.maars@marsgc.net